

Introduction

1.1 Overview

This specification provides the common definitions for its companion specifications, CAE Specification, **Commands and Utilities, Issue 5** and CAE Specification, **System Interfaces and Headers, Issue 5** (see **Referenced Documents** on page xiv). It defines general terms, concepts and interfaces used by both other volumes. Thus, this volume is a prerequisite for understanding either of the other two.

1.2 Terminology

The following terms are used in this specification:

can

This describes a permissible optional feature or behaviour available to the user or application; all systems support such features or behaviour as mandatory requirements.

implementation-dependent

The value or behaviour is not consistent across all implementations. The provider of an implementation normally documents the requirements for correct program construction and correct data in the use of that value or behaviour. When the value or behaviour in the implementation is designed to be variable or customisable on each instantiation of the system, the provider of the implementation normally documents the nature and permissible ranges of this variation. Applications that are intended to be portable must not rely on implementation-dependent values or behaviour.

legacy

Certain features are *legacy*, which means that they are being retained for compatibility with older applications, but have limitations which make them inappropriate for developing portable applications. New applications should use alternative means of obtaining equivalent functionality. Legacy features are marked **LEGACY**.

may

With respect to implementations, the feature or behaviour is optional. Applications should not rely on the existence of the feature. To avoid ambiguity, the reverse sense of *may* is expressed as *need not*, instead of *may not*.

must

This describes a requirement on the application or user.

should

With respect to implementations, the feature is recommended, but it is not mandatory. Applications should not rely on the existence of the feature.

With respect to users or applications, the word means recommended programming practice that is necessary for maximum portability.

undefined

A value or behaviour is undefined if this document imposes no portability requirements on applications for erroneous program constructs or erroneous data. Implementations may specify the result of using that value or causing that behaviour, but such specifications are not guaranteed to be consistent across all implementations. An application using such behaviour is not fully portable to all systems.

unspecified

A value or behaviour is unspecified if this document imposes no portability requirements on applications for correct program construct or correct data. Implementations may specify the result of using that value or causing that behaviour, but such specifications are not guaranteed to be consistent across all implementations. An application requiring a specific behaviour, rather than tolerating any behaviour when using that functionality, is not fully portable to all systems.

will

This means that the behaviour described is a requirement on the implementation and applications can rely on its existence.

1.3 Portability

Some of the utilities in CAE Specification, **Commands and Utilities, Issue 5** and functions in CAE Specification, **System Interfaces and Headers, Issue 5** describe functionality that might not be fully portable to systems based on the ISO POSIX-1 or ISO POSIX-2 standards. Where enhanced or reduced functionality is specified, the text is shaded and a code in the margin identifies the nature of the extension or warning (see **Codes**). For maximum portability, an application should avoid such functionality.

Unless the primary task of a utility is to produce textual material on its standard output, application developers should not rely on the format or content of any such material that may be produced. Where the primary task is to provide such material, but the output format is incompletely specified, the description is marked. Application developers are warned not to expect that the output of such an interface on one system will be any guide to its behaviour on another system.

Codes

The codes and their meanings are as follows:

EX

Extension.

The functionality described is an extension to the standards referenced above. Application writers may confidently make use of an extension as it will be supported on all XSI-conformant systems. These extensions are designed not to conflict with the published standards.

If an entire **SYNOPSIS** section is shaded and marked with one EX, all the functionality described in that entry is an extension.

Some behaviour which is allowed to be optional in the formal standards is mandated on XSI-conformant systems. Such behaviours (for example, those dependent on the availability of job control) might not be individually marked as extensions, but the mandatory nature of the feature is marked as an extension where the option is described, typically in the header where the corresponding symbolic constant is defined.

79	FIPS	FIPS Requirements.
80		The Federal Information Processing Standards (FIPS) are a series of U.S. government
81		procurement standards managed and maintained on behalf of the U.S. Department of
82		Commerce by the National Institute of Standards and Technology (NIST). Where restrictions
83		have been made in order to align with the FIPS requirements, they have the special mark shown
84		here, and appear in the index under FIPS alignment (as well as under EX).
85		The following restrictions are required by FIPS 151-2:
86		• The implementation will support {_POSIX_CHOWN_RESTRICTED}.
87		• The limit {NGROUPS_MAX} will be greater than or equal to 8.
88		• The implementation will support the setting of the group ID of a file (when it is created) to
89		that of the parent directory.
90		• The implementation will support {_POSIX_SAVED_IDS}.
91		• The implementation will support {_POSIX_VDISABLE}.
92		• The implementation will support {_POSIX_JOB_CONTROL}.
93		• The implementation will support {_POSIX_NO_TRUNC}.
94		• The <i>read()</i> call returns the number of bytes read when interrupted by a signal and will not
95		return -1.
96		• The <i>write()</i> call returns the number of bytes written when interrupted by a signal and will
97		not return -1.
98		• In the environment for the login shell, the environment variables <i>LOGNAME</i> and <i>HOME</i> will
99		be defined and have the properties described in Chapter 5 of this document.
100		• The value of {CHILD_MAX} will be greater than or equal to 25.
101		• The value of {OPEN_MAX} will be greater than or equal to 20.
102		• The implementation will support the functionality associated with the symbols CS7, CS8,
103		CSTOPB, PARODD and PARENB defined in <termios.h>.
104	JC	Job Control Extension.
105		Job control is an optional feature in the operating system described by the ISO POSIX-1
106		standard, but it is supported by all XSI-conformant systems. When interfaces rely on this
107		extension, they have the special mark shown here and appear in the index under JC (in addition
108		to being under EX).
109	OB	Obsolescent.
110		Some of the interfaces describe functionality that is obsolescent. Although these are fully
111		portable to all current XSI-conformant systems they may be withdrawn in future issues.
112	OF	Output format incompletely specified.
113		The format of the output produced by the utility is not fully specified. It is therefore not possible
114		to post-process this output in a consistent fashion. Typical problems include unknown length of
115		strings and unspecified field delimiters.

116	OH	Optional header.
117		In the SYNOPSIS section of some interfaces in CAE Specification, System Interfaces and
118		Headers, Issue 5 an included header is marked as in the following example:
119	OH	<code>#include <sys/types.h></code>
120		<code>#include <grp.h></code>
121		<code>struct group *getgrnam(const char *name);</code>
122		This indicates that the marked header is not required on XSI-conformant systems. This is an
123		extension to certain formal standards where the full synopsis is required.
124	OP	Dependent on optional service in XSI.
125		Typical implementations depend on an optional service and the functionality affected need not
126		be present if the optional service is not supported.
127	PI	The behaviour cannot be guaranteed to be consistent.
128		It is not possible to guarantee that the interface behaves in the same way on all XSI-conformant
129		systems. This is the case if it provides functionality that is system-defined or system-specific.
130		Options that are used to <i>select</i> alternative forms of system-specific behaviour are not marked, as
131		it is clear from their descriptions that their use is inherently non-portable.
132	RT	Realtime.
133		This identifies the interfaces and additional semantics in the Realtime Feature Group.
134	RTT	Realtime Threads.
135		This identifies the interfaces and additional semantics in the Realtime Threads Feature Group.
136	UN	Possibly unsupportable feature.
137		It need not be possible to implement the required functionality (as defined) on all XSI-
138		conformant systems and the functionality need not be present. This may, for example, be the
139		case where the XSI-conformant system is hosted and the underlying system provides the service
140		in an alternative way.

Glossary

141

142 **absolute pathname**

143 See **pathname resolution** on page 22.

144 **access mode**

145 A particular form of access permitted to a file.

146 **additional file access control mechanism**

147 See **file access permissions** on page 14.

148 **address space**

149 The memory locations that can be referenced by a process or the threads of a process.

150 **affirmative response**

151 An input string that matches one of the responses acceptable to the LC_MESSAGES category
 152 keyword **yesexpr**, matching an extended regular expression in the current locale; see Section
 153 5.3.6 on page 80.

154 **alert**

155 To cause the user's terminal to give some audible or visual indication that an error or some other
 156 event has occurred. When the standard output is directed to a terminal device, the method for
 157 alerting the terminal user is unspecified. When the standard output is not directed to a terminal
 158 device, the alert is accomplished by writing the alert character to standard output (unless the
 159 utility description indicates that the use of standard output produces undefined results in this
 160 case).

161 **alert character**

162 A character that in the output stream should cause a terminal to alert its user via a visual or
 163 audible notification. The alert character is the character designated by '\a' in the C language. It
 164 is unspecified whether this character is the exact sequence transmitted to an output device by
 165 the system to accomplish the alert function.

166 **alias name**

167 A word consisting solely of underscores, digits and alphabets from the portable character set
 168 (see Section 4.1 on page 43) and any of the following characters:

169 ! % , @

170 Implementations may allow other characters within alias names as an extension.

171 **alternate file access control mechanism**

172 See **file access permissions** on page 14.

173 **alternate signal stack**

174 EX Memory associated with a thread, established upon request by the implementation for a thread,
 175 separate from the thread signal stack, in which signal handlers responding to signals sent to that
 176 thread may be executed.

177 **angle brackets**

178 The characters "<" (left-angle-bracket) and ">" (right-angle-bracket). When used in the phrase
 179 "enclosed in angle brackets", the symbol "<" immediately precedes the object to be enclosed, and
 180 ">" immediately follows it. When describing these characters in the portable character set, the
 181 names <less-than-sign> and <greater-than-sign> are used.

appropriate privileges

An implementation-dependent means of associating privileges with a process with regard to the function calls and function call options defined in the **XSH** specification, and the commands in the **XCU** specification, that need special privileges. There may be zero or more such means.

argument

In the shell, a parameter passed to a utility as the equivalent of a single string in the *argv* array created by one of the *exec* functions. See Section 10.1 on page 133 and the **XCU** specification, **Command Search and Execution** in **Section 2.9.1**. An argument is one of the options, option-arguments or operands following the command name.

In the C language, an expression in a function call expression or a sequence of preprocessing tokens in a function-like macro invocation.

arm (a timer)

To start a timer measuring the passage of time, enabling notifying a process when the specified time or time interval has passed.

assignment

See **variable assignment** on page 35.

asterisk

The character "*".

async-cancel safe function

A function that may be safely invoked by an application while the asynchronous form of cancellation is enabled. No function is async-cancel-safe unless explicitly described as such.

async-signal safe function

A function that may be invoked, without restriction, from signal-catching functions. No function is async-signal safe unless explicitly described as such.

asynchronously generated signal

A signal that is not attributable to a specific thread. Examples are: signals sent via *kill()*, signals sent from the keyboard, and signals delivered to process groups. Being asynchronous is a property of how the signal was generated and not a property of the signal number. All signals may be generated asynchronously.

asynchronous I/O operation

An I/O operation that does not of itself cause the thread requesting the I/O to be blocked from further use of the processor.

This implies that the process and the I/O operation may be running concurrently.

asynchronous I/O completion

For an asynchronous read or write operation, when a corresponding synchronous read or write would have completed and when any associated status fields have been updated.

background job

See **background process group**.

background process

A process that is a member of a background process group.

background process group

(Or **background job**.) Any process group, other than a foreground process group, that is a member of a session that has established a connection with a controlling terminal.

225	backquote
226	The character ` , also known as a <i>grave accent</i> .
227	backslash
228	The character "\", also known as a <i>reverse solidus</i> .
229	backspace character
230	A character that, in the output stream, should cause printing (or displaying) to occur one column
231	position previous to the position about to be printed. If the position about to be printed is at the
232	beginning of the current line, the behaviour is unspecified. The backspace is the character
233	designated by '\b' in the C language. It is unspecified whether this character is the exact
234	sequence transmitted to an output device by the system to accomplish the backspace function.
235	The backspace character defined here is not necessarily the ERASE special character defined in
236	Section 9.1.9 on page 123.
237	base character
238	One of the set of characters defined in the Latin alphabet. In Western European languages other
239	than English, these characters are commonly used with diacritical marks (accents, cedilla, and so
240	on) to extend the range of characters in an alphabet.
241	basename
242	The final, or only, filename in a pathname.
243	basic regular expression
244	A pattern constructed according to the rules defined in Section 7.3 on page 104.
245	blank character
246	One of the characters that belong to the blank character class as defined via the LC_CTYPE
247	category in the current locale. In the POSIX locale, a blank character is either a tab or a space
248	character.
249	blank line
250	A line consisting solely of zero or more blank characters terminated by a newline character. See
251	also empty line on page 12.
252	blocked process (or thread)
253	A process (or thread) that is waiting for some condition (other than the availability of a
254	processor) to be satisfied before it can continue execution.
255	block-mode terminal
256	A terminal device operating in a mode incapable of the character-at-a-time input and output
257	operations described by some of the standard utilities. See Section 8.2 on page 118.
258	block special file
259	A file that refers to a device. A block special file is normally distinguished from a character
260	special file by providing access to the device in a manner such that the hardware characteristics
261	of the device are not visible.
262	braces
263	The characters "{" (left brace) and "}" (right brace), also known as <i>curly braces</i> . When used in the
264	phrase “enclosed in (curly) braces” the symbol "{" immediately precedes the object to be
265	enclosed, and "}" immediately follows it. When describing these characters in the portable
266	character set, the names <left-brace> and <right-brace> are used.
267	brackets
268	The characters "[" (left-bracket) and "]" (right-bracket), also known as <i>square brackets</i> . When used
269	in the phrase “enclosed in (square) brackets” the symbol "[" immediately precedes the object to
270	be enclosed, and "]" immediately follows it. When describing these characters in the portable

character set, the names <left-square-bracket> and <right-square-bracket> are used.

break value

The address at which dynamic memory allocation starts.

built-in utility

(Or **built-in**.) A utility implemented within a shell. The utilities referred to as *special built-ins* have special qualities, described in the XCU specification, **Section 2.14, Special Built-in Utilities**. Unless qualified, the term *built-in* includes the special built-in utilities. The utilities referred to as *regular built-ins* are those named in the XCU specification, **Command Search and Execution in Section 2.9.1**. There is no requirement that these utilities be actually built into the shell on the implementation, but they do have special command-search qualities.

byte

An individually addressable unit of data storage that is equal to or larger than an octet, used to store a character or a portion of a character; see **character**. A byte is composed of a contiguous sequence of bits, the number of which is implementation-dependent. The least significant bit is called the *low-order* bit; the most significant is called the *high-order* bit. Note that this definition of *byte* deviates intentionally from the usage of *byte* in some international standards, where it is used as a synonym for *octet* (always eight bits). On a system based on the ISO/IEC 9945-2: 1993 standard, a byte may be larger than eight bits so that it can be an integral portion of larger data objects that are not evenly divisible by eight bits (such as a 36-bit word that contains four 9-bit bytes).

carriage-return character

A character that in the output stream indicates that printing should start at the beginning of the same physical line in which the carriage-return character occurred. The carriage-return is the character designated by '\r' in the C language. It is unspecified whether this character is the exact sequence transmitted to an output device by the system to accomplish the movement to the beginning of the line.

character

A sequence of one or more bytes representing a single graphic symbol or control code. This term corresponds to the ISO C standard term *multibyte character* (multi-byte character), where a single-byte character is a special case of a multi-byte character. Unlike the usage in the ISO C standard, *character* here has no necessary relationship with storage space, and *byte* is used when storage space is discussed.

See Section 4.1 on page 43 for a further explanation of the graphical representations of characters, or *glyphs*, as opposed to character encodings.

character array

An array of type **char**.

character class

A named set of characters sharing an attribute associated with the name of the class. The classes and the characters that they contain are dependent on the value of the LC_CTYPE category in the current locale; see Section 5.3.1 on page 52.

character set

A finite set of different characters used for the representation, organisation or control of data.

character special file

A file that refers to a device. One specific type of character special file is a terminal device file, whose access is defined in Chapter 9 on page 119.

316	character string
317	A contiguous sequence of characters terminated by and including the first null byte.
318	child process
319	See process on page 25.
320	circumflex
321	The character "^".
322	clock
323	An object that measures the passage of time.
324	The current value of the time measured by a clock can be queried and, possibly, set to a value
325	within the legal range of the clock.
326	clock tick
327	An interval of time; an implementation-dependent number of these occur each second.
328	coded character set
329	A set of unambiguous rules that establishes a character set and the one-to-one relationship
330	between each character of the set and its bit representation.
331	codeset
332	The result of applying rules that map a numeric code value to each element of a character set.
333	An element of a character set may be related to more than one numeric code value but the
334	reverse is not true. However, for state-dependent encodings the relationship between numeric
335	code values to elements of a character set may be further controlled by state information; see
336	Section 4.2 on page 44. The character set may contain fewer elements than the total number of
337	possible numeric code values; that is, some code values may be unassigned.
338	collating element
339	The smallest entity used to determine the logical ordering of character or wide-character strings.
340	See collation sequence . A collating element consists of either a single character, or two or more
341	characters collating as a single entity. The value of the LC_COLLATE category in the current
342	locale determines the current set of collating elements.
343	collation
344	The logical ordering of character or wide-character strings according to defined precedence
345	rules. These rules identify a collation sequence between the collating elements, and such
346	additional rules that can be used to order strings consisting of multiple collating elements.
347	collation sequence
348	The relative order of collating elements as determined by the setting of the LC_COLLATE
349	category in the current locale. The character order, as defined for the LC_COLLATE category in
350	the current locale, defines the relative order of all collating elements, such that each element
351	occupies a unique position in the order. This is the order used in ranges of characters and
352	collating elements in regular expressions and pattern matching. In addition, the definition of the
353	collating weights of characters and collating elements uses collating elements to represent their
354	respective positions within the collation sequence.
355	Multi-level sorting is accomplished by assigning elements one or more collation weights, up to
356	the limit {COLL_WEIGHTS_MAX}; see < limits.h >. On each level, elements may be given the
357	same weight (at the primary level, called an equivalence class; see equivalence class on page 13)
358	or be omitted from the sequence. Strings that collate equal using the first assigned weight
359	(primary ordering) are then compared using the next assigned weight (secondary ordering), and
360	so on.

column position

A unit of horizontal measure related to characters in a line.

It is assumed that each character in a character set has an intrinsic column width independent of any output device. Each printable character in the portable character set has a column width of one. The standard utilities, when used as described in this specification set, assume that all characters have integral column widths. The column width of a character is not necessarily related to the internal representation of the character (numbers of bits or bytes).

The column position of a character in a line is defined as one plus the sum of the column widths of the preceding characters in the line. Column positions are numbered starting from 1.

command

A directive to the shell to perform a particular task; see the XCU specification, **Section 2.9, Shell Commands**.

command language interpreter

An interface that interprets sequences of text input as commands. It may operate on an input stream or it may interactively prompt and read commands from a terminal. It is possible for applications to invoke utilities through a number of interfaces, which are collectively considered to act as command interpreters. The most obvious of these are the *sh* utility and the *system()* function, although *popen()* and the various forms of *exec* may also be considered to behave as interpreters.

composite graphic symbol

A graphic symbol consisting of a combination of two or more other graphic symbols in a single character position, such as a diacritical mark and a basic letter.

condition variable

A synchronization object which allows a thread to suspend execution, repeatedly, until some associated predicate becomes true.

control character

A character, other than a graphic character, that affects the recording, processing, transmission or interpretation of text.

control operator

In the shell, a token that performs a control function. It is one of the following symbols:

```
&    &&    (    )    ;    ;;    newline    |    ||
```

The end-of-input indicator used internally by the shell is also considered a control operator. See the XCU specification, **Section 2.3, Token Recognition**.

On some systems, the symbol `((` is a control operator; its use produces unspecified results. Applications that wish to have nested subshells, such as:

```
((echo Hello);(echo World))
```

must separate the `((` characters into two tokens by including white space between them. Some systems may treat these as invalid arithmetic expressions instead of subshells.

The `((` and `)` symbols are control operators in the KornShell, used for an alternative syntax of an arithmetic expression command. A portable application cannot use `((` as a single token (with the exception of the `$((` form for shell arithmetic).

controlling process

The session leader that established the connection to the controlling terminal. If the terminal ceases to be a controlling terminal for this session, the session leader ceases to be the controlling process.

406		controlling terminal	
407		A terminal that is associated with a session. Each session may have at most one controlling	
408		terminal associated with it, and a controlling terminal is associated with exactly one session.	
409		Certain input sequences from the controlling terminal (see Chapter 9 on page 119) cause signals	
410		to be sent to all processes in the process group associated with the controlling terminal.	
411		conversion descriptor	
412	EX	A per-process unique value used to identify an open codeset conversion.	
413		core file	
414	EX	A file of unspecified format that may be generated when a process terminates abnormally.	
415		current working directory	
416		See working directory on page 36.	
417		cursor position	
418		The line and column position on the screen denoted by the terminal's cursor.	
419		data segment	
420	EX	Memory associated with a process, that may be used to contain dynamically allocated data.	
421		device	
422		A computer peripheral or an object that appears to the application as such.	
423		device ID	
424		A non-negative integer used to identify a device.	
425		direct I/O	
426		An operation that attempts to circumvent a system performance optimization for the	
427		optimization of the individual I/O operation.	
428		directory	
429		A file that contains directory entries. No two directory entries in the same directory have the	
430		same name.	
431		directory entry	
432		(Or link .) An object that associates a filename with a file. Several directory entries can associate	
433		names with the same file.	
434		directory stream	
435		A sequence of all the directory entries in a particular directory. An open directory stream may	
436		be implemented using a file descriptor.	
437		disarm (a timer)	
438		To stop a timer from measuring the passage of time, disabling any future process notifications	
439		(until the timer is armed again).	
440		display	
441		To output to the user's terminal. If the output is not directed to a terminal, the results are	
442		undefined.	
443		The XCU specification assigns precise requirements for the terms <i>display</i> and <i>write</i> . Some	
444		historical systems have chosen to implement certain utilities without using the traditional UNIX	
445		system file descriptor model. For example, the <i>vi</i> editor might employ direct screen memory	
446		updates on a personal computer, rather than a <i>write()</i> system call. An instance of user	
447		prompting might appear in a dialogue box, rather than with standard error. When the XCU	
448		specification uses the term <i>display</i> , the method of outputting to the terminal is unspecified; many	
449		historical implementations use <i>termcap</i> or <i>terminfo</i> , but this is not a requirement. The term <i>write</i>	
450		is used when the XCU specification mandates that a file descriptor be used and that the output	

can be redirected. However, it is assumed that when the writing is directly to the terminal (it has not been redirected elsewhere), there is no practical way for a user or test suite to determine whether a file descriptor is being used or not. Therefore, the use of a file descriptor is mandated only for the redirection case and the implementation is free to use any method when the output is not redirected. The verb *write* is used almost exclusively, with the very few exceptions of those utilities where output redirection need not be supported: *tabs*, *talk*, *tput* and *vi*.

dollar sign

The character "\$".

dot

The filename consisting of a single dot character (.). See **pathname resolution** on page 22. In the context of shell special built-in utilities, see *dot* in the XCU specification, **Section 2.14, Special Built-in Utilities**.

dot-dot

The filename consisting solely of two dot characters (.). See **pathname resolution** on page 22.

double-quote

The character "\"", also known as *quotation-mark*.

downshifting

The conversion of an upper-case character to its lower-case representation.

(clock) drift rate

The rate at which the time measured by a clock deviates from the actual passage of real time.

A positive drift rate causes a clock to gain time with respect to real time; a negative drift rate causes a clock to lose time with respect to real time.

driver

A module that controls data transferred to and received from peripheral devices. Drivers are traditionally written to be a part of the system implementation, although they are frequently written separately from the writing of the implementation. A driver may contain processor-specific code, and therefore be non-portable.

effective group ID

An attribute of a process that is used in determining various permissions, including file access permissions, described in **file access permissions** on page 14. See **group ID** on page 17. This value is subject to change during the process lifetime, as described in the *exec* family of functions and *setgid()*.

effective user ID

An attribute of a process that is used in determining various permissions, including file access permissions. See **user ID** on page 35. This value is subject to change during the process lifetime, as described in *exec* and *setuid()*.

eight-bit transparency

The ability of a software component to process 8-bit characters without modifying or utilising any part of the character in a way that is inconsistent with the rules of the current coded character set.

empty directory

A directory that contains, at most, directory entries for dot and dot-dot.

empty line

A line consisting of only a newline character. See also **blank line** on page 7.

495	empty string	
496	(Or null string .) A string whose first byte is a null byte.	
497	empty wide-character string	
498	A wide-character string whose first element is a null wide-character code.	
499	epoch	
500	The time zero hours, zero minutes, zero seconds, on January 1, 1970 Coordinated Universal	
501	Time. See seconds since the epoch on page 28.	
502	equivalence class	
503	A set of collating elements with the same primary collation weight.	
504	Elements in an equivalence class are typically elements that naturally group together, such as all	
505	accented letters based on the same base letter.	
506	The collation order of elements within an equivalence class is determined by the weights	
507	assigned on any subsequent levels after the primary weight.	
508	era	
509	An alternative method for counting and displaying years. See Section 5.3.5 on page 73.	
510	executable file	
511	A regular file acceptable as a new process image file by the equivalent of the <i>exec</i> family of	
512	functions, and thus usable as one form of a utility. The standard utilities described as compilers	
513	can produce executable files, but other unspecified methods of producing executable files may	
514	also be provided. The internal format of an executable file is unspecified, but a conforming	
515	application cannot assume an executable file is a text file.	
516	execute	
517	To perform the actions described in the XCU specification, Command Search and Execution in	
518	Section 2.9.1 . See also invoke on page 17.	
519	expand	
520	In the shell, when not qualified, the act of applying all the expansions described in the XCU	
521	specification, Section 2.6, Word Expansions .	
522	extended regular expression	
523	A pattern constructed according to the rules defined in Section 7.4 on page 109.	
524	extended signed integral type	
525	EX A signed integral type or an implementation-dependent type with similar properties.	
526	extended security controls	
527	The access control (see file access permissions on page 14) and privilege (see appropriate	
528	privileges on page 6) mechanisms have been defined to allow implementation-dependent	
529	extended security controls. These permit an implementation to provide security mechanisms to	
530	support different security policies from those described in this specification set. These	
531	mechanisms do not alter or override the defined semantics of any of the functions or utilities in	
532	this specification set.	
533	extended unsigned integral type	
534	EX An unsigned integral type or an implementation-dependent type with similar properties.	
535	feature test macro	
536	A macro used to determine whether a particular set of features will be included from a header.	
537	See the XSH specification, Section 2.2, The Compilation Environment .	

field

In the shell, a unit of text that is the result of parameter expansion (see the XCU specification, **Section 2.6.2, Parameter Expansion**), arithmetic expansion (see the XCU specification, **Section 2.6.4, Arithmetic Expansion**), command substitution (see the XCU specification, **Section 2.6.3, Command Substitution**), or field splitting (see the XCU specification, **Section 2.6.5, Field Splitting**). During command processing (see the XCU specification, **Section 2.9.1, Simple Commands**), the resulting fields are used as the command name and its arguments.

FIFO special file

(Or **FIFO**.) A type of file with the property that data written to such a file is read on a first-in-first-out basis. Other characteristics of FIFOs are described in *open()*, *read()*, *write()* and *lseek()*.

file

An object that can be written to, or read from, or both. A file has certain attributes, including access permissions and type. File types include regular file, character special file, block special file, FIFO special file and directory. Other types of files may be supported by the implementation.

file access permissions

The standard file access control mechanism uses the file permission bits, as described below. These bits are set at the time of file creation by functions such as *open()*, *creat()*, *mkdir()* and *mkfifo()* and are changed by *chmod()*. These bits are read by *stat()* or *fstat()*.

Implementations may provide *additional* or *alternate* file access control mechanisms, or both. An additional access control mechanism will only further restrict the access permissions defined by the file permission bits. An alternate file access control mechanism will:

- specify file permission bits for the file owner class, file group class, and file other class of that file, corresponding to the access permissions, to be returned by *stat()* or *fstat()*
- be enabled only by explicit user action, on a per-file basis by the file owner or a user with the appropriate privilege
- be disabled for a file after the file permission bits are changed for that file with *chmod()*. The disabling of the alternate mechanism need not disable any additional mechanisms supported by an implementation.

Whenever a process requests file access permission for read, write or execute/search, if no additional mechanism denies access, access is determined as follows:

- If a process has the appropriate privilege:
 - If read, write or directory search permission is requested, access is granted.
 - If execute permission is requested, access is granted if execute permission is granted to at least one user by the file permission bits or by an alternate access control mechanism; otherwise, access is denied.
- Otherwise:
 - The file permission bits of a file contain read, write and execute/search permissions for the file owner class, file group class and file other class.
 - Access is granted if an alternate access control mechanism is not enabled and the requested access permission bit is set for the class (file owner class, file group class, or file other class) to which the process belongs, or if an alternate access control mechanism is enabled and it allows the requested access; otherwise, access is denied.

581	file description
582	See open file description on page 21.
583	file descriptor
584	A per-process unique, non-negative integer used to identify an open file for the purpose of file
585	access. The value of a file descriptor is from zero to {OPEN_MAX}. A process can have no more
586	than {OPEN_MAX} file descriptors open simultaneously. File descriptors may also be used to
587	EX implement message catalogue descriptors and directory streams . See open file description on
588	page 21 and {OPEN_MAX} in <limits.h> .
589	file group class
590	The property of a file indicating access permissions for a process related to the group
591	identification of a process. A process is in the file group class of a file if the process is not in the
592	file owner class and if the effective group ID or one of the supplementary group IDs of the
593	process matches the group ID associated with the file. Other members of the class may be
594	implementation-dependent.
595	file hierarchy
596	Files in the system are organised in a hierarchical structure in which all of the non-terminal
597	nodes are directories and all of the terminal nodes are any other type of file. Because multiple
598	directory entries may refer to the same file, the hierarchy is properly described as a <i>directed graph</i> .
599	file mode
600	An object containing the <i>file mode bits</i> and file type of a file, as described in <sys/stat.h> .
601	file mode bits
602	A file's file permission bits, set-user-ID-on-execution bit (S_ISUID) and set-group-ID-on-
603	execution bit (S_ISGID); see <sys/stat.h> .
604	filename
605	A name consisting of 1 to {NAME_MAX} bytes used to name a file. The characters composing
606	the name may be selected from the set of all character values excluding the slash character and
607	the null byte. The filenames dot and dot-dot have special meaning; see pathname resolution on
608	page 22. A filename is sometimes referred to as a <i>pathname component</i> .
609	Filenames should be constructed from the portable filename character set because the use of
610	other characters can be confusing or ambiguous in certain contexts. (For instance, the use of a
611	colon (:) in a pathname could cause ambiguity if that pathname were included in a <i>PATH</i>
612	definition.)
613	file offset
614	The byte position in the file where the next I/O operation begins. Each open file description
615	associated with a regular file, block special file or directory has a file offset. A character special
616	file that does not refer to a terminal device may have a file offset. There is no file offset specified
617	for a pipe or FIFO.
618	file other class
619	The property of a file indicating access permissions for a process related to the user and group
620	identification of a process. A process is in the file other class of a file if the process is not in the
621	file owner class or file group class.
622	file owner class
623	The property of a file indicating access permissions for a process related to the user
624	identification of a process. A process is in the file owner class of a file if the effective user ID of
625	the process matches the user ID of the file.

file permission bits

Information about a file that is used, along with other information, to determine if a process has read, write or execute/search permission to a file. The bits are divided into three parts: owner, group and other. Each part is used with the corresponding file class of processes. These bits are contained in the file mode, as described in `<sys/stat.h>`. The detailed usage of the file permission bits in access decisions is described in **file access permissions** on page 14.

file serial number

A per-file-system unique identifier for a file.

file system

A collection of files and certain of their attributes. It provides a name space for file serial numbers referring to those files.

file times update

Each file has three associated time values that are updated when file data has been accessed, file data has been modified, or file status has been changed, respectively. These values are returned in the file characteristics structure, as described in `<sys/stat.h>`.

For each function or utility in this specification set that reads or writes file data or changes the file status, the appropriate time-related fields are noted as “marked for update”. At an update point in time, any marked fields are set to the current time and the update marks cleared. Two such update points are when the file is no longer open by any process and when `stat()` or `fstat()` is performed on the file. Additional update points are unspecified. Marks for update, and updates themselves, are not done for files on read-only file systems.

file type

See **file** on page 14.

filter

A command whose operation consists of reading data from standard input or a list of input files and writing data to standard output. Typically, its function is to perform some transformation on the data stream.

first open (of a file)

When a process opens a file that is not currently an open file within any process.

foreground job

See **foreground process group**.

foreground process

A process that is a member of a foreground process group.

foreground process group

(Or **foreground job**.) A process group whose member processes have certain privileges, denied to processes in background process groups, when accessing their controlling terminal. Each session that has established a connection with a controlling terminal has exactly one process group of the session as the foreground process group of that controlling terminal. See Chapter 9.

foreground process group ID

The process group ID of the foreground process group.

form-feed character

A character that in the output stream indicates that printing should start on the next page of an output device. The form-feed is the character designated by `'\f'` in the C language. If the form-feed is not the first character of an output line, the result is unspecified. It is unspecified whether this character is the exact sequence transmitted to an output device by the system to accomplish the movement to the next page.

672		graphic character
673		A character, other than a control character, that has a visual representation when handwritten,
674		printed or displayed.
675		group database
676		A system database of implementation-dependent format that contains at least the following
677		information for each group ID:
678		<ul style="list-style-type: none"> • Group Name
679		<ul style="list-style-type: none"> • Numerical Group ID
680		<ul style="list-style-type: none"> • List of users allowed in the group.
681		The list of users allowed in the group is used by the <i>newgrp</i> utility.
682		group ID
683		A non-negative integer that is used to identify a group of system users. Each system user is a
684		member of at least one group. When the identity of a group is associated with a process, a group
685	FIPS	ID value is referred to as a real group ID, an effective group ID, one of the supplementary group
686		IDs or a saved set-group-ID.
687		group name
688		A string that is used to identify a group, as described in group database . To be portable across
689		XSI-conformant systems, the value must be composed of characters from the portable filename
690		character set. The hyphen should not be used as the first character of a portable group name.
691		hard limit
692	EX	A system resource limitation that may be reset to a lesser or greater limit by a privileged process.
693		A non-privileged process is restricted to only lowering its hard limit.
694		hard link
695		The relationship between two directory entries that represent the same file; see directory entry
696		on page 11. This term is contrasted against symbolic link ; see symbolic link on page 31.
697		home directory
698		The current directory associated with a user at the time of login.
699		incomplete line
700		A sequence of one or more non-newline characters at the end of the file.
701		Inf
702		A value representing infinity that can be stored in a floating type. Not all systems support the
703		Inf value.
704		interactive shell
705		A processing mode of the shell that is suitable for direct user interaction.
706		internationalisation
707		The provision within a computer program of the capability of making itself adaptable to the
708		requirements of different native languages, local customs and coded character sets.
709		invoke
710		To perform the actions described in the XCU specification, Command Search and Execution in
711		Section 2.9.1 , except that searching for shell functions and special built-in utilities is suppressed.
712		See also execute on page 13.
713		ISO/IEC 646:1983
714		ISO 7-bit coded character set for information interchange. The reference version of the standard
715		contains 95 graphic characters, which are identical to the graphic characters defined in the ASCII

coded character set.

ISO 6937: 1983

ISO 7-bit or 8-bit coded character set for text communication using public communication networks, private communication networks, or interchange media such as magnetic tapes and discs.

ISO 8859-1: 1987

ISO 8-bit single-byte coded character set Part 1, Latin Alphabet No 1. This standard character set comprises 191 graphic characters covering the requirements of most of Western Europe.

job

A set of processes, comprising a shell pipeline, and any processes descended from it, that are all in the same process group. See the definition of **pipeline** in the XCU specification, **Section 2.9.2, Pipelines**.

job control

A facility that allows users selectively to stop (suspend) the execution of processes and continue (resume) their execution at a later point. The user typically employs this facility via the interactive interface jointly supplied by the terminal I/O driver and a command interpreter.

job control job ID

A handle that is used to refer to a job. The job control job ID can be any of the forms shown in the following table:

Job Control Job ID	Meaning
%%	Current job
%+	Current job
%–	Previous job
% <i>n</i>	Job number <i>n</i>
% <i>string</i>	Job whose command begins with <i>string</i>
%? <i>string</i>	Job whose command contains <i>string</i>

Table 2-1 Job Control Job ID Formats

last close (of a file)

When a process closes a file, resulting in the file not being an open file within any process.

line

A sequence of zero or more non-newline characters plus a terminating newline character.

link

See **directory entry** on page 11.

link count

The number of directory entries that refer to a particular file.

local customs

The conventions of a geographical area or territory for such things as date, time and currency formats.

locale

The definition of the subset of a user's environment that depends on language and cultural conventions; see Chapter 5 on page 49.

localisation

The process of establishing information within a computer system specific to the operation of particular native languages, local customs and coded character sets.

761		login
762		The unspecified activity by which a user gains access to the system. Each login is associated
763		with exactly one login name.
764		login name
765		A user name that is associated with a login.
766		map
767		To create an association between a page-aligned range of the address-space of a process and a
768		range of physical memory or some memory object, such that a reference to an address in that
769		range of the address-space results in a reference to the associated physical memory or memory
770		object. The mapped memory or memory object is not necessarily memory-resident.
771		marked message
772	EX	A STREAMs message on which a certain flag is set. Marking a message gives the application
773		protocol-specific information. An application can use <i>ioctl()</i> to determine whether a given
774		message is marked.
775		memory object
776	RT	Either a file or shared memory object.
777		When used in conjunction with <i>mmap()</i> , a memory object will appear in the address-space of the
778		calling process.
779		message
780		Information that can be transferred between processes or threads by being added to and
781		removed from a message queue. A message consists of a fixed-size message buffer.
782		message catalogue
783	EX	A file or storage area containing program messages, command prompts and responses to
784		prompts for a particular native language, territory and codeset.
785		message catalogue descriptor
786	EX	A per-process unique value used to identify an open message catalogue. A message catalogue
787		descriptor may be implemented using a file descriptor.
788		message queue
789		An object to which messages can be added and removed. Messages may be removed in the
790		order in which they were added or in priority order.
791		mode
792		A collection of attributes that specifies a file's type and its access permissions. See file access
793		permissions on page 14.
794		mount point
795		Either the system root directory or a directory for which the <i>st_dev</i> field of structure stat (see
796		<sys/stat.h>) differs from that of its parent directory.
797		multi-character collating element
798		A sequence of two or more characters that collate as an entity. For example, in some coded
799		character sets, an accented character is represented by a non-spacing accent, followed by the
800		letter. Other examples are the Spanish elements <i>ch</i> and <i>ll</i> .
801		mutex
802		A synchronization object used to allow multiple threads to serialize their access to shared data.
803		The name derives from the capability it provides; namely, mutual exclusion. The thread that has
804		locked a mutex becomes its owner and remains the owner until that same thread unlocks the
805		mutex.

name

In the shell, a word consisting solely of underscores, digits and alphabetic characters from the portable character set (see Section 4.1 on page 43). The first character of a name must not be a digit.

There are no explicit limits in this specification set on the sizes of names, words (see **word** on page 36), lines or other objects. However, other implicit limits do apply: shell script lines produced by many of the standard utilities cannot exceed {LINE_MAX} and the sum of exported variables comes under the {ARG_MAX} limit. Historical shells dynamically allocate memory for names and words and parse incoming lines a byte at a time. Lines cannot have an arbitrary {LINE_MAX} limit because of historical practice such as makefiles, where *make* removes the newline characters associated with the commands for a target and presents the shell with one very long line. The text on **INPUT FILES** in the XCU specification, **Section 1.9, Utility Description Defaults** does allow a shell to run out of memory, but it cannot have arbitrary programming limits.

named STREAM

EX A STREAMS-based file descriptor that is attached to a name in the file-system namespace. All subsequent operations on the named STREAM act on the STREAM that was associated with the file descriptor until the name is disassociated from the STREAM.

NaN (not a number)

A value that can be stored in a floating type but that is not a valid floating point number. Not all systems support the NaN value.

native language

A computer user's spoken or written language, such as American English, British English, Danish, Dutch, French, German, Italian, Japanese, Norwegian or Swedish.

negative response

An input string that matches one of the responses acceptable to the LC_MESSAGES category keyword **noexpr**, matching an extended regular expression in the current locale. See Section 5.3.6 on page 80.

newline character

A character that in the output stream indicates that printing should start at the beginning of the next line. The newline is the character designated by '\n' in the C language. It is unspecified whether this character is the exact sequence transmitted to an output device by the system to accomplish the movement to the next line.

non-spacing characters

A character, such as a character representing a diacritical mark in the ISO 6937:1983 standard coded character set, which is used in combination with other characters to form composite graphic symbols.

NUL

A character with all bits set to zero.

null byte

A byte with all bits set to zero.

null pointer

The value that is obtained by converting the number 0 into a pointer; for example, **(void *) 0**. The C language guarantees that this value will not match that of any legitimate pointer, so it is used by many functions that return pointers to indicate an error.

null string

See **empty string** on page 13.

852		null wide-character code
853		A wide-character code with all bits set to zero.
854		number sign
855		The character #, also known as <i>hash sign</i> .
856		object file
857		A regular file containing the output of a compiler, formatted as input to a linkage editor for
858		linking with other object files into an executable form. The methods of linking are unspecified
859		and may involve the dynamic linking of objects at run time. The internal format of an object file
860		is unspecified, but a conforming application cannot assume an object file is a text file.
861		offset maximum
862	EX	An attribute of an open file description representing the largest value that can be used as a file
863		offset.
864		open file
865		A file that is currently associated with a file descriptor.
866		open file description
867		A record of how a process or group of processes are accessing a file. Each file descriptor refers to
868		exactly one open file description, but an open file description can be referred to by more than
869		one file descriptor. A file offset, file status and file access modes are attributes of an open file
870		description.
871		operand
872		An argument to a command that is generally used as an object supplying information to a utility
873		necessary to complete its processing. Operands generally follow the options in a command line.
874		See Section 10.1 on page 133.
875		operator
876		In the shell, either a control operator or a redirection operator.
877		option
878		An argument to a command that is generally used to specify changes in the utility's default
879		behaviour; see Section 10.1 on page 133.
880		option-argument
881		A parameter that follows certain options. In some cases an option-argument is included within
882		the same argument string as the option; in most cases it is the next argument. See Section 10.1
883		on page 133.
884		orphaned process group
885		A process group in which the parent of every member is either itself a member of the group or is
886		not a member of the group's session.
887		page
888		The granularity of process memory mapping or locking.
889		Physical memory and memory objects can be mapped into the address-space of a process on
890		page boundaries and in integral multiples of pages. Process address-space can be locked into
891		memory (made memory-resident) on page boundaries and in integral multiples of pages.
892		page size
893	EX	The size, in bytes, of the system unit of memory allocation, protection and mapping. On systems
894		that have segment- rather than page-based memory architectures, the term "page" means a
895		segment.

parameter

In the shell, an entity that stores values. There are three types of parameters: variables (named parameters), positional parameters and special parameters. Parameter expansion is accomplished by introducing a parameter with the "\$" character. See the XCU specification, **Section 2.5, Parameters and Variables**.

In the C language, an object declared as part of a function declaration or definition that acquires a value on entry to the function, or an identifier following the macro name in a function-like macro definition.

parent directory

When discussing a given directory, the directory that both contains a directory entry for the given directory and is represented by the pathname dot-dot in the given directory.

When discussing other types of files, a directory containing a directory entry for the file under discussion.

This concept does not apply to dot and dot-dot.

parent process

See **process** on page 25.

parent process ID

An attribute of a new process identifying the parent of the process. The parent process ID of a process is the process ID of its creator, for the lifetime of the creator. After the creator's lifetime has ended, the parent process ID is the process ID of an implementation-dependent system process.

pathname

A character string that is used to identify a file. A pathname consists of, at most, {PATH_MAX} bytes, including the terminating null byte. It has an optional beginning slash, followed by zero or more filenames separated by slashes. If the pathname refers to a directory, it may also have one or more trailing slashes. Multiple successive slashes are considered to be the same as one slash. A pathname that begins with two successive slashes may be interpreted in an implementation-dependent manner, although more than two leading slashes are treated as a single slash. The interpretation of the pathname is described in **pathname resolution**.

pathname component

See **filename** on page 15.

pathname resolution

Pathname resolution is performed for a process to resolve a pathname to a particular file in a file hierarchy. There may be multiple pathnames that resolve to the same file.

Each filename in the pathname is located in the directory specified by its predecessor (for example, in the pathname fragment **a/b**, file **b** is located in directory **a**). Pathname resolution fails if this cannot be accomplished. If the pathname begins with a slash, the predecessor of the first filename in the pathname is taken to be the root directory of the process (such pathnames are referred to as absolute pathnames). If the pathname does not begin with a slash, the predecessor of the first filename of the pathname is taken to be the current working directory of the process (such pathnames are referred to as relative pathnames).

The interpretation of a pathname component is dependent on the values of {NAME_MAX} and {_POSIX_NO_TRUNC} associated with the path prefix of that component. If any pathname component is longer than {NAME_MAX}, because {_POSIX_NO_TRUNC} is in effect on all XSI-conformant systems for the path prefix of that component (see *pathconf()*), the implementation will consider this an error condition.

942	EX	If a symbolic link (see symbolic link on page 31) is encountered during pathname resolution,	
943		then pathname resolution is complete if all of the following are true:	
944		<ul style="list-style-type: none"> • This is the last component of the pathname. 	
945		<ul style="list-style-type: none"> • The pathname has no trailing slash. 	
946		<ul style="list-style-type: none"> • The function is required to act on the symbolic link itself, or certain arguments direct that the 	
947		function act on the symbolic link itself.	
948		In all other cases, the system prefixes the remaining pathname, if any, with the contents of the	
949		symbolic link. The function may fail, setting <i>errno</i> to [ENAMETOOLONG], if the combined	
950		length exceeds {PATH_MAX}. Otherwise, the resolved pathname is the resolution of the	
951		pathname just created. The result is either an absolute pathname that is resolved from the root	
952		directory of the process or a relative pathname that is resolved from the directory containing the	
953		symbolic link.	
954		The special filename dot refers to the directory specified by its predecessor. The special filename	
955		dot-dot refers to the parent directory of its predecessor directory. As a special case, in the root	
956		directory, dot-dot may refer to the root directory itself.	
957		A pathname consisting of a single slash resolves to the root directory of the process. A null	
958		pathname is invalid.	
959		path prefix	
960		A pathname, with an optional ending slash, that refers to a directory.	
961		pattern	
962		A sequence of characters used either with regular expression notation (see Chapter 7 on page	
963		101) or for pathname expansion (see the XCU specification, Section 2.6.6, Pathname Expansion),	
964		as a means of selecting various character strings or pathnames, respectively.	
965		The syntaxes of the two patterns are similar, but not identical; this specification set always	
966		indicates the type of pattern being referred to in the immediate context of the use of the term.	
967		period	
968		The character (.). The term <i>period</i> is contrasted against dot , which is used to describe a specific	
969		directory entry.	
970		permissions	
971		See file access permissions on page 14.	
972		persistence	
973		A mode for semaphores, shared memory and message queues requiring that the object and its	
974		state (including data, if any) are preserved after the object is no longer referenced by any process.	
975		Persistence of an object does not imply that the state of the object is maintained across a system	
976		crash or a system reboot.	
977		pipe	
978		An object accessed by one of the pair of file descriptors created by the <i>pipe()</i> function. Once	
979		created, the file descriptors can be used to manipulate it, and it behaves identically to a FIFO	
980		special file when accessed in this way. It has no name in the file hierarchy.	
981		positional parameter	
982		In the shell, a parameter denoted by a single digit or one or more digits in curly braces. See the	
983		XCU specification, Section 2.5.1, Positional Parameters .	

portable character set

The collection of characters that are required to be present in all locales supported by XSI-conformant systems:

```
A B C D E F G H I J K L M N O P Q R S T U V W X Y Z
a b c d e f g h i j k l m n o p q r s t u v w x y z
0 1 2 3 4 5 6 7 8 9 ! # % ^ & * ( ) _ + - = { } [ ]
: " ~ ; ' ` < > ? , . | \ / @ $
```

Also included are the alert, backspace, tab, newline, vertical-tab, form-feed, carriage-return and space characters and the null character, NUL.

This term is contrasted against the smaller **portable filename character set**. See Table 4-1 on page 43.

portable filename character set

The set of characters from which portable filenames are constructed. For a filename to be portable across implementations conforming to this specification set and the ISO POSIX-1 standard, it must consist only of the following characters:

```
A B C D E F G H I J K L M N O P Q R S T U V W X Y Z
a b c d e f g h i j k l m n o p q r s t u v w x y z
0 1 2 3 4 5 6 7 8 9 . _ -
```

The last three characters are the period, underscore and hyphen characters, respectively. The hyphen must not be used as the first character of a portable filename. Upper- and lower-case letters retain their unique identities between conforming implementations. In the case of a portable pathname, the slash character may also be used.

preallocation

The reservation of resources in a system for a particular use.

Preallocation does not imply that the resources are immediately allocated to that use, but merely indicates that they are guaranteed to be available in bounded time when needed.

preempted process (or thread)

A running thread whose execution is suspended due to another thread becoming runnable at a higher priority.

printable character

One of the characters included in the **print** character classification of the LC_CTYPE category in the current locale; see Section 5.3.1 on page 52.

printable file

A text file consisting only of the characters included in the **print** and **space** character classifications of the LC_CTYPE category and the backspace character, all in the current locale; see Section 5.3.1 on page 52.

priority

A non-negative integer associated with processes or threads whose value is constrained to a range defined by the applicable scheduling policy. Numerically higher values represent higher priorities.

priority band

The queueing order applied to normal priority STREAMS messages. High priority STREAMS messages are not grouped by priority bands. The only differentiation made by the STREAMS mechanism is between zero and non-zero bands, but specific protocol modules may differentiate between priority bands.

1029	priority-based scheduling	
1030	Scheduling in which the selection of a running thread is determined by the priorities of the	
1031	runnable threads.	
1032	privilege	
1033	See appropriate privileges on page 6.	
1034	process	
1035	An address space with one or more threads executing within that address space, and the	
1036	required system resources for those threads.	
1037	Many of the system resources defined by this specification are shared among all of the threads	
1038	within a process. These include: the process ID, the parent process ID, process group ID, session	
1039	membership, real, effective and saved-set user ID, real, effective and saved-set group ID,	
1040	supplementary group IDs, current working directory, root directory, file mode creation mask	
1041	and file descriptors.	
1042	A process is created by another process issuing the <i>fork()</i> function. The process that issues <i>fork()</i>	
1043	is known as the parent process, and the new process created by the <i>fork()</i> is known as the child	
1044	process.	
1045	process group	
1046	A collection of processes that permits the signalling of related processes. Each process in the	
1047	system is a member of a process group that is identified by a process group ID. A newly created	
1048	process joins the process group of its creator.	
1049	process group ID	
1050	The unique identifier representing a process group during its lifetime. A process group ID is a	
1051	positive integer. A process group ID will not be reused by the system until the process group	
1052	lifetime ends.	
1053	process group leader	
1054	A process whose process ID is the same as its process group ID.	
1055	process group lifetime	
1056	A period of time that begins when a process group is created and ends when the last remaining	
1057	process in the group leaves the group, due either to the end of the last process' lifetime or to the	
1058	last remaining process calling the <i>setsid()</i> or <i>setpgid()</i> functions.	
1059	process ID	
1060	The unique identifier representing a process. A process ID is a positive integer. A process ID	
1061	will not be reused by the system until the process lifetime ends. In addition, if there exists a	
1062	process group whose process group ID is equal to that process ID, the process ID will not be	
1063	reused by the system until the process group lifetime ends. A process that is not a system	
1064	process will not have a process ID of 1.	
1065	process lifetime	
1066	The period of time that begins when a process is created and ends when its process ID is	
1067	returned to the system. After a process is created with a <i>fork()</i> function, it is considered active.	
1068	At least one thread of control and address space exist until it terminates. It then enters an	
1069	inactive state where certain resources may be returned to the system, although some resources,	
1070	EX such as the process ID, are still in use. When another process executes a <i>wait()</i> , <i>wait3()</i> , <i>waitid()</i>	
1071	or <i>waitpid()</i> function for an inactive process, the remaining resources are returned to the system.	
1072	The last resource to be returned to the system is the process ID. At this time, the lifetime of the	
1073	process ends.	

1074		process list	
1075	EX	See thread list on page 34.	
1076		process virtual time	
1077	EX	The measurement of time in units elapsed by the system clock while a process is executing.	
1078		program	
1079		A prepared sequence of instructions to the system to accomplish a defined task. The term	
1080		<i>program</i> in this specification set encompasses applications written in the XSI Shell Command	
1081		Language, complex utility input languages (for example, <i>awk</i> , <i>lex</i> , <i>sed</i> , and so on), and high-level	
1082		languages.	
1083		pseudo-terminal	
1084	EX	A pseudo-terminal provides the process with an interface that is identical to the terminal	
1085		subsystem. A pseudo-terminal is composed of 2 devices, the master device and a slave device.	
1086		The slave device provides processes with an interface that is identical to the terminal interface,	
1087		although there need not be hardware behind that interface. Anything written on the master	
1088		device is presented to the slave as an input and anything written on the slave device is presented	
1089		as an input on the master side.	
1090		This specification requires a STREAMS-based implementation of pseudo-terminals to be	
1091		available, but does not preclude others also being available.	
1092		radix character	
1093		The character that separates the integer part of a number from the fractional part.	
1094		read-only file system	
1095		A file system that has implementation-dependent characteristics restricting modifications.	
1096		read-write lock	
1097	EX	Multiple readers, single writer (read-write) locks allow many threads to have simultaneous	
1098		read-only access to data while allowing only one thread to have write access at any given time.	
1099		They are typically used to protect data that is read-only more frequently than it is changed.	
1100		Read-write locks can be used to synchronise threads in the current process and other processes if	
1101		they are allocated in memory that is writable and shared among the cooperating processes and	
1102		have been initialised for this behaviour.	
1103		real group ID	
1104		The attribute of a process that, at the time of process creation, identifies the group of the user	
1105		who created the process. See group ID on page 17. This value is subject to change during the	
1106		process lifetime, as described in <i>setgid()</i> .	
1107		real time	
1108	EX	Time measured as total units elapsed by the system clock without regard to which thread is	
1109		executing.	
1110		real user ID	
1111		The attribute of a process that, at the time of process creation, identifies the user who created the	
1112		process. See user ID on page 35. This value is subject to change during the process lifetime, as	
1113		described in <i>setuid()</i> .	
1114		redirection	
1115		In the shell, a method of associating files with the input or output of commands. See the XCU	
1116		specification, Section 2.7, Redirection .	

1117	redirection operator
1118	In the shell, a token that performs a redirection function. It is one of the following symbols:
1119	< > > << >> <& >& <<- <>
1120	reentrant function
1121	A function whose effect, when called by two or more threads, is guaranteed to be as if the
1122	threads each executed the function one after another in an undefined order, even if the actual
1123	execution is interleaved.
1124	referenced shared memory object
1125	A shared memory object that is open or has one or more mappings defined on it.
1126	refresh
1127	To ensure that the information on the user's terminal screen is up-to-date.
1128	regular expression
1129	A pattern constructed according to the rules defined in Chapter 7 on page 101.
1130	region
1131	In the context of the address space of a process, a sequence of addresses.
1132	In the context of a file, a sequence of offsets.
1133	regular file
1134	A file that is a randomly accessible sequence of bytes, with no further structure imposed by the
1135	system.
1136	relative pathname
1137	See pathname resolution on page 22.
1138	(time) resolution
1139	The minimum time interval that a clock can measure or whose passage a timer can detect.
1140	root directory
1141	A directory, associated with a process, that is used in pathname resolution for pathnames that
1142	begin with a slash.
1143	runnable process (or thread)
1144	A thread that is capable of being a running thread, but for which no processor is available.
1145	running process (or thread)
1146	A thread currently executing on a processor. On multi-processor systems there may be more
1147	than one such thread in a system at a time.
1148	saved resource limits
1149	EX An attribute of a process that provides some flexibility in the handling of unrepresentable
1150	resource limits, as described in the <i>exec</i> family of functions and <i>setrlimit()</i> .
1151	saved set-group-ID
1152	An attribute of a process that allows some flexibility in the assignment of the effective group ID
1153	attribute, as described in the <i>exec</i> family of functions and <i>setgid()</i> .
1154	saved set-user-ID
1155	An attribute of a process that allows some flexibility in the assignment of the effective user ID
1156	attribute, as described in <i>exec</i> and <i>setuid()</i> .
1157	scheduling
1158	The application of a policy to select a runnable process or thread to become a running process or
1159	thread, or to alter one or more of the thread lists.

scheduling allocation domain

The set of processors on which an individual thread can be scheduled at any given time.

scheduling contention scope

A property of a thread that defines the set of threads against which that thread competes for resources.

For example, in a scheduling decision, threads sharing scheduling contention scope compete for processor resources. In this specification, a thread has scheduling contention scope of either `PTHREAD_SCOPE_SYSTEM` or `PTHREAD_SCOPE_PROCESS`.

scheduling policy

A set of rules that is used to determine the order of execution of threads to achieve some goal.

In the context of XSI, a scheduling policy affects thread ordering:

- when a thread is a running thread and it becomes a blocked thread
- when a thread is a running thread and it becomes a preempted thread
- when a thread is a blocked thread and it becomes a runnable thread
- when a running thread calls a function that can change the priority or scheduling policy of a thread
- in other scheduling policy-defined circumstances.

Conforming implementations are required to define the manner in which each of the scheduling policies may modify the priorities or otherwise affect the ordering of threads at each of the occurrences listed above. Additionally, conforming implementations will define at what other circumstances and in what manner each scheduling policy may modify the priorities or affect the ordering of threads.

screen

A rectangular region of columns and lines on a terminal display. A screen may be a portion of a physical display device or may occupy the entire physical area of the display device.

scroll

To move the representation of data vertically or horizontally relative to the terminal screen. There are two types of scrolling:

1. The cursor moves with the data.
2. The cursor remains stationary while the data moves.

seconds since the epoch

A value to be interpreted as the number of seconds between a specified time and the epoch. A Coordinated Universal Time name (specified in terms of seconds (*tm_sec*), minutes (*tm_min*), hours (*tm_hour*), days since January 1 of the year (*tm_yday*), and calendar year minus 1900 (*tm_year*)) is related to a time represented as seconds since the Epoch, according to the expression below.

If the year < 1970 or the value is negative, the relationship is undefined. If the year ≥ 1970 and the value is non-negative, the value is related to a Coordinated Universal Time name according to the expression:

$$tm_sec + tm_min*60 + tm_hour*3\,600 + tm_yday*86\,400 + (tm_year-70)*31\,536\,000 + ((tm_year-69)/4)*86\,400$$

1201	semaphore	
1202	A shareable resource that has a non-negative integral value. When the value is zero, there is a	
1203	(possibly empty) set of threads awaiting the availability of the semaphore.	
1204	semaphore lock operation	
1205	An operation that is applied to a semaphore. If, prior to the operation, the value of the	
1206	semaphore is zero, the semaphore lock operation causes the calling thread to be blocked and	
1207	added to the set of threads awaiting the semaphore. Otherwise, the value is decremented.	
1208	semaphore unlock operation	
1209	An operation that is applied to a semaphore. If, prior to the operation, there are any threads in	
1210	the set of threads awaiting the semaphore, then some thread from that set will be removed from	
1211	the set and become unblocked. Otherwise, the semaphore value is incremented.	
1212	session	
1213	A collection of process groups established for job control purposes. Each process group is a	
1214	member of a session. A process is considered to be a member of the session of which its process	
1215	group is a member. A newly created process joins the session of its creator. A process can alter	
1216	its session membership; see <i>setsid()</i> . There can be multiple process groups in the same session.	
1217	session leader	
1218	A process that has created a session; see <i>setsid()</i> .	
1219	session lifetime	
1220	The period between when a session is created and the end of the lifetime of all the process	
1221	groups that remain as members of the session.	
1222	shared memory object	
1223	An object that represents memory that can be mapped concurrently into the address space of	
1224	more than one process.	
1225	shell	
1226	A program that interprets sequences of text input as commands. It may operate on an input	
1227	stream or it may interactively prompt and read commands from a terminal.	
1228	shell, the	
1229	The XSI Shell Command Language Interpreter (see <i>sh</i>), a specific instance of a shell.	
1230	shell script	
1231	A file containing shell commands. If the file is made executable, it can be executed by specifying	
1232	its name as a simple command (see the XCU specification, Section 2.9.1, Simple Commands).	
1233	Execution of a shell script causes a shell to execute the commands within the script.	
1234	Alternatively, a shell can be requested to execute the commands in a shell script by specifying	
1235	the name of the shell script as the operand to the <i>sh</i> utility.	
1236	signal	
1237	A mechanism by which a process or thread may be notified of, or affected by, an event occurring	
1238	in the system. Examples of such events include hardware exceptions and specific actions by	
1239	processes. The term <i>signal</i> is also used to refer to the event itself.	
1240	signal stack	
1241	EX Memory established for a thread, in which signal handlers catching signals sent to that thread	
1242	are executed.	
1243	single-quote	
1244	The character "'", also known as <i>apostrophe</i> .	

1245		slash	
1246		The character "/", also known as <i>solidus</i> .	
1247		socket	
1248	EX	A communications endpoint associated with a file descriptor that provides communications	
1249		services using a specified communications protocol. See the Networking Services specification.	
1250		soft limit	
1251	EX	A resource limitation established for each process that the process may set to any value less than	
1252		or equal to the hard limit.	
1253		source code	
1254		When dealing with the XSI Shell Command Language, input to the command language	
1255		interpreter. The term <i>shell script</i> is synonymous with this meaning.	
1256		When dealing with the C language, input to a C compiler conforming to the ISO C standard.	
1257		When dealing with another XSI-compliant language, input to a compiler conforming to that	
1258		language standard.	
1259		Source code also refers to the input statements prepared for the following standard utilities:	
1260		<i>awk, bc, ed, lex, localedef, make, sed</i> and <i>yacc</i> .	
1261		Source code can also refer to a collection of sources meeting any or all of these meanings.	
1262		special parameter	
1263		In the shell, a parameter named by a single character from the following list:	
1264		* @ # ? ! - \$ 0	
1265		See the XCU specification, Section 2.5.2, Special Parameters .	
1266		space character	
1267		The character defined in the portable character set as <space>. The space character is a member	
1268		of the space character class of the current locale, but represents the single character, and not all	
1269		of the possible members of the class. (See white space on page 36.)	
1270		standard error	
1271		An output stream usually intended to be used for diagnostic messages.	
1272		standard input	
1273		An input stream usually intended to be used for primary data input.	
1274		standard output	
1275		An output stream usually intended to be used for primary data output.	
1276		standard utilities	
1277		The utilities described in the XCU specification.	
1278		stream	
1279		Appearing in lower case, a stream is a file access object that allows access to an ordered	
1280		sequence of characters, as described by the ISO C standard. Such objects can be created by the	
1281		<i>fdopen()</i> , <i>fopen()</i> or <i>popen()</i> functions, and are associated with a file descriptor. A stream	
1282		provides the additional services of user-selectable buffering and formatted input and output.	
1283		See the XSH specification, Section 2.4, Standard I/O Streams .	
1284		STREAM	
1285	EX	Appearing in upper case, STREAM refers to a full duplex connection between a process and an	
1286		open device or pseudo-device. It optionally includes one or more intermediate processing	
1287		modules that are interposed between the process end of the STREAM and the device driver (or	
1288		pseudo-device driver) end of the STREAM. See the XSH specification, Section 2.5, STREAMS .	

1289		STREAM end
1290	EX	The STREAM end is the driver end of the STREAM and is also known as the downstream end of the STREAM.
1291		
1292		STREAM head
1293	EX	The STREAM head is the beginning of the STREAM and is at the boundary between the system and the application process. This is also known as the upstream end of the STREAM.
1294		
1295		STREAMS multiplexor
1296	EX	A driver with multiple STREAMS connected to it. Multiplexing with STREAMS connected above is referred to as N-to-1, or upper multiplexing. Multiplexing with STREAMS connected below is referred to as 1-to-N or lower multiplexing.
1297		
1298		
1299		string
1300		A contiguous sequence of bytes terminated by and including the first null byte.
1301		subshell
1302		A shell execution environment, distinguished from the main or current shell execution environment by the attributes described in the XCU specification, Section 2.12, Shell Execution Environment .
1303		
1304		
1305		successfully transferred
1306		For a write operation to a regular file, when the system ensures that all data written is readable on any subsequent open of the file (even one that follows a system or power failure) in the absence of a failure of the physical storage medium.
1307		
1308		
1309		For a read operation, when an image of the data on the physical storage medium is available to the requesting process.
1310		
1311		supplementary group ID
1312		An attribute of a process used in determining file access permissions. A process has up to {NGROUPS_MAX} supplementary group IDs in addition to the effective group ID. The supplementary group IDs of a process are set to the supplementary group IDs of the parent process when the process is created. Whether a process' effective group ID is included in or omitted from its list of supplementary group IDs is unspecified.
1313		
1314		
1315		
1316		
1317		suspended job
1318		A job that has received a SIGSTOP, SIGTSTP, SIGTTIN or SIGTTOU signal that caused the process group to stop. A suspended job is a background job, but a background job is not necessarily a suspended job.
1319		
1320		
1321		symbolic link
1322	EX	A type of file that contains a pathname. The pathname is interpolated into a pathname being resolved, during pathname resolution, to create a new pathname when it is encountered.
1323		
1324		synchronised I/O completion
1325		The state of an I/O operation that has either been successfully transferred or diagnosed as unsuccessful.
1326		
1327		synchronised I/O data integrity completion
1328		• For read, when the operation has been completed or diagnosed if unsuccessful. The read is complete only when an image of the data has been successfully transferred to the requesting process. If there were any pending write requests affecting the data to be read at the time that the synchronised read operation was requested, these write requests shall be successfully transferred prior to reading the data.
1329		
1330		
1331		
1332		
1333		• For write, when the operation has been completed or diagnosed if unsuccessful. The write is complete only when the data specified in the write request is successfully transferred and all
1334		

1335		file system information required to retrieve the data is successfully transferred.	
1336		File attributes that are not necessary for data retrieval (access time, modification time, status	
1337		change time) need not be successfully transferred prior to returning to the calling process.	
1338		synchronised I/O file integrity completion	
1339		Identical to a synchronised I/O data integrity completion with the addition that all file attributes	
1340		relative to the I/O operation (including access time, modification time, status change time) will	
1341		be successfully transferred prior to returning to the calling process.	
1342		synchronised I/O operation	
1343		An I/O operation performed on a file that provides the application assurance of the integrity of	
1344		its data and files.	
1345		synchronous I/O operation	
1346		An I/O operation that causes the thread requesting the I/O to be blocked from further use of the	
1347		processor until that I/O operation completes.	
1348		Note that a synchronous I/O operation does not imply synchronised I/O data integrity	
1349		completion or synchronised I/O file integrity completion.	
1350		synchronously generated signal	
1351		A signal that is attributable to a specific thread.	
1352		For example, a thread executing an illegal instruction or touching invalid memory causes a	
1353		synchronously generated signal. Being synchronous is a property of how the signal was	
1354		generated and not a property of the signal number.	
1355		system	
1356		An implementation of the XSI.	
1357		system crash	
1358		An interval initiated by an unspecified circumstance that causes all processes (possibly other	
1359		than special system processes) to be terminated in an undefined manner, after which any	
1360		changes to the state and contents of files created or written to by an application prior to the	
1361	EX	interval are undefined, except as required elsewhere in this specification set .	
1362		system console	
1363	EX	An optional file that receives messages sent by <i>fmtmsg()</i> when the MM_CONSOLE flag is set.	
1364		system documentation	
1365		All documentation provided with an XSI-conformant implementation except for the	
1366		Conformance Statement Questionnaire (CSQ). Electronically distributed documents for an XSI-	
1367		conformant implementation are considered part of the system documentation.	
1368		system process	
1369		An implementation-dependent object, other than a process executing an application, that has a	
1370		process ID.	
1371		system scheduling priority	
1372		A number used as advice to the system to alter process scheduling priorities. Raising the value	
1373		should give a process additional preference when scheduling a process to run. Lowering the	
1374		value should reduce the preference and make a process less likely to run. Typically, a process	
1375		with higher system scheduling priority will run to completion more quickly than an equivalent	
1376		process with lower system scheduling priority. A scheduling priority of zero specifies the	
1377		default policy of the system.	
1378		This definition is not intended to suggest that all processes in a system have priorities that are	
1379		comparable. Scheduling policy extensions such as adding real-time priorities make the notion of	

a single underlying priority for all scheduling policies problematic. Some systems may implement the features related to *nice* to affect all processes on the system, others to affect just the general time-sharing activities implied by this specification set, and others may have no effect at all. Because of the use of “implementation-dependent” in *nice* and *renice*, a wide range of implementation strategies is possible.

system reboot

An implementation-dependent sequence of events that may result in the loss of transitory data; that is, data that is not saved in permanent storage. For example, message queues, shared memory, semaphores and processes.

tab character

A character that in the output stream indicates that printing or displaying should start at the next horizontal tabulation position on the current line. The tab is the character designated by ‘\t’ in the C language. If the current position is at or past the last defined horizontal tabulation position, the behaviour is unspecified. It is unspecified whether this character is the exact sequence transmitted to an output device by the system to accomplish the tabulation.

terminal

(Or **terminal device**.) A character special file that obeys the specifications of the general terminal interface as described in Chapter 9 on page 119.

text column

A roughly rectangular block of characters capable of being laid out side-by-side next to other text columns on an output page or terminal screen. The widths of text columns are measured in column positions.

text file

A file that contains characters organised into one or more lines. The lines must not contain NUL characters and none can exceed {LINE_MAX} bytes in length, including the newline character. Although the XSI does not distinguish between text files and binary files (see the ISO C standard), many utilities only produce predictable or meaningful output when operating on text files. The standard utilities that have such restrictions always specify *text files* in their **STDIN** or **INPUT FILES** sections.

The term *text file* does not prevent the inclusion of control or other non-printable characters (other than NUL). Therefore, standard utilities that list text files as inputs or outputs are either able to process the special characters gracefully or they explicitly describe their limitations within their individual sections. The only difference between text and binary files is that text files have lines of less than {LINE_MAX} bytes, with no NUL characters, each terminated by a newline character. The definition allows a file with a single newline character, but not a totally empty file, to be called a text file. If a file ends with an incomplete line it is not strictly a text file by this definition. The newline character referred to in this specification set is not some generic line separator, but a single character; files created on systems where they use multiple characters for ends of lines are not portable to all XSI-conformant systems without some translation process.

thread

A single flow of control within a process. Each thread has its own thread ID, scheduling priority and policy, *errno* value, thread-specific key/value bindings, and the required system resources to support a flow of control. Anything whose address may be determined by a thread, including but not limited to static variables, storage obtained via *malloc()*, directly addressable storage obtained through implementation-supplied functions and automatic variables, are accessible to all threads in the same process.

thread ID

Each thread in a process is uniquely identified during its lifetime by a value of type **pthread_t** called a thread ID.

thread list

An ordered set of runnable processes that all have the same ordinal value for their priority.

The ordering of processes on the list is determined by a scheduling policy or policies. The set of thread lists includes all runnable processes in the system.

thread-safe

A function that may be safely invoked concurrently by multiple threads. Examples are any “pure” function, a function which holds a mutex locked while it is accessing static storage, or objects shared among threads.

thread-specific data key

A process global handle of type **pthread_key_t** which is used for naming thread-specific data.

Although the same key value may be used by different threads, the values bound to the key by *pthread_setspecific()* and accessed by *pthread_getspecific()* are maintained on a per-thread basis and persist for the life of the calling thread.

tilde

The character ~.

timer

A mechanism that can notify a thread when the time as measured by a particular clock has reached or passed a specified value, or when a specified amount of time has passed.

timer overrun

A condition that occurs each time a timer, for which there is already an expiration signal queued to the process, expires.

token

A sequence of characters that the shell considers as a single unit when reading input, according to the rules in the XCU specification, **Section 2.3, Token Recognition**. A token is either an operator or a word.

upshifting

The conversion of a lower-case character to its upper-case representation.

user database

A system database of implementation-dependent format that contains at least the following information for each user ID:

- User name
- Numerical user ID
- Initial numerical group ID
- Initial working directory
- Initial user program.

The initial numerical group ID is used by the *newgrp* utility. Any other circumstances under which the initial values are operative are implementation-dependent.

If the initial user program field is null, an implementation-dependent program is used.

If the initial working directory field is null, the interpretation of that field is implementation-dependent.

user ID

A non-negative integer that is used to identify a system user. When the identity of a user is associated with a process, a user ID value is referred to as a real user ID, an effective user ID or a saved set-user-ID.

user name

A string that is used to identify a user, as described in **user database** on page 34. To be portable across XSI-conformant systems, the value must be composed of characters from the portable filename character set. The hyphen should not be used as the first character of a portable user name.

utility

A program that can be called by name from a shell to perform a specific task, or related set of tasks. This program is either an executable file, such as might be produced by a compiler or linker system from computer source code, or a file of shell source code, directly interpreted by the shell. The program may have been produced by the user, provided by the system implementor, or acquired from an independent distributor. The term *utility* does not apply to the special built-in utilities provided as part of the XSI Shell Command Language; see the XCU specification, **Section 2.14, Special Built-in Utilities**. The system may implement certain utilities as shell functions (see the XCU specification, **Section 2.9.5, Function Definition Command**) or built-in utilities, but only an application that is aware of the command search order described in the XCU specification, **Command Search and Execution in Section 2.9.1** or of performance characteristics can discern differences between the behaviour of such a function or built-in utility and that of a true executable file.

variable

In the shell, a named parameter. See the XCU specification, **Section 2.5, Parameters and Variables**.

variable assignment

In the shell, a word consisting of the following parts:

varname=value

When used in a context where assignment is defined to occur (see the XCU specification, **Section 2.9.1, Simple Commands**) and at no other time, the *value* (representing a word or field) will be assigned as the value of the variable denoted by *varname*. The *varname* and *value* parts meet the requirements for a name and a word, respectively, except that they are delimited by the embedded unquoted equals-sign in addition to the delimiting described in the XCU specification, **Section 2.3, Token Recognition**. In all cases, the variable will be created if it did not already exist. If *value* is not specified, the variable will be given a null value.

An alternative form of variable assignment:

symbol=value

(where *symbol* is a valid word delimited by an equals-sign, but not a valid name) produces unspecified results. This form is used by the KornShell *name[expression]=value* syntax.

vertical-tab character

A character that in the output stream indicates that printing should start at the next vertical tabulation position. The vertical-tab is the character designated by '\v' in the C language. If the current position is at or past the last defined vertical tabulation position, the behaviour is unspecified. It is unspecified whether this character is the exact sequence transmitted to an output device by the system to accomplish the tabulation.

1515 **white space**

1516 A sequence of one or more characters that belong to the **space** character class as defined via the
 1517 LC_CTYPE category in the current locale.

1518 In the POSIX locale, white space consists of one or more blank characters (space and tab
 1519 characters), newline characters, carriage-return characters, form-feed characters and vertical-tab
 1520 characters.

1521 **wide-character code (C language)**

1522 An integer value corresponding to a single graphic symbol or control code. See Section 4.3 on
 1523 page 45.

1524 **wide-character string**

1525 A contiguous sequence of wide-character codes terminated by and including the first null wide-
 1526 character code.

1527 **word**

1528 In the shell, a token other than an operator. In some cases a word is also a portion of a word
 1529 token: in the various forms of parameter expansion (see the XCU specification, **Section 2.6.2,**
 1530 **Parameter Expansion**), such as $\${name-word}$, and variable assignment, such as $name=word$, the
 1531 word is the portion of the token depicted by *word*. The concept of a word is no longer applicable
 1532 following word expansions only fields remain; see the XCU specification, **Section 2.6, Word**
 1533 **Expansions**.

1534 **working directory**

1535 (Or **current working directory**.) A directory, associated with a process, that is used in pathname
 1536 resolution for pathnames that do not begin with a slash.

1537 **world-wide portability interface**

1538 Functions for handling characters in a codeset-independent manner.

1539 **write**

1540 To output characters to a file, such as standard output or standard error. Unless otherwise
 1541 stated, standard output is the default output destination for all uses of the term *write*. See the
 1542 distinction between *display* and *write* in **display** on page 11.

1543 **XSI-conformant**

1544 A system which allows an application to be built using a set of services that are consistent across
 1545 all systems that conform to this specification set.

1546 **zombie process**

1547 An inactive process that will be deleted at some later time when its parent process executes
 1548 *wait()* or *waitpid()*.

1549 **[n, m] and [n, m)**

1550 Notations denoting mathematical ranges. The square brackets [and] include the limit; the
 1551 parentheses (and) exclude the limit; that is, if x is in $[0, 1]$, it can be from 0 to 1 inclusive, but if x
 1552 is in $[0, 1)$, it can be from 0 up to but not including 1.

1553 **± 0**

1554 The algebraic sign provides additional information about any variable that has the value zero.
 1555 Although all precisions have distinct representations for +0, -0, +Inf and -Inf, the signs are
 1556 significant in some circumstances, such as division by zero, and not in others.

1557 **CHANGE HISTORY**

1558 **Issue 4**

1559 Numerous changes and additions are made for alignment with the ISO C standard and the
1560 ISO POSIX-1 standard.

1561 **Issue 4, Version 2**

1562 The following terms are added to support the adoption of additional traditional UNIX
1563 interfaces: *alternate signal stack*, *break value*, *data segment*, *driver*, *hard limit*, *host byte order*,
1564 *named STREAM*, *network byte order*, *network host database*, *network net database*, *network*
1565 *protocol database*, *network service database*, *pad*, *parent window*, *priority band*, *process virtual*
1566 *time*, *pseudo-terminal*, *real time*, *signal stack*, *socket*, *soft limit*, *STREAM* (second definition),
1567 *STREAM end*, *STREAM head*, *STREAMS multiplexor*, *symbolic link*, *system console* and *timer*.

1568 **Issue 5**

1569 Numerous terms are added to support adoption of the ISO POSIX Threads Extension and
1570 the ISO POSIX Realtime Extension.

File Format Notation

1572

1573 The **STDIN**, **STDOUT**, **STDERR**, **INPUT FILES** and **OUTPUT FILES** sections of the utility
 1574 descriptions use a syntax to describe the data organisation within the files, when that
 1575 organisation is not otherwise obvious. The syntax is similar to that used by the **XSH**
 1576 specification *printf()* function, as described in this chapter. When used in **STDIN** or **INPUT**
 1577 **FILES** sections of the utility descriptions, this syntax describes the format that could have been
 1578 used to write the text to be read, not a format that could be used by the *scanf()* function to read
 1579 the input file.

1580 The description of an individual record is as follows:

1581 "*<format>*", [*<arg1>*, *<arg2>*, . . . , *<argn>*]

1582 The *format* is a character string that contains three types of objects defined below:

1583 *characters*

1584 Characters that are not *escape sequences* or *conversion specifications*, as described below, are
 1585 copied to the output.

1586 *escape sequences*

1587 Represent non-graphic characters.

1588 *conversion specifications*

1589 Specifies the output format of each argument. (See below.)

1590 The following characters have the following special meaning in the format string:

1591 " " (An empty character position.) One or more blank characters.

1592 Δ Exactly one space character.

1593 The notation for spaces allows some flexibility for application output. Note that an empty
 1594 character position in *format* represents one or more blank characters on the output (not *white*
 1595 *space*, which can include newline characters). Therefore, another utility that reads that output as
 1596 its input must be prepared to parse the data using *scanf()*, *awk*, and so forth. The Δ character is
 1597 used when exactly one space character is output.

1598 The following table lists escape sequences and associated actions on display devices capable of
 1599 the action.

Escape Sequence	Represents Character	Terminal Action
<code>\\</code>	backslash	None.
<code>\a</code>	alert	Attempts to alert the user through audible or visible notification.
<code>\b</code>	backspace	Moves the printing position to one column before the current position, unless the current position is the start of a line.
<code>\f</code>	form-feed	Moves the printing position to the initial printing position of the next logical page.
<code>\n</code>	newline	Moves the printing position to the start of the next line.
<code>\r</code>	carriage-return	Moves the printing position to the start of the current line.
<code>\t</code>	tab	Moves the printing position to the next tab position on the current line. If there are no more tab positions left on the line, the behaviour is undefined.
<code>\v</code>	vertical-tab	Moves the printing position to the start of the next vertical tab position. If there are no more vertical tab positions left on the page, the behaviour is undefined.

Table 3-1 Escape Sequences and Associated Actions

Each conversion specification is introduced by the percent-sign character (%). After the character %, the following appear in sequence:

flags Zero or more *flags*, in any order, that modify the meaning of the conversion specification.

field width An optional string of decimal digits to specify a minimum *field width*. For an output field, if the converted value has fewer bytes than the field width, it is padded on the left (or right, if the left-adjustment flag (-), described below, has been given to the field width).

precision Gives the minimum number of digits to appear for the d, o, i, u, x or X conversions (the field is padded with leading zeros), the number of digits to appear after the radix character for the e and f conversions, the maximum number of significant digits for the g conversion; or the maximum number of bytes to be written from a string in s conversion. The precision takes the form of a period (.) followed by a decimal digit string; a null digit string is treated as zero.

conversion characters

A conversion character (see below) that indicates the type of conversion to be applied.

The *flag* characters and their meanings are:

- The result of the conversion is left-justified within the field.

+ The result of a signed conversion always begins with a sign (+ or -).

<space> If the first character of a signed conversion is not a sign, a space character is prefixed to the result. This means that if the space character and + flags both appear, the space character flag is ignored.

1641	#	The value is to be converted to an alternative form. For c, d, i, u and s conversions, the behaviour is undefined. For o conversion, it increases the precision to force the first digit of the result to be a zero. For x or X conversion, a non-zero result has 0x or 0X prefixed to it, respectively. For e, E, f, g and G conversions, the result always contains a radix character, even if no digits follow the radix character. For g and G conversions, trailing zeros are not removed from the result as they usually are.
1642		
1643		
1644		
1645		
1646		
1647	0	For d, i, o, u, x, X, e, E, f, g and G conversions, leading zeros (following any indication of sign or base) are used to pad to the field width; no space padding is performed. If the 0 and – flags both appear, the 0 flag is ignored. For d, i, o, u, x and X conversions, if a precision is specified, the 0 flag is ignored. For other conversions, the behaviour is undefined.
1648		
1649		
1650		
1651		
1652		Each conversion character results in fetching zero or more arguments. The results are undefined if there are insufficient arguments for the format. If the format is exhausted while arguments remain, the excess arguments are ignored.
1653		
1654		
1655		The <i>conversion characters</i> and their meanings are:
1656	d,i,o,u,x,X	The integer argument is written as signed decimal (d or i), unsigned octal (o), unsigned decimal (u), or unsigned hexadecimal notation (x and X). The d and i specifiers convert to signed decimal in the style <code>[-]dddd</code> . The x conversion uses the numbers and letters 0123456789abcdef and the X conversion uses the numbers and letters 0123456789ABCDEF. The <i>precision</i> component of the argument specifies the minimum number of digits to appear. If the value being converted can be represented in fewer digits than the specified minimum, it is expanded with leading zeros. The default precision is 1. The result of converting a zero value with a precision of 0 is no characters. If both the field width and precision are omitted, the implementation may precede, follow or precede and follow numeric arguments of types d, i and u with blank characters; arguments of type o (octal) may be preceded with leading zeros.
1657		
1658		
1659		
1660		
1661		
1662		
1663		
1664		
1665		
1666		
1667		
1668		The treatment of integers and spaces is different from the <i>printf()</i> function in that they can be surrounded with blank characters. This was done so that, given a format such as:
1669		
1670		
1671		<code>"%d\n", <foo></code>
1672		the implementation could use a <i>printf()</i> call such as:
1673		<code>printf("%6d\n", foo);</code>
1674		and still conform. This notation is thus somewhat like <i>scanf()</i> in addition to <i>printf()</i> .
1675		
1676	f	The floating point number argument is written in decimal notation in the style <code>[-]ddd.ddd</code> , where the number of digits after the radix character (shown here as a decimal point) is equal to the <i>precision</i> specification. The LC_NUMERIC locale category determines the radix character to use in this format. If the <i>precision</i> is omitted from the argument, six digits are written after the radix character; if the <i>precision</i> is explicitly 0, no radix character appears.
1677		
1678		
1679		
1680		
1681		
1682	e,E	The floating point number argument is written in the style <code>[-]d.ddde±dd</code> (the symbol ± indicates either a plus or minus sign), where there is one digit before the radix character (shown here as a decimal point) and the number of digits after it is equal to the precision. The LC_NUMERIC locale category determines the radix character to use in this format. When the precision is missing, six digits are written after the radix character; if the precision is 0, no radix character appears. The E
1683		
1684		
1685		
1686		
1687		

1688		conversion character produces a number with E instead of e introducing the
1689		exponent. The exponent always contains at least two digits. However, if the value
1690		to be written requires an exponent greater than two digits, additional exponent
1691		digits are written as necessary.
1692	g,G	The floating point number argument is written in style f or e (or in style E in the
1693		case of a G conversion character), with the precision specifying the number of
1694		significant digits. The style used depends on the value converted: style g is used
1695		only if the exponent resulting from the conversion is less than -4 or greater than or
1696		equal to the precision. Trailing zeros are removed from the result. A radix
1697		character appears only if it is followed by a digit.
1698	c	The integer argument is converted to an unsigned char and the resulting byte is
1699		written.
1700	s	The argument is taken to be a string and bytes from the string are written until the
1701		end of the string or the number of bytes indicated by the <i>precision</i> specification of
1702		the argument is reached. If the precision is omitted from the argument, it is taken
1703		to be infinite, so all bytes up to the end of the string are written.
1704	%	Write a % character; no argument is converted.
1705		In no case does a non-existent or insufficient <i>field width</i> cause truncation of a field; if the result of
1706		a conversion is wider than the field width, the field is simply expanded to contain the conversion
1707		result. The term <i>field width</i> should not be confused with the term <i>precision</i> used in the
1708		description of %s.
1709		One difference from the C function <i>printf()</i> is that the l and h conversion characters are not used.
1710		As expressed by the XCU specification, there is no differentiation between decimal values for
1711		type int , type long or type short . The specifications %d or %i should be interpreted as an
1712		arbitrary length sequence of digits. Also, no distinction is made between single precision and
1713		double precision numbers (float or double in C). These are simply referred to as floating point
1714		numbers.
1715		Many of the output descriptions in the XCU specification use the term <i>line</i> , such as:
1716		"%s",<input line>
1717		Since the definition of <i>line</i> includes the trailing newline character already, there is no need to
1718		include a \n in the format; a double newline character would otherwise result.
1719	Examples	
1720		To represent the output of a program that prints a date and time in the form Sunday, July 3,
1721		10:02, where <weekday> and <month> are strings:
1722		"%s,%s%d,%d:%.2d\n",<weekday>,<month>,<day>,<hour>,<min>
1723		To show π written to 5 decimal places:
1724		"pi=%f\n",<value of π >
1725		To show an input file format consisting of five colon-separated fields:
1726		"%s:%s:%s:%s:%s\n",<arg1>,<arg2>,<arg3>,<arg4>,<arg5>
1727		

Character Set

1728

4.1 Portable Character Set

Conforming implementations support one or more coded character sets. Each supported locale includes the *portable character set* specified in the following table.

Symbolic Name	Glyph	Symbolic Name	Glyph	Symbolic Name	Glyph
<NUL>		<colon>	:	<circumflex>	^
<alert>		<semicolon>	;	<circumflex-accent>	^
<backspace>		<less-than-sign>	<	<underscore>	_
<tab>		<equals-sign>	=	<underline>	_
<newline>		<greater-than-sign>	>	<low-line>	_
<vertical-tab>		<question-mark>	?	<grave-accent>	`
<form-feed>		<commercial-at>	@	<a>	a
<carriage-return>		<A>	A		b
<space>			B	<c>	c
<exclamation-mark>	!	<C>	C	<d>	d
<quotation-mark>	"	<D>	D	<e>	e
<number-sign>	#	<E>	E	<f>	f
<dollar-sign>	\$	<F>	F	<g>	g
<percent-sign>	%	<G>	G	<h>	h
<ampersand>	&	<H>	H	<i>	i
<apostrophe>	'	<I>	I	<j>	j
<left-parenthesis>	(<J>	J	<k>	k
<right-parenthesis>)	<K>	K	<l>	l
<asterisk>	*	<L>	L	<m>	m
<plus-sign>	+	<M>	M	<n>	n
<comma>	,	<N>	N	<o>	o
<hyphen>	-	<O>	O	<p>	p
<hyphen-minus>	-	<P>	P	<q>	q
<period>	.	<Q>	Q	<r>	r
<full-stop>	.	<R>	R	<s>	s
<slash>	/	<S>	S	<t>	t
<solidus>	/	<T>	T	<u>	u
<zero>	0	<U>	U	<v>	v
<one>	1	<V>	V	<w>	w
<two>	2	<W>	W	<x>	x
<three>	3	<X>	X	<y>	y
<four>	4	<Y>	Y	<z>	z
<five>	5	<Z>	Z	<left-brace>	{
<six>	6	<left-square-bracket>	[<left-curly-bracket>	{
<seven>	7	<backslash>	\	<vertical-line>	
<eight>	8	<reverse-solidus>	\	<right-brace>	}
<nine>	9	<right-square-bracket>]	<right-curly-bracket>	}
				<tilde>	~

1772

Table 4-1 Portable Character Set

1773 Table 4-1 on page 43 defines the characters in the portable character set and the corresponding
 1774 symbolic character names used to identify each character in a character set description file. The
 1775 table contains more than one symbolic character name for characters whose traditional name
 1776 differs from the chosen name.

1777 This specification set places only the following requirements on the encoded values of the
 1778 characters in the portable character set:

- 1779 • If the encoded values associated with each member of the portable character set are not
 1780 invariant across all locales supported by the implementation, the results achieved by an
 1781 application accessing those locales are unspecified.
- 1782 • The encoded values associated with the digits 0 to 9 will be such that the value of each
 1783 character after 0 will be one greater than the value of the previous character.
- 1784 • A null character, NUL, which has all bits set to zero, will be in the set of characters.
- 1785 • The encoded values associated with the members of the portable character set are each
 1786 represented in a single byte. Moreover, if the value is stored in an object of C-language type
 1787 **char**, it is guaranteed to be positive (except the NUL, which is always zero).

1788 4.2 Character Encoding

1789 The POSIX locale contains the characters in Table 4-1 on page 43, which have the properties
 1790 listed in Section 5.3.1 on page 52. Implementations may also add other characters. In other
 1791 locales, the presence, meaning and representation of any additional characters is locale-specific.

1792 In locales other than the POSIX locale, a character may have a state-dependent encoding. There
 1793 are two types of these encodings:

- 1794 • A single-shift encoding (where each character not in the initial shift state is preceded by a
 1795 shift code) can be defined if each shift-code and character sequence is considered a multi-
 1796 byte character. This is done using the concatenated-constant format in a character set
 1797 description file, as described in Section 4.4 on page 45. If the implementation supports a
 1798 character encoding of this type, all of the standard utilities in the **XCU** specification will
 1799 support it. Use of a single-shift encoding with any of the functions in the **XSH** specification
 1800 that do not specifically mention the effects of state-dependent encoding is implementation-
 1801 dependent.
- 1802 • A locking-shift encoding (where the state of the character is determined by a shift code that
 1803 may affect more than the single character following it) cannot be defined with the current
 1804 character set description file format. Use of a locking-shift encoding with any of the standard
 1805 utilities in the **XCU** specification or with any of the functions in the **XSH** specification that do
 1806 not specifically mention the effects of state-dependent encoding is implementation-
 1807 dependent.

1808 While in the initial shift state, all characters in the portable character set retain their usual
 1809 interpretation and do not alter the shift state. The interpretation for subsequent bytes in the
 1810 sequence is a function of the current shift state. A byte with all bits zero is interpreted as the null
 1811 character independent of shift state. Thus a byte with all bits zero must never occur in the
 1812 second or subsequent bytes of a character.

1813 The maximum allowable number of bytes in a character in the current locale is indicated by
 1814 **MB_CUR_MAX**, defined in the **XSH** specification `<stdlib.h>`, and by the `<mb_cur_max>` value
 1815 in a character set description file; see Section 4.4 on page 45. The implementation's maximum
 1816 number of bytes in a character is defined by the C-language macro `{MB_LEN_MAX}`.

4.3 C Language Wide-character Codes

In the shell, the standard utilities are written so that the encodings of characters are described by the locale's LC_CTYPE definition (see Section 5.3.1 on page 52) and there is no differentiation between characters consisting of single octets (8-bit bytes), larger bytes, or multiple bytes. However, in the C language, a differentiation is made. To ease the handling of variable length characters, the C language has introduced the concept of wide character codes.

All wide-character codes in a given process consist of an equal number of bits. This is in contrast to characters, which can consist of a variable number of bytes. The byte or byte sequence that represents a character can also be represented as a wide-character code. Wide-character codes thus provide a uniform size for manipulating text data. A wide-character code having all bits zero is the null wide-character code (see **null wide-character code** on page 21), and terminates wide-character strings (see Section 4.3). The wide-character value for each member of the Portable Character Set will equal its value when used as the lone character in an integer character constant. Wide-character codes for other characters are locale- and implementation-dependent. State shift bytes do not have a wide-character code representation.

4.4 Character Set Description File

Implementations provide a character set description file for at least one coded character set supported by the implementation. These files are referred to elsewhere in this specification set as *charmap* files. It is implementation-dependent whether or not users or applications can provide additional character set description files.

This specification set does not require that multiple character sets or codesets be supported. Although multiple charmap files are supported, it is the responsibility of the implementation to provide the file or files; if only one is provided, only that one will be accessible using the *localedef* utility's *-f* option (although in the case of just one file on the system, *-f* is not useful).

Each character set description file defines characteristics for the coded character set and the encoding for the characters specified in Table 4-1 on page 43 and may define encoding for additional characters supported by the implementation. Other information about the coded character set may also be in the file. Coded character set character values are defined using symbolic character names followed by character encoding values.

The character set description file provides:

- The capability to describe character set attributes (such as collation order or character classes) independent of character set encoding, and using only the characters in the portable character set. This makes it possible to create generic *localedef* source files for all codesets that share the portable character set (such as the ISO 8859 family or IBM Extended ASCII).
- Standardised symbolic names for all characters in the portable character set, making it possible to refer to any such character regardless of encoding.

The charmap file was introduced to resolve problems with the portability of, especially, *localedef* sources. This specification set assumes that the portable character set is constant across all locales, but does not prohibit implementations from supporting two incompatible codings, such as both ASCII and EBCDIC. Such dual-support implementations should have all charmaps and *localedef* sources encoded using one portable character set, in effect cross-compiling for the other environment. Naturally, charmaps (and *localedef* sources) are only portable without transformation between systems using the same encodings for the portable character set. They can, however, be transformed between two sets using only a subset of the actual characters (the portable set). However, the particular coded character set used for an application or an

implementation does not necessarily imply different characteristics or collation; on the contrary, these attributes should in many cases be identical, regardless of codeset. The charmap provides the capability to define a common locale definition for multiple codesets (the same *localedef* source can be used for codesets with different extended characters; the ability in the charmap to define empty names allows for characters missing in certain codesets).

Each symbolic name specified in Table 4-1 on page 43 is included in the file and is mapped to a unique encoding value (except for those symbolic names that are shown with identical glyphs). If the control characters commonly associated with the symbolic names in the following table are supported by the implementation, the symbolic names and their corresponding encoding values are included in the file. Some of the encodings associated with the symbolic names in this table may be the same as characters in the portable character set table.

<ACK>	<DC2>	<ENQ>	<FS>	<IS4>	<SOH>
<BEL>	<DC3>	<EOT>	<GS>	<LF>	<STX>
<BS>	<DC4>	<ESC>	<HT>	<NAK>	<SUB>
<CAN>		<ETB>	<IS1>	<RS>	<SYN>
<CR>	<DLE>	<ETX>	<IS2>	<SI>	<US>
<DC1>		<FF>	<IS3>	<SO>	<VT>

Table 4-2 Control Character Set

The following declarations can precede the character definitions. Each must consist of the symbol shown in the following list, starting in column 1, including the surrounding brackets, followed by one or more blank characters, followed by the value to be assigned to the symbol.

<code_set_name>	The name of the coded character set for which the character set description file is defined. The characters of the name must be taken from the set of characters with visible glyphs defined in Table 4-1 on page 43.
<mb_cur_max>	The maximum number of bytes in a multi-byte character. This defaults to 1.
<mb_cur_min>	An unsigned positive integer value that defines the minimum number of bytes in a character for the encoded character set. On XSI-conformant systems, <mb_cur_min> is always 1.
<escape_char>	The escape character used to indicate that the characters following will be interpreted in a special way, as defined later in this section. This defaults to backslash (\), which is the character glyph used in all the following text and examples, unless otherwise noted.
<comment_char>	The character that when placed in column 1 of a charmap line, is used to indicate that the line is to be ignored. The default character is the number sign (#).

The character set mapping definitions will be all the lines immediately following an identifier line containing the string **CHARMAP** starting in column 1, and preceding a trailer line containing the string **END CHARMAP** starting in column 1. Empty lines and lines containing a **<comment_char>** in the first column will be ignored. Each non-comment line of the character set mapping definition (that is, between the **CHARMAP** and **END CHARMAP** lines of the file) must be in either of two forms:

```
"%s %s %s\n", <symbolic-name>, <encoding>, <comments>
```

or:

```
"%s...%s %s %s\n", <symbolic-name>, <symbolic-name>, <encoding>,  
<comments>
```

In the first format, the line in the character set mapping definition defines a single symbolic name and a corresponding encoding. A symbolic name is one or more characters from the set shown with visible glyphs in Table 4-1 on page 43, enclosed between angle brackets. A character following an escape character is interpreted as itself; for example, the sequence **<\\>** represents the symbolic name **>** enclosed between angle brackets.

In the second format, the line in the character set mapping definition defines a range of one or more symbolic names. In this form, the symbolic names must consist of zero or more non-numeric characters from the set shown with visible glyphs in Table 4-1 on page 43, followed by an integer formed by one or more decimal digits. The characters preceding the integer must be identical in the two symbolic names, and the integer formed by the digits in the second symbolic name must be equal to or greater than the integer formed by the digits in the first name. This is interpreted as a series of symbolic names formed from the common part and each of the integers between the first and the second integer, inclusive. As an example, **<j0101>...<j0104>** is interpreted as the symbolic names **<j0101>**, **<j0102>**, **<j0103>** and **<j0104>**, in that order.

A character set mapping definition line must exist for all symbolic names specified in Table 4-1 on page 43, and must define the coded character value that corresponds to the character glyph indicated in the table, or the coded character value that corresponds with the control character symbolic name. If the control characters commonly associated with the symbolic names in Table 4-2 on page 46 are supported by the implementation, the symbolic name and the corresponding encoding value must be included in the file. Additional unique symbolic names may be included. A coded character value can be represented by more than one symbolic name.

The encoding part is expressed as one (for single-byte character values) or more concatenated decimal, octal or hexadecimal constants in the following formats:

```
"%cd%d", <escape_char>, <decimal byte value>
```

```
"%cx%x", <escape_char>, <hexadecimal byte value>
```

```
"%co", <escape_char>, <octal byte value>
```

Decimal constants must be represented by two or three decimal digits, preceded by the escape character and the lower-case letter d; for example, **\d05**, **\d97** or **\d143**. Hexadecimal constants must be represented by two hexadecimal digits, preceded by the escape character and the lower-case letter x; for example, **\x05**, **\x61** or **\x8f**. Octal constants must be represented by two or three octal digits, preceded by the escape character; for example, **\05**, **\141** or **\217**. In a portable charmap file, each constant must represent an 8-bit byte. Implementations supporting other byte sizes may allow constants to represent values larger than those that can be represented in 8-bit bytes, and to allow additional digits in constants. When constants are concatenated for multi-byte character values, they must be of the same type, and interpreted in byte order from first to last with the least significant byte of the multi-byte character specified by the last constant. The manner in which these constants are represented in the character stored in

1946 the system is implementation-dependent. (This big endian notation was chosen for reasons of
1947 portability. There is no requirement that the internal representation in the computer memory be
1948 in this same order.) Omitting bytes from a multi-byte character definition produces undefined
1949 results.

1950 In lines defining ranges of symbolic names, the encoded value is the value for the first symbolic
1951 name in the range (the symbolic name preceding the ellipsis). Subsequent symbolic names
1952 defined by the range will have encoding values in increasing order. For example, the line:

1953 <j0101>...<j0104> \d129\d254

1954 will be interpreted as:

1955 <j0101> \d129\d254

1956 <j0102> \d129\d255

1957 <j0103> \d130\d0

1958 <j0104> \d130\d1

1959 Note that this line will be interpreted as the example even on systems with bytes larger than 8
1960 bits.

1961 The comment is optional.

1962 For the interpretation of the dollar sign and the number sign, see **dollar sign** on page 12 and
1963 **number sign** on page 21.

1964

5.1 General

A *locale* is the definition of the subset of a user's environment that depends on language and cultural conventions. It is made up from one or more categories. Each category is identified by its name and controls specific aspects of the behaviour of components of the system. Category names correspond to the following environment variable names:

LC_CTYPE Character classification and case conversion.

LC_COLLATE Collation order.

LC_TIME Date and time formats.

LC_NUMERIC Numeric, non-monetary formatting.

LC_MONETARY Monetary formatting.

LC_MESSAGES Formats of informative and diagnostic messages and interactive responses.

The standard utilities in the **XCU** specification base their behaviour on the current locale, as defined in the **ENVIRONMENT VARIABLES** section for each utility. The behaviour of some of the C-language functions defined in the **XSH** specification will also be modified based on the current locale, as defined by the last call to *setlocale()*.

EX Locales other than those supplied by the implementation can be created by the application via the *localedef* utility, if it is provided; see the **XCU** specification. This capability is supported on all X/Open systems where the {**POSIX2_LOCALEDEF**} or {**XOPEN_XCU_VERSION**} options are supported; see the **XSH** specification **<unistd.h>**. Even if *localedef* is not provided, all implementations conforming to the **XSH** specification provide one or more locales that behave as described in this chapter. The input to the utility is described in Section 5.3 on page 50. The value that is used to specify a locale when using environment variables will be the string specified as the *name* operand to the *localedef* utility when the locale was created. The strings "C" and "POSIX" are reserved as identifiers for the POSIX locale (see Section 5.2 on page 50). When the value of a locale environment variable begins with a slash (/), it is interpreted as the pathname of the locale definition; the type of file (regular, directory, and so forth) used to store the locale definition is implementation-dependent. If the value does not begin with a slash, the mechanism used to locate the locale is implementation-dependent.

If different character sets are used by the locale categories, the results achieved by an application utilising these categories are undefined. Likewise, if different codesets are used for the data being processed by interfaces whose behaviour is dependent on the current locale, or the codeset is different from the codeset assumed when the locale was created, the result is also undefined.

Applications can select the desired locale by invoking the *setlocale()* function (or equivalent) with the appropriate value. If the function is invoked with an empty string, such as:

```
setlocale(LC_ALL, "");
```

the value of the corresponding environment variable is used. If the environment variable is unset or is set to the empty string, the implementation sets the appropriate environment as defined in Chapter 6 on page 93.

2003 5.2 POSIX Locale

2004 All systems provide a *POSIX locale*, also known as the C locale. The behaviour of standard
 2005 utilities and functions in the POSIX locale is as if the locale was defined via the *localedef* utility
 2006 with input data from the POSIX locale tables in Section 5.3.

2007 The tables in Section 5.3 describe the characteristics and behaviour of the POSIX locale for data
 2008 consisting entirely of characters from the portable character set and the control character set. For
 2009 other characters, the behaviour is unspecified. For C-language programs, the POSIX locale is the
 2010 default locale when the *setlocale()* function is not called.

2011 The POSIX locale can be specified by assigning to the appropriate environment variables the
 2012 values "C" or "POSIX".

2013 All implementations define a locale as the default locale, to be invoked when no environment
 2014 variables are set, or set to the empty string. This default locale can be the POSIX locale or any
 2015 other, implementation-dependent locale. Some implementations may provide facilities for local
 2016 installation administrators to set the default locale, customising it for each location. This
 2017 specification set does not require such a facility.

2018 5.3 Locale Definition

2019 Locales can be described with the file format presented in this section. The file format is that
 2020 accepted by the *localedef* utility. For the purposes of this section, the file is referred to as the *locale*
 2021 *definition file*, but no locales are affected by this file unless it is processed by *localedef* or some
 2022 similar mechanism. Any requirements in this section imposed upon the utility apply to *localedef*
 2023 or to any other similar utility used to install locale information using the locale definition file
 2024 format described here.

2025 The locale definition file must contain one or more locale category source definitions, and must
 2026 not contain more than one definition for the same locale category. If the file contains source
 2027 definitions for more than one category, implementation-dependent categories, if present, must
 2028 appear after the categories defined by Section 5.1 on page 49. A category source definition must
 2029 contain either the definition of a category or a **copy** directive. For a description of the **copy**
 2030 directive, see *localedef*. In the event that some of the information for a locale category, as
 2031 specified in this specification, is missing from the locale source definition, the behaviour of that
 2032 category, if it is referenced, is unspecified.

2033 A category source definition consists of a category header, a category body and a category
 2034 trailer. A category header consists of the character string naming of the category, beginning with
 2035 the characters LC_. The category trailer consists of the string END, followed by one or more
 2036 blank characters and the string used in the corresponding category header.

2037 The category body consists of one or more lines of text. Each line contains an identifier,
 2038 optionally followed by one or more operands. Identifiers are either keywords, identifying a
 2039 particular locale element, or collating elements. In addition to the keywords defined in this
 2040 specification, the source can contain implementation-dependent keywords. Each keyword
 2041 within a locale must have a unique name (that is, two categories cannot have a commonly-
 2042 named keyword); no keyword can start with the characters LC_. Identifiers must be separated
 2043 from the operands by one or more blank characters.

2044 Operands must be characters, collating elements or strings of characters. Strings must be
 2045 enclosed in double-quotes. Literal double-quotes within strings must be preceded by the *<escape*
 2046 *character>*, described below. When a keyword is followed by more than one operand, the
 2047 operands must be separated by semicolons; blank characters are allowed both before and after a
 2048 semicolon.

2049 The first category header in the file can be preceded by a line modifying the comment character.
 2050 It has the following format, starting in column 1:

2051 "comment_char %c\n",*<comment character>*

2052 The comment character defaults to the number sign (#). Blank lines and lines containing the
 2053 *<comment character>* in the first position are ignored.

2054 The first category header in the file can be preceded by a line modifying the escape character to
 2055 be used in the file. It has the following format, starting in column 1:

2056 "escape_char %c\n",*<escape character>*

2057 The escape character defaults to backslash, which is the character used in all examples shown in
 2058 this specification.

2059 A line can be continued by placing an escape character as the last character on the line; this
 2060 continuation character will be discarded from the input. Although the implementation need not
 2061 accept any one portion of a continued line with a length exceeding {LINE_MAX} bytes, it places
 2062 no limits on the accumulated length of the continued line. Comment lines cannot be continued
 2063 on a subsequent line using an escaped newline character.

2064 Individual characters, characters in strings, and collating elements must be represented using
 2065 symbolic names, as defined below. In addition, characters can be represented using the
 2066 characters themselves or as octal, hexadecimal or decimal constants. When non-symbolic
 2067 notation is used, the resultant locale definitions will in many cases not be portable between
 2068 systems. The left angle bracket (<) is a reserved symbol, denoting the start of a symbolic name;
 2069 when used to represent itself it must be preceded by the escape character. The following rules
 2070 apply to character representation:

- 2071 1. A character can be represented via a symbolic name, enclosed within angle brackets "<"
 2072 and ">". The symbolic name, including the angle brackets, must exactly match a symbolic
 2073 name defined in the charmap file specified via the *localedef -f* option, and will be replaced
 2074 by a character value determined from the value associated with the symbolic name in the
 2075 charmap file. The use of a symbolic name not found in the charmap file constitutes an
 2076 error, unless the category is LC_CTYPE or LC_COLLATE, in which case it constitutes a
 2077 warning condition (see *localedef* for a description of action resulting from errors and
 2078 warnings). The specification of a symbolic name in a **collating-element** or
 2079 **collating-symbol** section that duplicates a symbolic name in the charmap file (if present)
 2080 is an error. Use of the escape character or a right angle bracket within a symbolic name is
 2081 invalid unless the character is preceded by the escape character.

2082 **Example:**

2083 <c>;<c-cedilla> " <M><a><y> "

2. A character can be represented by the character itself, in which case the value of the character is implementation-dependent. Within a string, the double-quote character, the escape character and the right angle bracket character must be escaped (preceded by the escape character) to be interpreted as the character itself. Outside strings, the characters:

`, ; < > escape_char`

must be escaped to be interpreted as the character itself.

Example:

`c β "May"`

3. A character can be represented as an octal constant. An octal constant is specified as the escape character followed by two or more octal digits. Each constant represents a byte value. Multi-byte values can be represented by concatenated constants specified in byte order with the last constant specifying the least significant byte of the character.

Example:

`\143;\347;\143\150 "\115\141\171"`

4. A character can be represented as a hexadecimal constant. A hexadecimal constant is specified as the escape character followed by an x followed by two or more hexadecimal digits. Each constant represents a byte value. Multi-byte values can be represented by concatenated constants specified in byte order with the last constant specifying the least significant byte of the character.

Example:

`\x63;\xe7;\x63\x68 "\x4d\x61\x79"`

5. A character can be represented as a decimal constant. A decimal constant is specified as the escape character followed by a d followed by two or more decimal digits. Each constant represents a byte value. Multi-byte values can be represented by concatenated constants specified in byte order with the last constant specifying the least significant byte of the character.

Example:

`\d99;\d231;\d99\d104 "\d77\d97\d121"`

Implementations may accept single-digit octal, decimal or hexadecimal constants following the escape character. Only characters existing in the character set for which the locale definition is created can be specified, whether using symbolic names, the characters themselves, or octal, decimal or hexadecimal constants. If a charmap file is present, only characters defined in the charmap can be specified using octal, decimal or hexadecimal constants. Symbolic names not present in the charmap file can be specified and will be ignored, as specified under item 1 above.

5.3.1 LC_CTYPE

The LC_CTYPE category defines character classification, case conversion and other character attributes. In addition, a series of characters can be represented by three adjacent periods representing an ellipsis symbol (...). The ellipsis specification is interpreted as meaning that all values between the values preceding and following it represent valid characters. The ellipsis specification is valid only within a single encoded character set; that is, within a group of characters of the same size. An ellipsis is interpreted as including in the list all characters with an encoded value higher than the encoded value of the character preceding the ellipsis and lower than the encoded value of the character following the ellipsis.

2127 **Example:**

2128 `\x30; . . . ; \x39;`

2129 includes in the character class all characters with encoded values between the endpoints.

2130 The following keywords are recognised. In the descriptions, the term “automatically included”

2131 means that it is not an error either to include or omit any of the referenced characters; the

2132 implementation will provide them if missing (even if the entire keyword is missing) and accept

2133 them silently if present. When the implementation automatically includes a missing character, it

2134 will have an encoded value dependent on the charmap file in effect (see the description of the

2135 *localedef* `-f` option); otherwise, it will have a value derived from an implementation-dependent

2136 character mapping.

2137 The character classes **digit**, **xdigit**, **lower**, **upper** and **space** have a set of automatically included

2138 characters. These only need to be specified if the character values (that is, encoding) differ from

2139 the implementation default values. It is not possible to define a locale without these

2140 automatically included characters unless some implementation extension is used to prevent

2141 their inclusion. Such a definition would not be a proper superset of the C or POSIX locale and

2142 thus, it might not be possible for applications conforming to the XSI to work properly.

2143 **upper** Define characters to be classified as upper-case letters.

2144 In the POSIX locale, the 26 upper-case letters are included:

2145 `A B C D E F G H I J K L M N O P Q R S T U V W X Y Z`

2146 In a locale definition file, no character specified for the keywords **cntrl**, **digit**,

2147 **punct** or **space** can be specified. The upper-case letters A to Z, as defined in

2148 Section 4.4 on page 45 (the portable character set), are automatically included

2149 in this class.

2150 **lower** Define characters to be classified as lower-case letters.

2151 In the POSIX locale, the 26 lower-case letters are included:

2152 `a b c d e f g h i j k l m n o p q r s t u v w x y z`

2153 In a locale definition file, no character specified for the keywords **cntrl**, **digit**,

2154 **punct** or **space** can be specified. The lower-case letters a to z of the portable

2155 character set are automatically included in this class.

2156 **alpha** Define characters to be classified as letters.

2157 In the POSIX locale, all characters in the classes **upper** and **lower** are included.

2158 In a locale definition file, no character specified for the keywords **cntrl**, **digit**,

2159 **punct** or **space** can be specified. Characters classified as either **upper** or **lower**

2160 are automatically included in this class.

2161 **digit** Define the characters to be classified as numeric digits.

2162 In the POSIX locale, only:

2163 `0 1 2 3 4 5 6 7 8 9`

2164 are included.

2165 In a locale definition file, only the digits 0, 1, 2, 3, 4, 5, 6, 7, 8 and 9 can be

2166 specified, and in contiguous ascending sequence by numerical value. The

2167 digits 0 to 9 of the portable character set are automatically included in this

2168 class.

2169		The definition of character class digit requires that only ten characters the
2170		ones defining digits can be specified; alternative digits (for example, Hindi or
2171		Kanji) cannot be specified here. However, the encoding may vary if an
2172		implementation supports more than one encoding.
2173	space	Define characters to be classified as white-space characters.
2174		In the POSIX locale, at a minimum, the characters space, form-feed, newline,
2175		carriage-return, tab and vertical-tab are included.
2176		In a locale definition file, no character specified for the keywords upper ,
2177		lower , alpha , digit , graph or xdigit can be specified. The characters space,
2178		form-feed, newline, carriage-return, tab and vertical-tab of the portable
2179		character set, and any characters included in the class blank are automatically
2180		included in this class.
2181	cntrl	Define characters to be classified as control characters.
2182		In the POSIX locale, no characters in classes alpha or print are included.
2183		In a locale definition file, no character specified for the keywords upper ,
2184		lower , alpha , digit , punct , graph , print or xdigit can be specified.
2185	punct	Define characters to be classified as punctuation characters.
2186		In the POSIX locale, neither the space character nor any characters in classes
2187		alpha , digit or cntrl are included.
2188		In a locale definition file, no character specified for the keywords upper ,
2189		lower , alpha , digit , cntrl , xdigit or as the space character can be specified.
2190	graph	Define characters to be classified as printable characters, not including the
2191		space character.
2192		In the POSIX locale, all characters in classes alpha , digit and punct are
2193		included; no characters in class cntrl are included.
2194		In a locale definition file, characters specified for the keywords upper , lower ,
2195		alpha , digit , xdigit and punct are automatically included in this class. No
2196		character specified for the keyword cntrl can be specified.
2197	print	Define characters to be classified as printable characters, including the space
2198		character.
2199		In the POSIX locale, all characters in class graph are included; no characters in
2200		class cntrl are included.
2201		In a locale definition file, characters specified for the keywords upper , lower ,
2202		alpha , digit , xdigit , punct and the space character are automatically included
2203		in this class. No character specified for the keyword cntrl can be specified.
2204	xdigit	Define the characters to be classified as hexadecimal digits.
2205		In the POSIX locale, only:
2206		0 1 2 3 4 5 6 7 8 9 A B C D E F a b c d e f
2207		are included.
2208		In a locale definition file, only the characters defined for the class digit can be
2209		specified, in contiguous ascending sequence by numerical value, followed by
2210		one or more sets of six characters representing the hexadecimal digits 10 to 15

2211		inclusive, with each set in ascending order (for example A, B, C, D, E, F, a, b, c,
2212		d, e, f). The digits 0 to 9, the upper-case letters A to F and the lower-case
2213		letters a to f of the portable character set are automatically included in this
2214		class.
2215		The definition of character class xdigit requires that the characters included in
2216		character class digit be included here also.
2217	blank	Define characters to be classified as blank characters.
2218		In the POSIX locale, only the space and tab characters are included.
2219		In a locale definition file, the characters space and tab are automatically
2220		included in this class.
2221	EX charclass	Define one or more locale-specific character class names as strings separated
2222		by semicolons. Each named character class can then be defined subsequently
2223		in the LC_CTYPE definition. A character class name consists of at least one
2224		and at most {CHARCLASS_NAME_MAX} bytes of alphanumeric characters
2225		from the portable filename character set. The first character of a character
2226		class name cannot be a digit. The name cannot match any of the LC_CTYPE
2227		keywords defined in this specification.
2228	charclass-name	Define characters to be classified as belonging to the named locale-specific
2229		character class. In the POSIX locale, the locale-specific named character
2230		classes need not exist.
2231		If a class name is defined by a charclass keyword, but no characters are
2232		subsequently assigned to it, this is not an error; it represents a class without
2233		any characters belonging to it.
2234		The <i>charclass-name</i> can be used as the <i>property</i> argument to the <i>wctype()</i>
2235		function, in regular expression and shell pattern-matching bracket
2236		expressions, and by the <i>tr</i> command.
2237	toupper	Define the mapping of lower-case letters to upper-case letters.
2238		In the POSIX locale, at a minimum, the 26 lower-case characters:
2239		a b c d e f g h i j k l m n o p q r s t u v w x y z
2240		are mapped to the corresponding 26 upper-case characters:
2241		A B C D E F G H I J K L M N O P Q R S T U V W X Y Z
2242		In a locale definition file, the operand consists of character pairs, separated by
2243		semicolons. The characters in each character pair are separated by a comma
2244		and the pair enclosed by parentheses. The first character in each pair is the
2245		lower-case letter, the second the corresponding upper-case letter. Only
2246		characters specified for the keywords lower and upper can be specified. The
2247		lower-case letters a to z, and their corresponding upper-case letters A to Z, of
2248		the portable character set are automatically included in this mapping, but only
2249		when the toupper keyword is omitted from the locale definition.
2250	tolower	Define the mapping of upper-case letters to lower-case letters.
2251		In the POSIX locale, at a minimum, the 26 upper-case characters:
2252		A B C D E F G H I J K L M N O P Q R S T U V W X Y Z

2253 are mapped to the corresponding 26 lower-case characters:

2254 a b c d e f g h i j k l m n o p q r s t u v w x y z

2255 In a locale definition file, the operand consists of character pairs, separated by
 2256 semicolons. The characters in each character pair are separated by a comma
 2257 and the pair enclosed by parentheses. The first character in each pair is the
 2258 upper-case letter, the second the corresponding lower-case letter. Only
 2259 characters specified for the keywords **lower** and **upper** can be specified. If the
 2260 **tolower** keyword is omitted from the locale definition, the mapping will be
 2261 the reverse mapping of the one specified for **toupper**.

2262 **copy** Specify the name of an existing locale to be used as the definition of this
 2263 category. If this keyword is specified, no other keyword can be specified.

2264 The following table shows the character class combinations allowed.

In Class	Can Also Belong To										
	upper	lower	alpha	digit	space	cntrl	punct	graph	print	xdigit	blank
2268 upper		-	A	x	x	x	x	A	A	-	x
2269 lower	-		A	x	x	x	x	A	A	-	x
2270 alpha	-	-		x	x	x	x	A	A	-	x
2271 digit	x	x	x		x	x	x	A	A	A	x
2272 space	x	x	x	x		-	*	*	*	x	-
2273 cntrl	x	x	x	x	-		x	x	x	x	-
2274 punct	x	x	x	x	-	x		A	A	x	-
2275 graph	-	-	-	-	-	x	-		A	-	-
2276 print	-	-	-	-	-	x	-	-		-	-
2277 xdigit	-	-	-	-	x	x	x	A	A		x
2278 blank	x	x	x	x	A	-	*	*	*	x	

2279 **Table 5-1** Valid Character Class Combinations

2280 **Notes:**

- 2281 1. Explanation of codes:
- 2282 A Automatically included; see text.
- 2283 - Permitted.
- 2284 x Mutually exclusive.
- 2285 * See note 2.
- 2286 2. The space character, which is part of the **space** and **blank** classes, cannot belong
- 2287 to **punct** or **graph**, but automatically belongs to the **print** class. Other **space** or
- 2288 **blank** characters can be classified as any of **punct**, **graph** or **print**.

2289 The character classifications for the POSIX locale follow; the code listing depicting the *localedef*
 2290 input, the table representing the same information, sorted by character.

```

2291 LC_CTYPE
2292 # The following is the POSIX locale LC_CTYPE.
2293 # "alpha" is by default "upper" and "lower"
2294 # "alnum" is by definition "alpha" and "digit"
2295 # "print" is by default "alnum", "punct" and the <space> character
2296 # "graph" is by default "alnum" and "punct"
2297 #
2298 upper    <A>;<B>;<C>;<D>;<E>;<F>;<G>;<H>;<I>;<J>;<K>;<L>;<M>;\
2299          <N>;<O>;<P>;<Q>;<R>;<S>;<T>;<U>;<V>;<W>;<X>;<Y>;<Z>
2300 #
2301 lower    <a>;<b>;<c>;<d>;<e>;<f>;<g>;<h>;<i>;<j>;<k>;<l>;<m>;\
2302          <n>;<o>;<p>;<q>;<r>;<s>;<t>;<u>;<v>;<w>;<x>;<y>;<z>
2303 #
2304 digit    <zero>;<one>;<two>;<three>;<four>;<five>;<six>;\
2305          <seven>;<eight>;<nine>
2306 #
2307 space    <tab>;<newline>;<vertical-tab>;<form-feed>;\
2308          <carriage-return>;<space>
2309 #
2310 cntrl    <alert>;<backspace>;<tab>;<newline>;<vertical-tab>;\
2311          <form-feed>;<carriage-return>;\
2312          <NUL>;<SOH>;<STX>;<ETX>;<EOT>;<ENQ>;<ACK>;<SO>;\
2313          <SI>;<DLE>;<DC1>;<DC2>;<DC3>;<DC4>;<NAK>;<SYN>;\
2314          <ETB>;<CAN>;<EM>;<SUB>;<ESC>;<IS4>;<IS3>;<IS2>;\
2315          <IS1>;<DEL>
2316 #
2317 punct    <exclamation-mark>;<quotation-mark>;<number-sign>;\
2318          <dollar-sign>;<percent-sign>;<ampersand>;<apostrophe>;\
2319          <left-parenthesis>;<right-parenthesis>;<asterisk>;\
2320          <plus-sign>;<comma>;<hyphen>;<period>;<slash>;\
2321          <colon>;<semicolon>;<less-than-sign>;<equals-sign>;\
2322          <greater-than-sign>;<question-mark>;<commercial-at>;\
2323          <left-square-bracket>;<backslash>;<right-square-bracket>;\
2324          <circumflex>;<underscore>;<grave-accent>;<left-curly-bracket>;\
2325          <vertical-line>;<right-curly-bracket>;<tilde>
2326 #
2327 xdigit    <zero>;<one>;<two>;<three>;<four>;<five>;<six>;<seven>;\
2328          <eight>;<nine>;<A>;<B>;<C>;<D>;<E>;<F>;<a>;<b>;<c>;<d>;<e>;<f>
2329 #
2330 blank    <space>;<tab>
2331 #
2332 toupper  (<a>,<A>);(<b>,<B>);(<c>,<C>);(<d>,<D>);(<e>,<E>);\
2333          (<f>,<F>);(<g>,<G>);(<h>,<H>);(<i>,<I>);(<j>,<J>);\
2334          (<k>,<K>);(<l>,<L>);(<m>,<M>);(<n>,<N>);(<o>,<O>);\
2335          (<p>,<P>);(<q>,<Q>);(<r>,<R>);(<s>,<S>);(<t>,<T>);\
2336          (<u>,<U>);(<v>,<V>);(<w>,<W>);(<x>,<X>);(<y>,<Y>);(<z>,<Z>)
2337 #
2338 tolower  (<A>,<a>);(<B>,<b>);(<C>,<c>);(<D>,<d>);(<E>,<e>);\
2339          (<F>,<f>);(<G>,<g>);(<H>,<h>);(<I>,<i>);(<J>,<j>);\
2340          (<K>,<k>);(<L>,<l>);(<M>,<m>);(<N>,<n>);(<O>,<o>);\

```

```

2341      (<P>,<p>);(<Q>,<q>);(<R>,<r>);(<S>,<s>);(<T>,<t>);\
2342      (<U>,<u>);(<V>,<v>);(<W>,<w>);(<X>,<x>);(<Y>,<y>);(<Z>,<z>)
2343  END LC_CTYPE

```

	Symbolic Name	Other Case	Character Classes
2344	<NUL>		cntrl
2345	<SOH>		cntrl
2346	<STX>		cntrl
2347	<ETX>		cntrl
2348	<EOT>		cntrl
2349	<ENQ>		cntrl
2350	<ACK>		cntrl
2351	<alert>		cntrl
2352	<backspace>		cntrl
2353	<tab>		cntrl, space, blank
2354	<newline>		cntrl, space
2355	<vertical-tab>		cntrl, space
2356	<form-feed>		cntrl, space
2357	<carriage-return>		cntrl, space
2358	<SO>		cntrl
2359	<SI>		cntrl
2360	<DLE>		cntrl
2361	<DC1>		cntrl
2362	<DC2>		cntrl
2363	<DC3>		cntrl
2364	<DC4>		cntrl
2365	<NAK>		cntrl
2366	<SYN>		cntrl
2367	<ETB>		cntrl
2368	<CAN>		cntrl
2369			cntrl
2370	<SUB>		cntrl
2371	<ESC>		cntrl
2372	<IS4>		cntrl
2373	<IS3>		cntrl
2374	<IS2>		cntrl
2375	<IS1>		cntrl
2376	<space>		space, print, blank
2377	<exclamation-mark>		punct, print, graph
2378	<quotation-mark>		punct, print, graph
2379	<number-sign>		punct, print, graph
2380	<dollar-sign>		punct, print, graph
2381	<percent-sign>		punct, print, graph
2382	<ampersand>		punct, print, graph
2383	<apostrophe>		punct, print, graph
2384	<left-parenthesis>		punct, print, graph
2385	<right-parenthesis>		punct, print, graph

2389			
2390			
2391	Symbolic Name	Other Case	Character Classes
2392	<asterisk>		punct, print, graph
2393	<plus-sign>		punct, print, graph
2394	<comma>		punct, print, graph
2395	<hyphen>		punct, print, graph
2396	<period>		punct, print, graph
2397	<slash>		punct, print, graph
2398	<zero>		digit, xdigit, print, graph
2399	<one>		digit, xdigit, print, graph
2400	<two>		digit, xdigit, print, graph
2401	<three>		digit, xdigit, print, graph
2402	<four>		digit, xdigit, print, graph
2403	<five>		digit, xdigit, print, graph
2404	<six>		digit, xdigit, print, graph
2405	<seven>		digit, xdigit, print, graph
2406	<eight>		digit, xdigit, print, graph
2407	<nine>		digit, xdigit, print, graph
2408	<colon>		punct, print, graph
2409	<semicolon>		punct, print, graph
2410	<less-than-sign>		punct, print, graph
2411	<equals-sign>		punct, print, graph
2412	<greater-than-sign>		punct, print, graph
2413	<question-mark>		punct, print, graph
2414	<commercial-at>		punct, print, graph
2415	<A>	<a>	upper, xdigit, alpha, print, graph
2416			upper, xdigit, alpha, print, graph
2417	<C>	<c>	upper, xdigit, alpha, print, graph
2418	<D>	<d>	upper, xdigit, alpha, print, graph
2419	<E>	<e>	upper, xdigit, alpha, print, graph
2420	<F>	<f>	upper, xdigit, alpha, print, graph
2421	<G>	<g>	upper, alpha, print, graph
2422	<H>	<h>	upper, alpha, print, graph
2423	<I>	<i>	upper, alpha, print, graph
2424	<J>	<j>	upper, alpha, print, graph
2425	<K>	<k>	upper, alpha, print, graph
2426	<L>	<l>	upper, alpha, print, graph
2427	<M>	<m>	upper, alpha, print, graph
2428	<N>	<n>	upper, alpha, print, graph
2429	<O>	<o>	upper, alpha, print, graph
2430	<P>	<p>	upper, alpha, print, graph
2431	<Q>	<q>	upper, alpha, print, graph
2432	<R>	<r>	upper, alpha, print, graph
2433	<S>	<s>	upper, alpha, print, graph
2434	<T>	<t>	upper, alpha, print, graph
2435	<U>	<u>	upper, alpha, print, graph
2436	<V>	<v>	upper, alpha, print, graph
2437	<W>	<w>	upper, alpha, print, graph

2438
2439
2440
2441
2442
2443
2444
2445
2446
2447
2448
2449
2450
2451
2452
2453
2454
2455
2456
2457
2458
2459
2460
2461
2462
2463
2464
2465
2466
2467
2468
2469
2470
2471
2472
2473
2474
2475
2476
2477
2478
2479
2480

Symbolic Name	Other Case	Character Classes
<X>	<x>	upper, alpha, print, graph
<Y>	<y>	upper, alpha, print, graph
<Z>	<z>	upper, alpha, print, graph
<left-square-bracket>		punct, print, graph
<backslash>		punct, print, graph
<right-square-bracket>		punct, print, graph
<circumflex>		punct, print, graph
<underscore>		punct, print, graph
<grave-accent>		punct, print, graph
<a>	<A>	lower, xdigit, alpha, print, graph
		lower, xdigit, alpha, print, graph
<c>	<C>	lower, xdigit, alpha, print, graph
<d>	<D>	lower, xdigit, alpha, print, graph
<e>	<E>	lower, xdigit, alpha, print, graph
<f>	<F>	lower, xdigit, alpha, print, graph
<g>	<G>	lower, alpha, print, graph
<h>	<H>	lower, alpha, print, graph
<i>	<I>	lower, alpha, print, graph
<j>	<J>	lower, alpha, print, graph
<k>	<K>	lower, alpha, print, graph
<l>	<L>	lower, alpha, print, graph
<m>	<M>	lower, alpha, print, graph
<n>	<N>	lower, alpha, print, graph
<o>	<O>	lower, alpha, print, graph
<p>	<P>	lower, alpha, print, graph
<q>	<Q>	lower, alpha, print, graph
<r>	<R>	lower, alpha, print, graph
<s>	<S>	lower, alpha, print, graph
<t>	<T>	lower, alpha, print, graph
<u>	<U>	lower, alpha, print, graph
<v>	<V>	lower, alpha, print, graph
<w>	<W>	lower, alpha, print, graph
<x>	<X>	lower, alpha, print, graph
<y>	<Y>	lower, alpha, print, graph
<z>	<Z>	lower, alpha, print, graph
<left-curly-bracket>		punct, print, graph
<vertical-line>		punct, print, graph
<right-curly-bracket>		punct, print, graph
<tilde>		punct, print, graph
		cntrl

2481 5.3.2 LC_COLLATE

2482 The LC_COLLATE category provides a collation sequence definition for numerous utilities in
 2483 the XCU specification (*sort*, *uniq*, and so forth), regular expression matching (see Chapter 7 on
 2484 page 101) and the *strcoll()*, *strxfrm()*, *wscoll()* and *wcsxfrm()* functions in the XSH specification.

2485 A collation sequence definition defines the relative order between collating elements (characters
 2486 and multi-character collating elements) in the locale. This order is expressed in terms of
 2487 collation values; that is, by assigning each element one or more collation values (also known as
 2488 collation weights). This does not imply that implementations assign such values, but that
 2489 ordering of strings using the resultant collation definition in the locale will behave as if such
 2490 assignment is done and used in the collation process. At least the following capabilities are
 2491 provided:

- 2492 1. **Multi-character collating elements.** Specification of multi-character collating elements
 2493 (that is, sequences of two or more characters to be collated as an entity).
- 2494 2. **User-defined ordering of collating elements.** Each collating element is assigned a
 2495 collation value defining its order in the character (or basic) collation sequence. This
 2496 ordering is used by regular expressions and pattern matching and, unless collation weights
 2497 are explicitly specified, also as the collation weight to be used in sorting.
- 2498 3. **Multiple weights and equivalence classes.** Collating elements can be assigned one or
 2499 more (up to the limit {COLL_WEIGHTS_MAX}) collating weights for use in sorting. The
 2500 first weight is hereafter referred to as the primary weight.
- 2501 4. **One-to-Many mapping.** A single character is mapped into a string of collating elements.
- 2502 5. **Equivalence class definition.** Two or more collating elements have the same collation
 2503 value (primary weight).
- 2504 6. **Ordering by weights.** When two strings are compared to determine their relative order,
 2505 the two strings are first broken up into a series of collating elements; the elements in each
 2506 successive pair of elements are then compared according to the relative primary weights
 2507 for the elements. If equal, and more than one weight has been assigned, then the pairs of
 2508 collating elements are recompared according to the relative subsequent weights, until
 2509 either a pair of collating elements compare unequal or the weights are exhausted.

2510 The following keywords are recognised in a collation sequence definition. They are described in
 2511 detail in the following sections.

2512	collating-element	Define a collating-element symbol representing a multi-character collating element. This keyword is optional.
2513		
2514	collating-symbol	Define a collating symbol for use in collation order statements. This keyword is optional.
2515		
2516	order_start	Define collation rules. This statement is followed by one or more collation order statements, assigning character collation values and collation weights to collating elements.
2517		
2518		
2519	order_end	Specify the end of the collation-order statements.
2520	copy	Specify the name of an existing locale to be used as the definition of this category. If this keyword is specified, no other keyword can be specified.
2521		

The collating-element Keyword

In addition to the collating elements in the character set, the **collating-element** keyword is used to define multi-character collating elements. The syntax is:

```
"collating-element %s from \"%s\"\\n\",<collating-symbol>,<string>
```

The *<collating-symbol>* operand is a symbolic name, enclosed between angle brackets (< and >), and must not duplicate any symbolic name in the current charmap file (if any), or any other symbolic name defined in this collation definition. The string operand is a string of two or more characters that collates as an entity. A *<collating-element>* defined via this keyword is only recognised with the LC_COLLATE category.

Example:

```
collating-element <ch> from "<c><h>"
collating-element <e-acute> from "<acute><e>"
collating-element <ll> from "ll"
```

The collating-symbol Keyword

This keyword will be used to define symbols for use in collation sequence statements; that is, between the **order_start** and the **order_end** keywords. The syntax is:

```
"collating-symbol %s\\n\",<collating-symbol>
```

The *<collating-symbol>* is a symbolic name, enclosed between angle brackets (< and >), and must not duplicate any symbolic name in the current charmap file (if any), or any other symbolic name defined in this collation definition. A *<collating-symbol>* defined via this keyword is only recognised with the LC_COLLATE category.

Example:

```
collating-symbol <UPPER_CASE>
collating-symbol <HIGH>
```

The **collating-symbol** keyword defines a symbolic name that can be associated with a relative position in the character order sequence. While such a symbolic name does not represent any collating element, it can be used as a weight.

The order_start Keyword

The **order_start** keyword must precede collation order entries and also defines the number of weights for this collation sequence definition and other collation rules.

The syntax of the **order_start** keyword is:

```
"order_start %s;%s;...;%s\\n\",<sort-rules>,<sort-rules>
```

The operands to the **order_start** keyword are optional. If present, the operands define rules to be applied when strings are compared. The number of operands define how many weights each element is assigned; if no operands are present, one **forward** operand is assumed. If present, the first operand defines rules to be applied when comparing strings using the first (primary) weight; the second when comparing strings using the second weight, and so on. Operands are separated by semicolons (;). Each operand consists of one or more collation directives, separated by commas (.). If the number of operands exceeds the {COLL_WEIGHTS_MAX} limit, the utility will issue a warning message. The following directives will be supported:

forward Specifies that comparison operations for the weight level proceed from start of string towards the end of string.

2564 **backward** Specifies that comparison operations for the weight level proceed from end of
 2565 string towards the beginning of string.

2566 **position** Specifies that comparison operations for the weight level will consider the relative
 2567 position of elements in the strings not subject to **IGNORE**. The string containing
 2568 an element not subject to **IGNORE** after the fewest collating elements subject to
 2569 **IGNORE** from the start of the compare will collate first. If both strings contain a
 2570 character not subject to **IGNORE** in the same relative position, the collating values
 2571 assigned to the elements will determine the ordering. In case of equality,
 2572 subsequent characters not subject to **IGNORE** are considered in the same manner.

2573 The directives **forward** and **backward** are mutually exclusive.

2574 **Example:**

2575 order_start forward;backward

2576 If no operands are specified, a single **forward** operand is assumed.

2577 The character (and collating element) order is defined by the order in which characters and
 2578 elements are specified between the **order_start** and **order_end** keywords. This character order is
 2579 used in range expressions in regular expressions (see Chapter 7). Weights assigned to the
 2580 characters and elements define the collation sequence; in the absence of weights, the character
 2581 order is also the collation sequence.

2582 The **position** keyword provides the capability to consider, in a compare, the relative position of
 2583 characters not subject to **IGNORE**. As an example, consider the two strings “o-ring” and “or-
 2584 ing”. Assuming the hyphen is subject to **IGNORE** on the first pass, the two strings will compare
 2585 equal, and the position of the hyphen is immaterial. On second pass, all characters except the
 2586 hyphen are subject to **IGNORE**, and in the normal case the two strings would again compare
 2587 equal. By taking position into account, the first collates before the second.

2588 **Collation Order**

2589 The **order_start** keyword is followed by collating identifier entries. The syntax for the collating
 2590 element entries is:

2591 "%s %s;%s;...;%s\n", <collating-identifier>, <weight>, <weight>, ...

2592 Each *collating-identifier* consists of either a character (in any of the forms defined in Section 5.3 on
 2593 page 50), a <collating-element>, a <collating-symbol>, an ellipsis or the special symbol
 2594 **UNDEFINED**. The order in which collating elements are specified determines the character
 2595 order sequence, such that each collating element compares less than the elements following it.
 2596 The NUL character compares lower than any other character.

2597 A <collating-element> is used to specify multi-character collating elements, and indicates that the
 2598 character sequence specified via the <collating-element> is to be collated as a unit and in the
 2599 relative order specified by its place.

2600 A <collating-symbol> is used to define a position in the relative order for use in weights. No
 2601 weights are specified with a <collating-symbol>.

2602 The ellipsis symbol specifies that a sequence of characters will collate according to their encoded
 2603 character values. It is interpreted as indicating that all characters with a coded character set
 2604 value higher than the value of the character in the preceding line, and lower than the coded
 2605 character set value for the character in the following line, in the current coded character set, will
 2606 be placed in the character collation order between the previous and the following character in
 2607 ascending order according to their coded character set values. An initial ellipsis is interpreted as
 2608 if the preceding line specified the NUL character, and a trailing ellipsis as if the following line

specified the highest coded character set value in the current coded character set. An ellipsis is treated as invalid if the preceding or following lines do not specify characters in the current coded character set. The use of the ellipsis symbol ties the definition to a specific coded character set and may preclude the definition from being portable between implementations.

The symbol **UNDEFINED** is interpreted as including all coded character set values not specified explicitly or via the ellipsis symbol. Such characters are inserted in the character collation order at the point indicated by the symbol, and in ascending order according to their coded character set values. If no **UNDEFINED** symbol is specified, and the current coded character set contains characters not specified in this section, the utility will issue a warning message and place such characters at the end of the character collation order.

The optional operands for each collation-element are used to define the primary, secondary, or subsequent weights for the collating element. The first operand specifies the relative primary weight, the second the relative secondary weight, and so on. Two or more collation-elements can be assigned the same weight; they belong to the same *equivalence class* if they have the same primary weight. Collation behaves as if, for each weight level, elements subject to **IGNORE** are removed, unless the **position** collation directive is specified for the corresponding level with the **order_start** keyword. Then each successive pair of elements is compared according to the relative weights for the elements. If the two strings compare equal, the process is repeated for the next weight level, up to the limit {COLL_WEIGHTS_MAX}.

Weights are expressed as characters (in any of the forms specified in Section 5.3 on page 50), *<collating-symbol>s*, *<collating-element>s*, an ellipsis, or the special symbol **IGNORE**. A single character, a *<collating-symbol>* or a *<collating-element>* represent the relative position in the character collating sequence of the character or symbol, rather than the character or characters themselves. Thus, rather than assigning absolute values to weights, a particular weight is expressed using the relative order value assigned to a collating element based on its order in the character collation sequence.

One-to-many mapping is indicated by specifying two or more concatenated characters or symbolic names. For example, if the character *<eszet>* is given the string "*<s><s>*" as a weight, comparisons are performed as if all occurrences of the character *<eszet>* are replaced by *<s><s>* (assuming that *<s>* has the collating weight *<s>*). If it is necessary to define *<eszet>* and *<s><s>* as an equivalence class, then a collating element must be defined for the string *ss*.

All characters specified via an ellipsis will by default be assigned unique weights, equal to the relative order of characters. Characters specified via an explicit or implicit **UNDEFINED** special symbol will by default be assigned the same primary weight (that is, belong to the same equivalence class). An ellipsis symbol as a weight is interpreted to mean that each character in the sequence has unique weights, equal to the relative order of their character in the character collation sequence. The use of the ellipsis as a weight is treated as an error if the collating element is neither an ellipsis nor the special symbol **UNDEFINED**.

The special keyword **IGNORE** as a weight indicates that when strings are compared using the weights at the level where **IGNORE** is specified, the collating element is ignored; that is, as if the string did not contain the collating element. In regular expressions and pattern matching, all characters that are subject to **IGNORE** in their primary weight form an equivalence class.

An empty operand is interpreted as the collating element itself.

2652 For example, the order statement:

2653 <a> <a> ; <a>

2654 is equal to:

2655 <a>

2656 An ellipsis can be used as an operand if the collating element was an ellipsis, and is interpreted
2657 as the value of each character defined by the ellipsis.

2658 The collation order as defined in this section defines the interpretation of bracket expressions in
2659 regular expressions (see Section 7.3.5 on page 105).

2660 **Example:**

```

2661     order_start    forward;backward
2662     UNDEFINED     IGNORE; IGNORE
2663     <LOW>
2664     <space>       <LOW> ; <space>
2665     ...           <LOW> ; ...
2666     <a>            <a> ; <a>
2667     <a-acute>      <a> ; <a-acute>
2668     <a-grave>      <a> ; <a-grave>
2669     <A>            <a> ; <A>
2670     <A-acute>      <a> ; <A-acute>
2671     <A-grave>      <a> ; <A-grave>
2672     <ch>           <ch> ; <ch>
2673     <Ch>           <ch> ; <Ch>
2674     <s>            <s> ; <s>
2675     <eszet>        "<s><s>" ; "<eszet><eszet>"
2676     order_end

```

2677 This example is interpreted as follows:

- 2678 1. The **UNDEFINED** means that all characters not specified in this definition (explicitly or
2679 via the ellipsis) are ignored for collation purposes; for regular expression purposes they are
2680 ordered first.
- 2681 2. All characters between <space> and <a> have the same primary equivalence class and
2682 individual secondary weights based on their ordinal encoded values.
- 2683 3. All characters based on the upper- or lower-case character a belong to the same primary
2684 equivalence class.
- 2685 4. The multi-character collating element <ch> is represented by the collating symbol <ch>
2686 and belongs to the same primary equivalence class as the multi-character collating element
2687 <Ch>.

2688 **The order_end Keyword**

2689 The collating order entries must be terminated with an **order_end** keyword.

2690 The collation sequence definition of the POSIX locale follows; the code listing depicts the
2691 *localedef* input.

2692 LC_COLLATE
2693 # This is the POSIX locale definition for the LC_COLLATE category.
2694 # The order is the same as in the ASCII codeset.
2695 order_start forward
2696 <NUL>
2697 <SOH>
2698 <STX>
2699 <ETX>
2700 <EOT>
2701 <ENQ>
2702 <ACK>
2703 <alert>
2704 <backspace>
2705 <tab>
2706 <newline>
2707 <vertical-tab>
2708 <form-feed>
2709 <carriage-return>
2710 <SO>
2711 <SI>
2712 <DLE>
2713 <DC1>
2714 <DC2>
2715 <DC3>
2716 <DC4>
2717 <NAK>
2718 <SYN>
2719 <ETB>
2720 <CAN>
2721
2722 <SUB>
2723 <ESC>
2724 <IS4>
2725 <IS3>
2726 <IS2>
2727 <IS1>
2728 <space>
2729 <exclamation-mark>
2730 <quotation-mark>
2731 <number-sign>
2732 <dollar-sign>
2733 <percent-sign>
2734 <ampersand>
2735 <apostrophe>
2736 <left-parenthesis>
2737 <right-parenthesis>
2738 <asterisk>
2739 <plus-sign>
2740 <comma>
2741 <hyphen>

2742	<period>
2743	<slash>
2744	<zero>
2745	<one>
2746	<two>
2747	<three>
2748	<four>
2749	<five>
2750	<six>
2751	<seven>
2752	<eight>
2753	<nine>
2754	<colon>
2755	<semicolon>
2756	<less-than-sign>
2757	<equals-sign>
2758	<greater-than-sign>
2759	<question-mark>
2760	<commercial-at>
2761	<A>
2762	
2763	<C>
2764	<D>
2765	<E>
2766	<F>
2767	<G>
2768	<H>
2769	<I>
2770	<J>
2771	<K>
2772	<L>
2773	<M>
2774	<N>
2775	<O>
2776	<P>
2777	<Q>
2778	<R>
2779	<S>
2780	<T>
2781	<U>
2782	<V>
2783	<W>
2784	<X>
2785	<Y>
2786	<Z>
2787	<left-square-bracket>
2788	<backslash>
2789	<right-square-bracket>
2790	<circumflex>
2791	<underscore>
2792	<grave-accent>
2793	<a>

```

2794      <b>
2795      <c>
2796      <d>
2797      <e>
2798      <f>
2799      <g>
2800      <h>
2801      <i>
2802      <j>
2803      <k>
2804      <l>
2805      <m>
2806      <n>
2807      <o>
2808      <p>
2809      <q>
2810      <r>
2811      <s>
2812      <t>
2813      <u>
2814      <v>
2815      <w>
2816      <x>
2817      <y>
2818      <z>
2819      <left-curly-bracket>
2820      <vertical-line>
2821      <right-curly-bracket>
2822      <tilde>
2823      <DEL>
2824      order_end
2825      #
2826      END LC_COLLATE

```

2827 5.3.3 LC_MONETARY

2828 The LC_MONETARY category defines the rules and symbols that are used to format monetary
 2829 EX numeric information. This information is available through the *localeconv()* function and is used
 2830 by the *strfmon()* function.

2831 EX Some of the information is also available in an alternative form via the *nl_langinfo()* function
 2832 (see CRNCYSTR in <langinfo.h>).

2833 The following items are defined in this category of the locale. The item names are the keywords
 2834 recognised by the *localedef* utility when defining a locale. They are also similar to the member
 2835 names of the **lconv** structure defined in <locale.h>; see the **XSH** specification for the exact
 2836 symbols in the header. The *localeconv()* function returns {CHAR_MAX} for unspecified integer
 2837 items and the empty string ("") for unspecified or size zero string items.

2838 In a locale definition file, the operands are strings, formatted as indicated by the grammar in
 2839 Section 5.4 on page 82. For some keywords, the strings can contain only integers. Keywords
 2840 that are not provided, string values set to the empty string (""), or integer keywords set to -1, are
 2841 used to indicate that the value is not available in the locale.

2842	int_curr_symbol	The international currency symbol. The operand is a four-character string, with the first three characters containing the alphabetic international currency symbol in accordance with those specified in the ISO 4217:1987 standard. The fourth character is the character used to separate the international currency symbol from the monetary quantity.																		
2843																				
2844																				
2845																				
2846																				
2847	currency_symbol	The string used as the local currency symbol.																		
2848	mon_decimal_point	The operand is a string containing the symbol that is used as the decimal delimiter (radix character) in monetary formatted quantities. In contexts where standards (such as the ISO C standard) limit the mon_decimal_point to a single byte, the result of specifying a multi-byte operand is unspecified.																		
2849																				
2850																				
2851																				
2852																				
2853	mon_thousands_sep	The operand is a string containing the symbol that is used as a separator for groups of digits to the left of the decimal delimiter in formatted monetary quantities. In contexts where standards limit the mon_thousands_sep to a single byte, the result of specifying a multi-byte operand is unspecified.																		
2854																				
2855																				
2856																				
2857																				
2858	mon_grouping	Define the size of each group of digits in formatted monetary quantities. The operand is a sequence of integers separated by semicolons. Each integer specifies the number of digits in each group, with the initial integer defining the size of the group immediately preceding the decimal delimiter, and the following integers defining the preceding groups. If the last integer is not -1, then the size of the previous group (if any) will be repeatedly used for the remainder of the digits. If the last integer is -1, then no further grouping will be performed.																		
2859																				
2860																				
2861																				
2862																				
2863																				
2864																				
2865																				
2866		The following is an example of the interpretation of the mon_grouping keyword. Assuming that the value to be formatted is 123456789 and the mon_thousands_sep is ', then the following table shows the result. The third column shows the equivalent string in the ISO C standard that would be used by the <i>localeconv()</i> function to accommodate this grouping.																		
2867																				
2868																				
2869																				
2870																				
2871																				
2872																				
2873																				
2874																				
2875																				
2876																				
2877																				
2878																				
<table border="1"> <thead> <tr> <th>mon_grouping</th><th>Formatted Value</th><th>ISO C String</th></tr> </thead> <tbody> <tr> <td>3;-1</td><td>123456'789</td><td>"\3\177"</td></tr> <tr> <td>3</td><td>123'456'789</td><td>"\3"</td></tr> <tr> <td>3;2;-1</td><td>1234'56'789</td><td>"\3\2\177"</td></tr> <tr> <td>3;2</td><td>12'34'56'789</td><td>"\3\2"</td></tr> <tr> <td>-1</td><td>123456789</td><td>"\177"</td></tr> </tbody> </table>			mon_grouping	Formatted Value	ISO C String	3;-1	123456'789	"\3\177"	3	123'456'789	"\3"	3;2;-1	1234'56'789	"\3\2\177"	3;2	12'34'56'789	"\3\2"	-1	123456789	"\177"
mon_grouping	Formatted Value	ISO C String																		
3;-1	123456'789	"\3\177"																		
3	123'456'789	"\3"																		
3;2;-1	1234'56'789	"\3\2\177"																		
3;2	12'34'56'789	"\3\2"																		
-1	123456789	"\177"																		
2879		In these examples, the octal value of {CHAR_MAX} is 177.																		
2880	positive_sign	A string used to indicate a non-negative-valued formatted monetary quantity.																		
2881																				
2882	negative_sign	A string used to indicate a negative-valued formatted monetary quantity.																		
2883	int_frac_digits	An integer representing the number of fractional digits (those to the right of the decimal delimiter) to be written in a formatted monetary quantity using int_curr_symbol .																		
2884																				
2885																				

2886	frac_digits	An integer representing the number of fractional digits (those to the right of the decimal delimiter) to be written in a formatted monetary quantity using currency_symbol .
2887		
2888		
2889	p_cs_precedes	An integer set to 1 if the currency_symbol or int_curr_symbol precedes the value for a monetary quantity with a non-negative value, and set to 0 if the symbol succeeds the value.
2890		
2891		
2892	p_sep_by_space	An integer set to 0 if no space separates the currency_symbol or int_curr_symbol from the value for a monetary quantity with a non-negative value, set to 1 if a space separates the symbol from the value, and set to 2 if a space separates the symbol and the sign string, if adjacent.
2893		
2894		
2895		
2896	n_cs_precedes	An integer set to 1 if the currency_symbol or int_curr_symbol precedes the value for a monetary quantity with a negative value, and set to 0 if the symbol succeeds the value.
2897		
2898		
2899	n_sep_by_space	An integer set to 0 if no space separates the currency_symbol or int_curr_symbol from the value for a monetary quantity with a negative value, set to 1 if a space separates the symbol from the value, and set to 2 if a space separates the symbol and the sign string, if adjacent.
2900		
2901		
2902		
2903	p_sign_posn	An integer set to a value indicating the positioning of the positive_sign for a monetary quantity with a non-negative value. The following integer values are recognised for both p_sign_posn and n_sign_posn :
2904		
2905		
2906		0 Parentheses enclose the quantity and the currency_symbol or int_curr_symbol .
2907		
2908		1 The sign string precedes the quantity and the currency_symbol or int_curr_symbol .
2909		
2910		2 The sign string succeeds the quantity and the currency_symbol or int_curr_symbol .
2911		
2912		3 The sign string precedes the currency_symbol or int_curr_symbol .
2913		4 The sign string succeeds the currency_symbol or int_curr_symbol .
2914	n_sign_posn	An integer set to a value indicating the positioning of the negative_sign for a negative formatted monetary quantity.
2915		
2916	copy	Note: This is a <i>localedef</i> utility keyword, unavailable through <i>localeconv()</i> .
2917		
2918		Specify the name of an existing locale to be used as the definition of this category. If this keyword is specified, no other keyword can be specified.
2919		

The following table shows the result of various combinations:

		p_sep_by_space		
		2	1	0
p_cs_precedes = 1	p_sign_posn = 0	(\$1.25)	(\$ 1.25)	(\$1.25)
	p_sign_posn = 1	+ \$1.25	+\$ 1.25	+\$1.25
	p_sign_posn = 2	\$1.25 +	\$ 1.25+	\$1.25+
	p_sign_posn = 3	+ \$1.25	+\$ 1.25	+\$1.25
	p_sign_posn = 4	\$ +1.25	\$+ 1.25	\$+1.25
p_cs_precedes = 0	p_sign_posn = 0	(1.25 \$)	(1.25 \$)	(1.25\$)
	p_sign_posn = 1	+1.25 \$	+1.25 \$	+1.25\$
	p_sign_posn = 2	1.25\$ +	1.25 \$+	1.25\$+
	p_sign_posn = 3	1.25+ \$	1.25 +\$	1.25+\$
	p_sign_posn = 4	1.25\$ +	1.25 \$+	1.25\$+

The monetary formatting definitions for the POSIX locale follow; the code listing depicting the *localedef* input, the table representing the same information with the addition of *localeconv()* and *nl_langinfo()* formats. All values are unspecified in the POSIX locale.

```

LC_MONETARY
# This is the POSIX locale definition for
# the LC_MONETARY category.
#
int_curr_symbol      " "
currency_symbol      " "
mon_decimal_point    " "
mon_thousands_sep   " "
mon_grouping         -1
positive_sign        " "
negative_sign        " "
int_frac_digits      -1
p_cs_precedes        -1
p_sep_by_space       -1
n_cs_precedes        -1
n_sep_by_space       -1
p_sign_posn          -1
n_sign_posn          -1
#
END LC_MONETARY

```

Item	POSIX locale Value	langinfo Constant	localeconv() Value	localedef Value
currency_symbol	n/a	CRNCYSTR	""	""
frac_digits	n/a	-	CHAR_MAX	-1
int_curr_symbol	n/a	-	""	""
int_frac_digits	n/a	-	CHAR_MAX	-1
mon_decimal_point	n/a	-	""	""
mon_thousands_sep	n/a	-	""	""
mon_grouping	n/a	-	""	""
positive_sign	n/a	-	""	""
negative_sign	n/a	-	""	""
p_cs_precedes	n/a	CRNCYSTR	CHAR_MAX	-1
n_cs_precedes	n/a	CRNCYSTR	CHAR_MAX	-1
p_sep_by_space	n/a	-	CHAR_MAX	-1
n_sep_by_space	n/a	-	CHAR_MAX	-1
p_sign_posn	n/a	-	CHAR_MAX	-1
n_sign_posn	n/a	-	CHAR_MAX	-1

EX In the preceding table, the **langinfo Constant** column represents an X/Open extension. The entry **n/a** indicates that the value is not available in the POSIX locale.

5.3.4 LC_NUMERIC

EX The LC_NUMERIC category defines the rules and symbols that will be used to format non-monetary numeric information. This information is available through the *localeconv()* function. Some of the information is also available in an alternative form via the *nl_langinfo()* function.

The following items are defined in this category of the locale. The item names are the keywords recognised by the *localedef* utility when defining a locale. They are also similar to the member names of the *lconv* structure defined in `<locale.h>`; see the **XSH** specification for the exact symbols in the header. The *localeconv()* function returns {CHAR_MAX} for unspecified integer items and the empty string ("") for unspecified or size zero string items.

In a locale definition file, the operands are strings, formatted as indicated by the grammar in Section 5.4 on page 82. For some keywords, the strings only can contain integers. Keywords that are not provided, string values set to the empty string (""), or integer keywords set to -1, will be used to indicate that the value is not available in the locale. The following keywords are recognised:

decimal_point	The operand is a string containing the symbol that is used as the decimal delimiter (radix character) in numeric, non-monetary formatted quantities. This keyword cannot be omitted and cannot be set to the empty string. In contexts where standards limit the decimal_point to a single byte, the result of specifying a multi-byte operand is unspecified.
thousands_sep	The operand is a string containing the symbol that is used as a separator for groups of digits to the left of the decimal delimiter in numeric, non-monetary formatted monetary quantities. In contexts where standards limit the thousands_sep to a single byte, the result of specifying a multi-byte operand is unspecified.
grouping	Define the size of each group of digits in formatted non-monetary quantities. The operand is a sequence of integers separated by semicolons. Each integer specifies the number of digits in each group, with the initial integer defining the size of the group immediately preceding the decimal delimiter, and the

3005 following integers defining the preceding groups. If the last integer is not -1,
 3006 then the size of the previous group (if any) will be repeatedly used for the
 3007 remainder of the digits. If the last integer is -1, then no further grouping will
 3008 be performed.

3009 **copy** **Note:** This is a *localedef* utility keyword, unavailable through *localeconv()*.

3010 Specify the name of an existing locale to be used as the definition of this
 3011 category. If this keyword is specified, no other keyword can be specified.

3012 The non-monetary numeric formatting definitions for the POSIX locale follow; the code listing
 3013 depicting the *localedef* input, the table representing the same information with the addition of
 3014 EX *localeconv()* values and *nl_langinfo()* constants.

```
3015 LC_NUMERIC
3016 # This is the POSIX locale definition for
3017 # the LC_NUMERIC category.
3018 #
3019 decimal_point      "<period>"
3020 thousands_sep      " "
3021 grouping           -1
3022 #
3023 END LC_NUMERIC
```

Item	POSIX locale Value	langinfo Constant	localeconv() Value	localedef Value
decimal_point	"."	RADIXCHAR	"."	"."
thousands_sep	n/a	THOUSEP	""	""
grouping	n/a	-	""	-1

3030 EX In the preceding table, the **langinfo Constant** column represents an X/Open extension. The
 3031 entry **n/a** indicates that the value is not available in the POSIX locale.

3032 5.3.5 LC_TIME

3033 The LC_TIME category defines the interpretation of the field descriptors supported by the *date*
 3034 EX utility and affects the behaviour of the *strptime()*, *wcsftime()*, *strptime()* and *nl_langinfo()*
 3035 functions. Because the interfaces for C-language access and locale definition differ significantly,
 3036 they are described separately.

3037 LC_TIME Locale Definition

3038 For locale definition, the following mandatory keywords are recognised:

3039 **abday** Define the abbreviated weekday names, corresponding to the %a field
 3040 descriptor (conversion specification in the *strptime()*, *wcsftime()* and *strptime()*
 3041 functions). The operand consists of seven semicolon-separated strings, each
 3042 surrounded by double-quotes. The first string is the abbreviated name of the
 3043 day corresponding to Sunday, the second the abbreviated name of the day
 3044 corresponding to Monday, and so on.

3045	day	Define the full weekday names, corresponding to the %A field descriptor. The operand consists of seven semicolon-separated strings, each surrounded by double-quotes. The first string is the full name of the day corresponding to Sunday, the second the full name of the day corresponding to Monday, and so on.
3046		
3047		
3048		
3049		
3050	abmon	Define the abbreviated month names, corresponding to the %b field descriptor. The operand consists of twelve semicolon-separated strings, each surrounded by double-quotes. The first string is the abbreviated name of the first month of the year (January), the second the abbreviated name of the second month, and so on.
3051		
3052		
3053		
3054		
3055	mon	Define the full month names, corresponding to the %B field descriptor. The operand consists of twelve semicolon-separated strings, each surrounded by double-quotes. The first string is the full name of the first month of the year (January), the second the full name of the second month, and so on.
3056		
3057		
3058		
3059	d_t_fmt	Define the appropriate date and time representation, corresponding to the %c field descriptor. The operand consists of a string, and can contain any combination of characters and field descriptors. In addition, the string can contain escape sequences defined in the table in Table 3-1 on page 40 (\, \a, \b, \f, \n, \r, \t, \v).
3060		
3061		
3062		
3063		
3064	d_fmt	Define the appropriate date representation, corresponding to the %x field descriptor. The operand consists of a string, and can contain any combination of characters and field descriptors. In addition, the string can contain escape sequences defined in the table in Table 3-1 on page 40.
3065		
3066		
3067		
3068	t_fmt	Define the appropriate time representation, corresponding to the %X field descriptor. The operand consists of a string, and can contain any combination of characters and field descriptors. In addition, the string can contain escape sequences defined in the table in Table 3-1 on page 40.
3069		
3070		
3071		
3072	am_pm	Define the appropriate representation of the <i>ante meridiem</i> and <i>post meridiem</i> strings, corresponding to the %p field descriptor. The operand consists of two strings, separated by a semicolon, each surrounded by double-quotes. The first string represents the <i>ante meridiem</i> designation, the last string the <i>post meridiem</i> designation.
3073		
3074		
3075		
3076		
3077	t_fmt_ampm	Define the appropriate time representation in the 12-hour clock format with am_pm , corresponding to the %r field descriptor. The operand consists of a string and can contain any combination of characters and field descriptors. If the string is empty, the 12-hour format is not supported in the locale.
3078		
3079		
3080		
3081	EX era	Define how years are counted and displayed for each era in a locale. The operand consists of semicolon-separated strings. Each string is an era description segment with the format:
3082		
3083		
3084		<i>direction:offset:start_date:end_date:era_name:era_format</i>
3085		according to the definitions below. There can be as many era description segments as are necessary to describe the different eras.
3086		
3087		Note: The start of an era might not be the earliest point in the era it may be the latest. For example, the Christian era BC starts on the day before January 1, AD 1, and increases with earlier time.
3088		
3089		

3090		<i>direction</i>	Either a + or a – character. The + character indicates that years closer to the <i>start_date</i> have lower numbers than those closer to the <i>end_date</i> . The – character indicates that years closer to the <i>start_date</i> have higher numbers than those closer to the <i>end_date</i> .
3091			
3092			
3093			
3094		<i>offset</i>	The number of the year closest to the <i>start_date</i> in the era, corresponding to the %Ey field descriptor.
3095			
3096		<i>start_date</i>	A date in the form <i>yyyy/mm/dd</i> , where <i>yyyy</i> , <i>mm</i> and <i>dd</i> are the year, month and day numbers respectively of the start of the era. Years prior to AD 1 are represented as negative numbers.
3097			
3098			
3099		<i>end_date</i>	The ending date of the era, in the same format as the <i>start_date</i> , or one of the two special values –* or +*. The value –* indicates that the ending date is the beginning of time. The value +* indicates that the ending date is the end of time.
3100			
3101			
3102			
3103		<i>era_name</i>	A string representing the name of the era, corresponding to the %EC field descriptor.
3104			
3105		<i>era_format</i>	A string for formatting the year in the era, corresponding to the %EY field descriptor.
3106			
3107		era_d_fmt	Define the format of the date in alternative era notation, corresponding to the %Ex field descriptor.
3108			
3109		era_t_fmt	Define the locale's appropriate alternative time format, corresponding to the %EX field descriptor.
3110			
3111		era_d_t_fmt	Define the locale's appropriate alternative date and time format, corresponding to the %Ec field descriptor.
3112			
3113		alt_digits	Define alternative symbols for digits, corresponding to the %O field descriptor modifier. The operand consists of semicolon-separated strings, each surrounded by double-quotes. The first string is the alternative symbol corresponding with zero, the second string the symbol corresponding with one, and so on. Up to 100 alternative symbol strings can be specified. The %O modifier indicates that the string corresponding to the value specified via the field descriptor will be used instead of the value.
3114			
3115			
3116			
3117			
3118			
3119			
3120		copy	Specify the name of an existing locale to be used as the definition of this category. If this keyword is specified, no other keyword can be specified.
3121			
3122		LC_TIME C-language Access	
3123	EX	The following information can be accessed. These correspond to constants defined in <langinfo.h> and used as arguments to the <i>nl_langinfo()</i> function.	
3124			
3125		ABDAY_x	The abbreviated weekday names (for example Sun), where <i>x</i> is a number from 1 to 7.
3126			
3127		DAY_x	The full weekday names (for example Sunday), where <i>x</i> is a number from 1 to 7.
3128			
3129		ABMON_x	The abbreviated month names (for example Jan), where <i>x</i> is a number from 1 to 12.
3130			
3131		MON_x	The full month names (for example January), where <i>x</i> is a number from 1 to 12.
3132			

3133	D_T_FMT	The appropriate date and time representation.
3134	D_FMT	The appropriate date representation.
3135	T_FMT	The appropriate time representation.
3136	AM_STR	The appropriate ante-meridiem affix.
3137	PM_STR	The appropriate post-meridiem affix.
3138	T_FMT_AMPM	The appropriate time representation in the 12-hour clock format with AM_STR and PM_STR.
3139		
3140	ERA	The era description segments, which describe how years are counted and displayed for each era in a locale. Each era description segment has the format:
3141		
3142		
3143		<i>direction:offset:start_date:end_date:era_name:era_format</i>
3144		according to the definitions below. There will be as many era description segments as are necessary to describe the different eras. Era description segments are separated by semicolons.
3145		
3146		
3147		Note: The start of an era might not be the earliest point in the era it may be the latest. For example, the Christian era BC starts on the day before January 1, AD 1, and increases with earlier time.
3148		
3149		
3150		<i>direction</i> Either a + or a – character. The + character indicates that years closer to the <i>start_date</i> have lower numbers than those closer to the <i>end_date</i> . The – character indicates that years closer to the <i>start_date</i> have higher numbers than those closer to the <i>end_date</i> .
3151		
3152		
3153		
3154		<i>offset</i> The number of the year closest to the <i>start_date</i> in the era.
3155		<i>start_date</i> A date in the form <i>yyyy/mm/dd</i> , where <i>yyyy</i> , <i>mm</i> and <i>dd</i> are the year, month and day numbers respectively of the start of the era. Years prior to AD 1 are represented as negative numbers.
3156		
3157		
3158		<i>end_date</i> The ending date of the era, in the same format as the <i>start_date</i> , or one of the two special values –* or +*. The value –* indicates that the ending date is the beginning of time. The value +* indicates that the ending date is the end of time.
3159		
3160		
3161		
3162		<i>era_name</i> The era, corresponding to the %EC conversion specification.
3163		<i>era_format</i> The format of the year in the era, corresponding to the %EY conversion specification.
3164		
3165	ERA_D_FMT	The era date format.
3166	EX ERA_T_FMT	The locale's appropriate alternative time format, corresponding to the %EX field descriptor.
3167		
3168	ERA_D_T_FMT	The locale's appropriate alternative date and time format, corresponding to the %Ec field descriptor.
3169		
3170	ALT_DIGITS	The alternative symbols for digits, corresponding to the %O conversion specification modifier. The value consists of semicolon-separated symbols. The first is the alternative symbol corresponding to zero, the second is the symbol corresponding to one, and so on. Up to 100 alternative symbols may be specified.
3171		
3172		
3173		
3174		

The following table displays the correspondence between the items described above and the conversion specifiers used by the *date* utility and the *strftime()*, *wcsftime()* and *strptime()* functions.

localedef Keyword	langinfo Constant	Conversion Specifier
abday	ABDAY_ <i>x</i>	%a
day	DAY_ <i>x</i>	%A
abmon	ABMON_ <i>x</i>	%b
mon	MON	%B
d_t_fmt	D_T_FMT	%c
d_fmt	D_FMT	%x
t_fmt	T_FMT	%X
am_pm	AM_STR	%p
am_pm	PM_STR	%p
t_fmt_ampm	T_FMT_AMPM	%r
era	ERA	%EC, %Ey, %EY
era_d_fmt	ERA_D_FMT	%Ex
era_t_fmt	ERA_T_FMT	%EX
era_d_t_fmt	ERA_D_T_FMT	%Ec
alt_digits	ALT_DIGITS	%O

In the preceding table, the **langinfo Constant** column represents an X/Open extension.

LC_TIME General Information

Although certain of the field descriptors in the POSIX locale (such as the name of the month) are shown with initial capital letters, this need not be the case in other locales. Programs using these fields may need to adjust the capitalisation if the output is going to be used at the beginning of a sentence.

The LC_TIME descriptions of **abday**, **day**, **mon** and **abmon** imply a Gregorian style calendar (7-day weeks, 12-month years, leap years, and so forth). Formatting time strings for other types of calendars is outside the scope of this specification set.

As specified under *date* in the Locale Definition and *strftime()*, in the **XSH** specification, the field descriptors corresponding to the optional keywords consist of a modifier followed by a traditional field descriptor (for instance %Ex). If the optional keywords are not supported by the implementation or are unspecified for the current locale, these field descriptors are treated as the traditional field descriptor. For instance, assume the following keywords:

```
alt_digits    "0th"; "1st"; "2nd"; "3rd"; "4th"; "5th"; \
              "6th"; "7th"; "8th"; "9th"; "10th"
```

```
d_fmt        "The %Od day of %B in %Y"
```

On 7/4/1776, the %x field descriptor would result in “The 4th day of July in 1776”, while 7/14/1789 would come out as “The 14 day of July in 1789”. It can be noted that the above example is for illustrative purposes only; the %O modifier is primarily intended to provide for Kanji or Hindi digits in *date* formats.

3216 EX The following is an example for Japan that supports the current plus last three Emperors and
 3217 reverts to Western style numbering for years prior to the Meiji era. The example also allows for
 3218 the custom of using a special name for the first year of an era instead of using 1. (The examples
 3219 substitute romaji where kanji should be used.)

3220 era_d_fmt "%EY%mgatsu%dnichi (%a)"

```
3221 era      "+:2:1990/01/01:+*:Heisei:%EC%Eynen";\
3222          "+:1:1989/01/08:1989/12/31:Heisei:%ECgannen";\
3223          "+:2:1927/01/01:1989/01/07:Shouwa:%EC%Eynen";\
3224          "+:1:1926/12/25:1926/12/31:Shouwa:%ECgannen";\
3225          "+:2:1913/01/01:1926/12/24:Taishou:%EC%Eynen";\
3226          "+:1:1912/07/30:1912/12/31:Taishou:%ECgannen";\
3227          "+:2:1869/01/01:1912/07/29:Meiji:%EC%Eynen";\
3228          "+:1:1868/09/08:1868/12/31:Meiji:%ECgannen";\
3229          "-:1868:1868/09/07:-*:%Ey"
```

3230 Assuming that the current date is September 21, 1991, a request to *date* or *strftime()* would yield
 3231 the following results:

```
3232 %Ec - Heisei3nen9gatsu21nichi (Sat) 14:39:26
3233 %EC - Heisei
3234 %Ex - Heisei3nen9gatsu21nichi (Sat)
3235 %Ey - 3
3236 %EY - Heisei3nen
```

3237 Example era definitions for the Republic of China:

```
3238 era      "+:2:1913/01/01:+*:ChungHwaMingGuo:%EC%EyNen";\
3239          "+:1:1912/1/1:1912/12/31:ChungHwaMingGuo:%ECYuenNen";\
3240          "+:1:1911/12/31:-*:MingChien:%EC%EyNen"
```

3241 Example definitions for the Christian Era:

```
3242 era      "+:0:0000/01/01:+*:AD:%EC %Ey";\
3243          "+:1:-0001/12/31:-*:BC:%Ey %EC"
```

3244 The LC_TIME category definition of the POSIX locale follows; the code listing depicts the
 3245 EX *localedef* input; the table depicts the *langinfo* items defined in this category.

```
3246 LC_TIME
3247 # This is the POSIX locale definition for
3248 # the LC_TIME category.
3249 #
3250 # Abbreviated weekday names (%a)
3251 abday      "<S><u><n>" "<M><o><n>" "<T><u><e>" "<W><e><d>" \
3252            "<T><h><u>" "<F><r><i>" "<S><a><t>"
3253 #
3254 # Full weekday names (%A)
3255 day        "<S><u><n><d><a><y>" "<M><o><n><d><a><y>" \
3256            "<T><u><e><s><d><a><y>" "<W><e><d><n><e><s><d><a><y>" \
3257            "<T><h><u><r><s><d><a><y>" "<F><r><i><d><a><y>" \
3258            "<S><a><t><u><r><d><a><y>"
3259 #
```

```

3260      # Abbreviated month names (%b)
3261      abmon      "<J><a><n>" ; "<F><e><b>" ; "<M><a><r>" ; \
3262                "<A><p><r>" ; "<M><a><y>" ; "<J><u><n>" ; \
3263                "<J><u><l>" ; "<A><u><g>" ; "<S><e><p>" ; \
3264                "<O><c><t>" ; "<N><o><v>" ; "<D><e><c>"
3265      #
3266      # Full month names (%B)
3267      mon        "<J><a><n><u><a><r><y>" ; "<F><e><b><r><u><a><r><y>" ; \
3268                "<M><a><r><c><h>" ; "<A><p><r><i><l>" ; \
3269                "<M><a><y>" ; "<J><u><n><e>" ; \
3270                "<J><u><l><y>" ; "<A><u><g><u><s><t>" ; \
3271                "<S><e><p><t><e><m><b><e><r>" ; "<O><c><t><o><b><e><r>" ; \
3272                "<N><o><v><e><m><b><e><r>" ; "<D><e><c><e><m><b><e><r>"
3273      #
3274      # Equivalent of AM/PM (%p)      "AM" ; "PM"
3275      am_pm      "<A><M>" ; "<P><M>"
3276      #
3277      # Appropriate date and time representation (%c)
3278      #      "%a %b %e %H:%M:%S %Y"
3279      d_t_fmt     "<percent-sign><a><space><percent-sign><b>\
3280                <space><percent-sign><e><space><percent-sign><H>\
3281                <colon><percent-sign><M><colon><percent-sign><S>\
3282                <space><percent-sign><Y>"
3283      #
3284      # Appropriate date representation (%x)      "%m/%d/%y"
3285      d_fmt       "<percent-sign><m><slash><percent-sign><d>\
3286                <slash><percent-sign><y>"
3287      #
3288      # Appropriate time representation (%X)      "%H:%M:%S"
3289      t_fmt       "<percent-sign><H><colon><percent-sign><M>\
3290                <colon><percent-sign><S>"
3291      #
3292      # Appropriate 12-hour time representation (%r) "%I:%M:%S %p"
3293      t_fmt_ampm  "<percent-sign><I><colon><percent-sign><M><colon>\
3294                <percent-sign><S> <percent_sign><p>"
3295      #
3296      END LC_TIME

```

Item	POSIX Locale Value	Item	POSIX Locale Value
D_T_FMT	"%a %b %e %H:%M:%S %Y"	MON_3	"March"
D_FMT	"%m/%d/%y"	MON_4	"April"
T_FMT	"%H:%M:%S"	MON_5	"May"
AM_STR	"AM"	MON_6	"June"
PM_STR	"PM"	MON_7	"July"
T_FMT_AMP	"%I:%M:%S %p"	MON_8	"August"
DAY_1	"Sunday"	MON_9	"September"
DAY_2	"Monday"	MON_10	"October"
DAY_3	"Tuesday"	MON_11	"November"
DAY_4	"Wednesday"	MON_12	"December"
DAY_5	"Thursday"	ABMON_1	"Jan"
DAY_6	"Friday"	ABMON_2	"Feb"
DAY_7	"Saturday"	ABMON_3	"Mar"
ABDAY_1	"Sun"	ABMON_4	"Apr"
ABDAY_2	"Mon"	ABMON_5	"May"
ABDAY_3	"Tue"	ABMON_6	"Jun"
ABDAY_4	"Wed"	ABMON_7	"Jul"
ABDAY_5	"Thu"	ABMON_8	"Aug"
ABDAY_6	"Fri"	ABMON_9	"Sep"
ABDAY_7	"Sat"	ABMON_10	"Oct"
MON_1	"January"	ABMON_11	"Nov"
MON_2	"February"	ABMON_12	"Dec"

5.3.6 LC_MESSAGES

The LC_MESSAGES category defines the format and values for affirmative and negative responses.

The message catalogue used by the standard utilities and selected by the *catopen()* function is determined by the setting of *NLSPATH*; see Chapter 6 on page 93. The LC_MESSAGES category can be specified as part of an *NLSPATH* substitution field.

The following keywords are recognised as part of the locale definition file. The *nl_langinfo()* function accepts upper-case versions of the first four keywords.

yesexpr The operand consists of an extended regular expression (see Section 7.4 on page 109) that describes the acceptable affirmative response to a question expecting an affirmative or negative response.

noexpr The operand consists of an extended regular expression that describes the acceptable negative response to a question expecting an affirmative or negative response.

yesstr (LEGACY)
The operand consists of a fixed string (not a regular expression) that can be used by an application for composition of a message that lists an acceptable affirmative response, such as in a prompt.

nostr (LEGACY)
The operand consists of a fixed string that can be used by an application for composition of a message that lists an acceptable negative response.

copy Specify the name of an existing locale to be used as the definition of this category. If this keyword is specified, no other keyword can be specified.

3344 Note that the **yesstr** and **nostr** values have different uses from those in Issue 3.

3345 The format and values for affirmative and negative responses of the POSIX locale follow; the

3346 code listing depicting the *localedef* input, the table representing the same information with the

3347 EX addition of *nl_langinfo()* constants.

3348 LC_MESSAGES

3349 # This is the POSIX locale definition for

3350 # the LC_MESSAGES category.

3351 #

3352 yesexpr "<circumflex><left-square-bracket><y><Y><right-square-bracket>"

3353 #

3354 noexpr "<circumflex><left-square-bracket><n><N><right-square-bracket>"

3355 #

3356 EX yesstr "yes"

3357 nostr "no"

3358 END LC_MESSAGES

3359

3360

localedef Keyword	langinfo Constant	POSIX Locale Value
yesexpr	YESEXPR	"^[yY]"
noexpr	NOEXPR	"^[nN]"
yesstr	YESSTR	"yes" (LEGACY)
nostr	NOSTR	"no" (LEGACY)

3363 EX

3364 EX

3365 LC_MESSAGES Application Usage

3366 EX The **yesstr** and **nostr** locale keywords and the YESSTR and NOSTR *langinfo* items formerly were

3367 used to match user affirmative and negative responses. In this issue, the **yesexpr**, **noexpr**,

3368 YESEXPR and NOEXPR extended regular expressions have replaced them. However, they have

3369 been retained for backward compatibility to allow an application to include a sample desired

3370 response in a prompting message. They are marked **LEGACY**. Applications should use the

3371 general locale-based messaging facilities (see the **Internationalisation Guide**) to issue such

3372 prompting messages.

3373 5.4 Locale Definition Grammar

3374 The grammar and lexical conventions in this section together describe the syntax for the locale
 3375 definition source. The general conventions for this style of grammar are described in the XCU
 3376 specification, **Section 1.8, Grammar Conventions**. The grammar takes precedence over the text.

3377 5.4.1 Locale Lexical Conventions

3378 The lexical conventions for the locale definition grammar are described in this section.

3379 The following tokens are processed (in addition to those string constants shown in the
 3380 grammar):

3381	LOC_NAME	A string of characters representing the name of a locale.
3382	CHAR	Any single character.
3383	NUMBER	A decimal number, represented by one or more decimal digits.
3384	COLLSYMBOL	A symbolic name, enclosed between angle brackets. The string cannot duplicate any charmap symbol defined in the current charmap (if any), or a COLLELEMENT symbol.
3385		
3386		
3387	COLLELEMENT	A symbolic name, enclosed between angle brackets, which cannot duplicate either any charmap symbol or a COLLSYMBOL symbol.
3388		
3389 EX	CHARCLASS	A string of alphanumeric characters from the portable character set, the first of which is not a digit, consisting of at least one and at most {CHARCLASS_NAME_MAX} bytes, and optionally surrounded by double-quotes.
3390		
3391		
3392		
3393	CHARSYMBOL	A symbolic name, enclosed between angle brackets, from the current charmap (if any).
3394		
3395	OCTAL_CHAR	One or more octal representations of the encoding of each byte in a single character. The octal representation consists of an escape character (normally a backslash) followed by two or more octal digits.
3396		
3397		
3398		
3399	HEX_CHAR	One or more hexadecimal representations of the encoding of each byte in a single character. The hexadecimal representation consists of an escape character followed by the constant x and two or more hexadecimal digits.
3400		
3401		
3402		
3403	DECIMAL_CHAR	One or more decimal representations of the encoding of each byte in a single character. The decimal representation consists of an escape character followed by a character d and two or more decimal digits.
3404		
3405		
3406	ELLIPSIS	The string ...
3407	EXTENDED_REG_EXP	An extended regular expression as defined in the grammar in Section 7.5 on page 112.
3408		
3409	EOL	The line termination character newline.

3410 **5.4.2 Locale Grammar**

3411 This section presents the grammar for the locale definition.

```

3412 %token          LOC_NAME
3413 %token          CHAR
3414 %token          NUMBER
3415 %token          COLLSYMBOL COLLELEMENT
3416 %token          CHARSYMBOL OCTAL_CHAR HEX_CHAR DECIMAL_CHAR
3417 %token          ELLIPSIS
3418 %token          EXTENDED_REG_EXP
3419 %token          EOL
3420 %start          locale_definition
3421 %%
3422 locale_definition : global_statements locale_categories
3423                  | locale_categories
3424                  ;
3425 global_statements : global_statements symbol_redefine
3426                  | symbol_redefine
3427                  ;
3428 symbol_redefine   : 'escape_char' CHAR EOL
3429                  | 'comment_char' CHAR EOL
3430                  ;
3431 locale_categories : locale_categories locale_category
3432                  | locale_category
3433                  ;
3434 locale_category   : lc_ctype | lc_collate | lc_messages
3435                  | lc_monetary | lc_numeric | lc_time
3436                  ;
3437 /* The following grammar rules are common to all categories */
3438 char_list         : char_list char_symbol
3439                  | char_symbol
3440                  ;
3441 char_symbol       : CHAR | CHARSYMBOL
3442                  | OCTAL_CHAR | HEX_CHAR | DECIMAL_CHAR
3443                  ;
3444 elem_list        : elem_list char_symbol
3445                  | elem_list COLLSYMBOL
3446                  | elem_list COLLELEMENT
3447                  | char_symbol
3448                  | COLLSYMBOL
3449                  | COLLELEMENT
3450                  ;
3451 symb_list        : symb_list COLLSYMBOL
3452                  | COLLSYMBOL
3453                  ;

```

```

3454     locale_name      : LOC_NAME
3455                       | ' ' LOC_NAME ' '
3456                       ;

3457     /* The following is the LC_CTYPE category grammar */

3458     lc_ctype           : ctype_hdr ctype_keywords      ctype_tlr
3459                       | ctype_hdr 'copy' locale_name EOL ctype_tlr
3460                       ;

3461     ctype_hdr          : 'LC_CTYPE' EOL
3462                       ;

3463     ctype_keywords     : ctype_keywords ctype_keyword
3464                       | ctype_keyword
3465                       ;

3466     ctype_keyword      : charclass_keyword charclass_list EOL
3467                       | charconv_keyword charconv_list EOL
3468 EX   | 'charclass' charclass_namelist EOL
3469                       ;
3470     charclass_namelist : charclass_namelist ' ;' CHARCLASS
3471                       | CHARCLASS
3472                       ;

3473     charclass_keyword  : 'upper' | 'lower' | 'alpha' | 'digit'
3474                       | 'punct' | 'xdigit' | 'space' | 'print'
3475                       | 'graph' | 'blank' | 'cntrl'
3476 EX   | CHARCLASS
3477                       ;

3478     charclass_list     : charclass_list ' ;' char_symbol
3479                       | charclass_list ' ;' ELLIPSIS ';' char_symbol
3480                       | char_symbol
3481                       ;

3482     charconv_keyword   : 'toupper'
3483                       | 'tolower'
3484                       ;

3485     charconv_list      : charconv_list ' ;' charconv_entry
3486                       | charconv_entry
3487                       ;

3488     charconv_entry     : '(' char_symbol ',' char_symbol ')'
3489                       ;

3490     ctype_tlr          : 'END' 'LC_CTYPE' EOL
3491                       ;

3492     /* The following is the LC_COLLATE category grammar */

3493     lc_collate         : collate_hdr collate_keywords   collate_tlr
3494                       | collate_hdr 'copy' locale_name EOL collate_tlr
3495                       ;

3496     collate_hdr        : 'LC_COLLATE' EOL
3497                       ;

```

3498	collate_keywords	:	order_statements	
3499			opt_statements order_statements	
3500		;		
3501	opt_statements	:	opt_statements collating_symbols	
3502			opt_statements collating_elements	
3503			collating_symbols	
3504			collating_elements	
3505		;		
3506	collating_symbols	:	'collating-symbol' COLLSYMBOL EOL	
3507		;		
3508	collating_elements	:	'collating-element' COLLELEMENT	
3509			'from' ''' elem_list ''' EOL	
3510		;		
3511	order_statements	:	order_start collation_order order_end	
3512		;		
3513	order_start	:	'order_start' EOL	
3514			'order_start' order_opts EOL	
3515		;		
3516	order_opts	:	order_opts ' ;' order_opt	
3517			order_opt	
3518		;		
3519	order_opt	:	order_opt ',' opt_word	
3520			opt_word	
3521		;		
3522	opt_word	:	'forward' 'backward' 'position'	
3523		;		
3524	collation_order	:	collation_order collation_entry	
3525			collation_entry	
3526		;		
3527	collation_entry	:	COLLSYMBOL EOL	
3528			collation_element weight_list EOL	
3529			collation_element EOL	
3530		;		
3531	collation_element	:	char_symbol	
3532			COLLELEMENT	
3533			ELLIPSIS	
3534			'UNDEFINED'	
3535		;		
3536	weight_list	:	weight_list ' ;' weight_symbol	
3537			weight_list ' ;'	
3538			weight_symbol	
3539		;		

```

3540     weight_symbol      : /* empty */
3541                          | char_symbol
3542                          | COLLSYMBOL
3543                          | ''' elem_list '''
3544                          | ''' symb_list '''
3545                          | ELLIPSIS
3546                          | 'IGNORE'
3547                          ;

3548     order_end           : 'order_end' EOL
3549                          ;

3550     collate_tlr         : 'END' 'LC_COLLATE' EOL
3551                          ;

3552     /* The following is the LC_MESSAGES category grammar */

3553     lc_messages         : messages_hdr messages_keywords      messages_tlr
3554                          | messages_hdr 'copy' locale_name EOL messages_tlr
3555                          ;

3556     messages_hdr        : 'LC_MESSAGES' EOL
3557                          ;

3558     messages_keywords    : messages_keywords messages_keyword
3559                          | messages_keyword
3560                          ;

3561     messages_keyword     : 'yesexpr' ''' EXTENDED_REG_EXP ''' EOL
3562                          | 'noexpr'  ''' EXTENDED_REG_EXP ''' EOL
3563                          | 'yesstr'  ''' char_list ''' EOL
3564                          | 'nostr'   ''' char_list ''' EOL
3565                          ;

3566     messages_tlr        : 'END' 'LC_MESSAGES' EOL
3567                          ;

3568     /* The following is the LC_MONETARY category grammar */

3569     lc_monetary          : monetary_hdr monetary_keywords      monetary_tlr
3570                          | monetary_hdr 'copy' locale_name EOL monetary_tlr
3571                          ;

3572     monetary_hdr         : 'LC_MONETARY' EOL
3573                          ;

3574     monetary_keywords    : monetary_keywords monetary_keyword
3575                          | monetary_keyword
3576                          ;

3577     monetary_keyword     : mon_keyword_string mon_string EOL
3578                          | mon_keyword_char NUMBER EOL
3579                          | mon_keyword_char '-1' EOL
3580                          | mon_keyword_grouping mon_group_list EOL
3581                          ;

```

```

3582     mon_keyword_string : 'int_curr_symbol' | 'currency_symbol' |
3583                          | 'mon_decimal_point' | 'mon_thousands_sep'
3584                          | 'positive_sign' | 'negative_sign'
3585                          ;
3586     mon_string          : ''' char_list '''
3587                          | '""'
3588                          ;
3589     mon_keyword_char    : 'int_frac_digits' | 'frac_digits'
3590                          | 'p_cs_precedes' | 'p_sep_by_space'
3591                          | 'n_cs_precedes' | 'n_sep_by_space'
3592                          | 'p_sign_posn' | 'n_sign_posn'
3593                          ;
3594     mon_keyword_grouping : 'mon_grouping'
3595                          ;
3596     mon_group_list      : NUMBER
3597                          | mon_group_list ' ;' NUMBER
3598                          ;
3599     monetary_tlr       : 'END' 'LC_MONETARY' EOL
3600                          ;
3601     /* The following is the LC_NUMERIC category grammar */
3602     lc_numeric          : numeric_hdr numeric_keywords numeric_tlr
3603                          | numeric_hdr 'copy' locale_name EOL numeric_tlr
3604                          ;
3605     numeric_hdr         : 'LC_NUMERIC' EOL
3606                          ;
3607     numeric_keywords    : numeric_keywords numeric_keyword
3608                          | numeric_keyword
3609                          ;
3610     numeric_keyword     : num_keyword_string num_string EOL
3611                          | num_keyword_grouping num_group_list EOL
3612                          ;
3613     num_keyword_string  : 'decimal_point'
3614                          | 'thousands_sep'
3615                          ;
3616     num_string          : ''' char_list '''
3617                          | '""'
3618                          ;
3619     num_keyword_grouping : 'grouping'
3620                          ;
3621     num_group_list      : NUMBER
3622                          | num_group_list ' ;' NUMBER
3623                          ;

```

3624	numeric_tlr	:	'END' 'LC_NUMERIC' EOL	
3625		;		
3626	/* The following is the LC_TIME category grammar */			
3627	lc_time	:	time_hdr time_keywords time_tlr	
3628			time_hdr 'copy' locale_name EOL time_tlr	
3629		;		
3630	time_hdr	:	'LC_TIME' EOL	
3631		;		
3632	time_keywords	:	time_keywords time_keyword	
3633			time_keyword	
3634		;		
3635	time_keyword	:	time_keyword_name time_list EOL	
3636			time_keyword_fmt time_string EOL	
3637			time_keyword_opt time_list EOL	
3638		;		
3639	time_keyword_name	:	'abday' 'day' 'abmon' 'mon'	
3640		;		
3641	time_keyword_fmt	:	'd_t_fmt' 'd_fmt' 't_fmt'	
3642			'am_pm' 't_fmt_ampm'	
3643		;		
3644	time_keyword_opt	:	'era' 'era_d_fmt' 'era_t_fmt'	
3645			'era_d_t_fmt' 'alt_digits'	
3646		;		
3647	time_list	:	time_list ' ;' time_string	
3648			time_string	
3649		;		
3650	time_string	:	'"' char_list '"'	
3651		;		
3652	time_tlr	:	'END' 'LC_TIME' EOL	
3653		;		

3654 5.5 Locale Definition Example

3655 The following is an example of a locale definition file that could be used as input to the *localedef*
 3656 utility. It assumes that the utility is executed with the *-f* option, naming a *charmap* file with (at
 3657 least) the following content:

```
3658 CHARMAP
3659 <space>      \x20
3660 <dollar>     \x24
3661 <A>          \101
3662 <a>          \141
3663 <A-acute>    \346
3664 <a-acute>    \365
3665 <A-grave>    \300
3666 <a-grave>    \366
3667 <b>          \142
3668 <C>          \103
3669 <c>          \143
3670 <c-cedilla>  \347
3671 <d>          \x64
3672 <H>          \110
3673 <h>          \150
3674 <eszet>     \xb7
3675 <s>          \x73
3676 <z>          \x7a
3677 END CHARMAP
```

3678 It should not be taken as complete or to represent any actual locale, but only to illustrate the
 3679 syntax.

3680 A further set of examples is offered as part of the **Internationalisation Guide**.

```
3681 #
3682 LC_CTYPE
3683 lower  <a>;<b>;<c>;<c-cedilla>;<d>;...;<z>
3684 upper  A;B;C;Ç;...;Z
3685 space  \x20;\x09;\x0a;\x0b;\x0c;\x0d
3686 blank  \040;\011
3687 toupper (<a>,<A>);(b,B);(c,C);(ç,Ç);(d,D);(z,Z)
3688 END LC_CTYPE
3689 #
3690 LC_COLLATE
3691 #
3692 # The following example of collation is based on the proposed
3693 # Canadian standard Z243.4.1-1990, "Canadian Alphanumeric
3694 # Ordering Standard For Character sets of CSA Z234.4 Standard".
3695 # (Other parts of this example locale definition file do not
3696 # purport to relate to Canada, or to any other real culture.)
3697 # The proposed standard defines a 4-weight collation, such that
3698 # in the first pass, characters are compared without regard to
3699 # case or accents; in second pass, backwards compare without
3700 # regard to case; in the third pass, forward compare without
3701 # regard to diacriticals. In the 3 first passes, non-alphabetic
3702 # characters are ignored; in the fourth pass, only special
3703 # characters are considered, such that "The string that has a
```

```

3704      # special character in the lowest position comes first.  If two
3705      # strings have a special character in the same position, the
3706      # collation value of the special character determines ordering.
3707      #
3708      # Only a subset of the character set is used here; mostly to
3709      # illustrate the set-up.
3710      #
3711      #
3712      collating-symbol <LOW_VALUE>
3713      collating-symbol <LOWER-CASE>
3714      collating-symbol <SUBSCRIPT-LOWER>
3715      collating-symbol <SUPERSCRIPT-LOWER>
3716      collating-symbol <UPPER-CASE>
3717      collating-symbol <NO-ACCENT>
3718      collating-symbol <PECULIAR>
3719      collating-symbol <LIGATURE>
3720      collating-symbol <ACUTE>
3721      collating-symbol <GRAVE>
3722      # Further collating-symbols follow.
3723      #
3724      # Properly, the standard does not include any multi-character
3725      # collating elements; the one below is added for completeness.
3726      #
3727      collating_element <ch> from "<c><h>"
3728      collating_element <CH> from "<C><H>"
3729      collating_element <Ch> from "<C><h>"
3730      #
3731      order_start forward;backward;forward;forward,position
3732      #
3733      # Collating symbols are specified first in the sequence to allocate
3734      # basic collation values to them, lower than that of any character.
3735      <LOW_VALUE>
3736      <LOWER-CASE>
3737      <SUBSCRIPT-LOWER>
3738      <SUPERSCRIPT-LOWER>
3739      <UPPER-CASE>
3740      <NO-ACCENT>
3741      <PECULIAR>
3742      <LIGATURE>
3743      <ACUTE>
3744      <GRAVE>
3745      <RING-ABOVE>
3746      <DIAERESIS>
3747      <TILDE>
3748      # Further collating symbols are given a basic collating value here.
3749      #
3750      # Here follow special characters.
3751      <space>          IGNORE;IGNORE;IGNORE;<space>
3752      # Other special characters follow here.
3753      #

```

```

3754      # Here follow the regular characters.
3755      <a>          <a>;<NO-ACCENT>;<LOWER-CASE>;IGNORE
3756      <A>          <a>;<NO-ACCENT>;<UPPER-CASE>;IGNORE
3757      <a-acute>     <a>;<ACUTE>;<LOWER-CASE>;IGNORE
3758      <A-acute>     <a>;<ACUTE>;<UPPER-CASE>;IGNORE
3759      <a-grave>     <a>;<GRAVE>;<LOWER-CASE>;IGNORE
3760      <A-grave>     <a>;<GRAVE>;<UPPER-CASE>;IGNORE
3761      <ae>          "<a><e>";"<LIGATURE><LIGATURE>";\
3762                  "<LOWER-CASE><LOWER-CASE>";IGNORE
3763      <AE>          "<a><e>";"<LIGATURE><LIGATURE>";\
3764                  "<UPPER-CASE><UPPER-CASE>";IGNORE
3765      <b>          <b>;<NO-ACCENT>;<LOWER-CASE>;IGNORE
3766      <B>          <b>;<NO-ACCENT>;<UPPER-CASE>;IGNORE
3767      <c>          <c>;<NO-ACCENT>;<LOWER-CASE>;IGNORE
3768      <C>          <c>;<NO-ACCENT>;<UPPER-CASE>;IGNORE
3769      <ch>         <ch>;<NO-ACCENT>;<LOWER-CASE>;IGNORE
3770      <Ch>         <ch>;<NO-ACCENT>;<PECULIAR>;IGNORE
3771      <CH>         <ch>;<NO-ACCENT>;<UPPER-CASE>;IGNORE
3772      #
3773      # As an example, the strings "Bach" and "bach" could be encoded (for
3774      # compare purposes) as:
3775      # "Bach"      <b>;<a>;<ch>;<LOW_VALUE>;<NO-ACCENT>;<NO-ACCENT>;\
3776                  <NO-ACCENT>;<LOW_VALUE>;<UPPER>;<LOWER>;<LOWER>;<NULL>
3777      # "bach"     <b>;<a>;<ch>;<LOW_VALUE>;<NO-ACCENT>;<NO-ACCENT>;\
3778                  <NO-ACCENT>;<LOW_VALUE>;<LOWER>;<LOWER>;<LOWER>;<NULL>
3779      #
3780      # The two strings are equal in pass 1 and 2, but differ in pass 3.
3781      #
3782      # Further characters follow.
3783      #
3784      UNDEFINED      IGNORE;IGNORE;IGNORE;IGNORE
3785      #
3786      order_end
3787      #
3788      END LC_COLLATE
3789      #
3790      LC_MONETARY
3791      int_curr_symbol      "USD "
3792      currency_symbol      "$"
3793      mon_decimal_point    "."
3794      mon_grouping         3;0
3795      positive_sign        ""
3796      negative_sign        "- "
3797      p_cs_precedes        1
3798      n_sign_posn          0
3799      END LC_MONETARY
3800      #
3801      LC_NUMERIC
3802      copy "US_en.ASCII"
3803      END LC_NUMERIC
3804      #

```

```

3805      LC_TIME
3806      abday  "Sun"; "Mon"; "Tue"; "Wed"; "Thu"; "Fri"; "Sat "
3807      #
3808      day    "Sunday"; "Monday"; "Tuesday"; "Wednesday"; \
3809           "Thursday"; "Friday"; "Saturday"
3810      #
3811      abmon  "Jan"; "Feb"; "Mar"; "Apr"; "May"; "Jun"; \
3812           "Jul"; "Aug"; "Sep"; "Oct"; "Nov"; "Dec "
3813      #
3814      mon    "January"; "February"; "March"; "April"; \
3815           "May"; "June"; "July"; "August"; "September"; \
3816           "October"; "November"; "December"
3817      #
3818      d_t_fmt "%a %b %d %T %Z %Y\n"
3819      END LC_TIME
3820      #
3821      LC_MESSAGES
3822      yesexpr  "^[yY][[:alpha:]]*" | (OK) "
3823      #
3824      noexpr   "^[nN][[:alpha:]]*"
3825      END LC_MESSAGES

```

Environment Variables

3826

6.1 Environment Variable Definition

Environment variables defined in this chapter affect the operation of multiple utilities, functions and applications. There are other environment variables that are of interest only to specific utilities. Environment variables that apply to a single utility only are defined as part of the utility description. See the **ENVIRONMENT VARIABLES** section of the utility descriptions in the **XCU** specification for information on environment variable usage.

The value of an environment variable is a string of characters. For a C-language program, an array of strings called the environment is made available when a process begins. The array is pointed to by the external variable *environ*, which is defined as:

```
extern char **environ;
```

These strings have the form *name=value*; *names* do not contain the character =. For values to be portable across XSI-conformant systems, the value must be composed of characters from the portable character set (except NUL and as indicated below). There is no meaning associated with the order of strings in the environment. If more than one string in a process' environment has the same *name*, the consequences are undefined.

Environment variable names used by the utilities in the **XCU** specification consist solely of upper-case letters, digits and the "_" (underscore) from the characters defined in Table 4-1 on page 43. Other characters may be permitted by an implementation; applications must tolerate the presence of such names. Upper- and lower-case letters retain their unique identities and are not folded together. The name space of environment variable names containing lower-case letters is reserved for applications. Applications can define any environment variables with names from this name space without modifying the behaviour of the standard utilities.

The *values* that the environment variables may be assigned are not restricted except that they are considered to end with a null byte and the total space used to store the environment and the arguments to the process is limited to {ARG_MAX} bytes.

EX Other *name=value* pairs may be placed in the environment by, for example, calling the *putenv()* function, manipulating the *environ* variable, or by using *envp* arguments when creating a process; see *exec* in the **XSH** specification.

It is unwise to conflict with certain variables that are frequently exported by widely used command interpreters and applications:

3857
3858
3859
3860
3861
3862
3863
3864
3865
3866
3867
3868
3869
3870
3871
3872
3873
3874
3875
3876
3877

ARFLAGS	IFS	MAILPATH	PS1
CC	LANG	MAILRC	PS2
CDPATH	LC_ALL	MAKEFLAGS	PS3
CFLAGS	LC_COLLATE	MAKESHELL	PS4
CHARSET	LC_CTYPE	MANPATH	PWD
COLUMNS	LC_MESSAGES	MBOX	RANDOM
DATMSK	LC_MONETARY	MORE	SECONDS
DEAD	LC_NUMERIC	MSGVERB	SHELL
EDITOR	LC_TIME	NLSPATH	TERM
ENV	LDFLAGS	NPROC	TERMCAP
EXINIT	LEX	OLDPWD	TERMINFO
FC	LFLAGS	OPTARG	TMPDIR
FCEDIT	LINENO	OPTERR	TZ
FFLAGS	LINES	OPTIND	USER
GET	LISTER	PAGER	VISUAL
GFLAGS	LOGNAME	PATH	YACC
HISTFILE	LPDEST	PPID	YFLAGS
HISTORY	MAIL	PRINTER	
HISTSIZE	MAILCHECK	PROCLANG	
HOME	MAILER	PROJECTDIR	

3878
3879
3880
3881
3882
3883
3884
3885
3886
3887

If the variables in the following two sections are present in the environment during the execution of an application or utility, they are given the meaning described below. Some are placed into the environment by the implementation at the time the user logs in; all can be added or changed by the user or any ancestor of the current process. The implementation will add or change environment variables named in this specification set only as specified in this specification set. If they are defined in the application's environment, the utilities in the **XCU** specification and the functions in the **XSH** specification assume they have the specified meaning. Conforming applications must not set these environment variables to have meanings other than as described. See *getenv()* and the **XCU** specification, **Section 2.12, Shell Execution Environment** for methods of accessing these variables.

6.2 Internationalisation Variables

This section describes environment variables that are relevant to the operation of internationalised interfaces described in the CAE Specification, **System Interfaces and Headers, Issue 5** and the CAE Specification, **Commands and Utilities, Issue 5**.

Users may use the following environment variables to announce specific localisation requirements to applications. Applications must retrieve this information using the *setlocale()* function to initialise the correct behaviour of the internationalised interfaces. The descriptions of the internationalisation environment variables describe the resulting behaviour only when the application locale is initialised in this way.

LANG

This variable determines the locale category for native language, local customs and coded character set in the absence of the *LC_ALL* and other *LC_** (*LC_COLLATE*, *LC_CTYPE*, *LC_MESSAGES*, *LC_MONETARY*, *LC_NUMERIC*, *LC_TIME*) environment variables. This can be used by applications to determine the language to use for error messages and instructions, collating sequences, date formats, and so forth.

LC_ALL

This variable determines the values for all locale categories. The value of the *LC_ALL* environment variable has precedence over any of the other environment variables starting with *LC_* (*LC_COLLATE*, *LC_CTYPE*, *LC_MESSAGES*, *LC_MONETARY*, *LC_NUMERIC*, *LC_TIME*) and the *LANG* environment variable.

LC_COLLATE

This variable determines the locale category for character collation. It determines collation information for regular expressions and sorting, including equivalence classes and multi-character collating elements, in various utilities and the *strcoll()* and *strxfrm()* functions. Additional semantics of this variable, if any, are implementation-dependent.

LC_CTYPE

This variable determines the locale category for character handling functions, such as *tolower()*, *toupper()* and *isalpha()*. This environment variable determines the interpretation of sequences of bytes of text data as characters (for example, single- as opposed to multi-byte characters), the classification of characters (for example, alpha, digit, graph) and the behaviour of character classes. Additional semantics of this variable, if any, are implementation-dependent.

LC_MESSAGES

This variable determines the locale category for processing affirmative and negative responses and the language and cultural conventions in which messages should be written. **It also affects the behaviour of the *catopen()* function in determining the message catalogue.** Additional semantics of this variable, if any, are implementation-dependent. The language and cultural conventions of diagnostic and informative messages whose format is unspecified by this specification set should be affected by the setting of *LC_MESSAGES*.

LC_MONETARY

This variable determines the locale category for monetary-related numeric formatting information. Additional semantics of this variable, if any, are implementation-dependent.

LC_NUMERIC

This variable determines the locale category for numeric formatting (for example, thousands separator and radix character) information in various utilities as well as the formatted I/O operations in *printf()* and *scanf()* and the string conversion functions in *strtod()*. Additional semantics of this variable, if any, are implementation-dependent.

LC_TIME

This variable determines the locale category for date and time formatting information. It affects the behaviour of the time functions in *strftime()*. Additional semantics of this variable, if any, are implementation-dependent.

NLSPATH

This variable contains a sequence of templates that the *catopen()* function uses when attempting to locate message catalogues. Each template consists of an optional prefix, one or more substitution fields, a filename and an optional suffix.

For example:

```
NLSPATH= "/system/nlslib/%N.cat"
```

defines that *catopen()* should look for all message catalogues in the directory */system/nlslib*, where the catalogue name should be constructed from the *name* parameter passed to *catopen()* (%N), with the suffix *.cat*.

Substitution fields consist of a "%" symbol, followed by a single-letter keyword. The following keywords are currently defined:

%N The value of the *name* parameter passed to *catopen()*.

%L The value of the LC_MESSAGES category.

%l The *language* element from the LC_MESSAGES category.

%t The *territory* element from the LC_MESSAGES category.

%c The *codeset* element from the LC_MESSAGES category.

%% A single % character.

An empty string is substituted if the specified value is not currently defined. The separators underscore (_) and period (.) are not included in %t and %c substitutions.

Templates defined in *NLSPATH* are separated by colons (:). A leading or two adjacent colons :: is equivalent to specifying %N. For example:

```
NLSPATH= " : %N.cat : /nlslib/%L/%N.cat "
```

indicates to *catopen()* that it should look for the requested message catalogue in *name*, *name.cat* and */nlslib/category/name.cat*, where *category* is the value of the LC_MESSAGES category of the current locale.

Users should not set the *NLSPATH* variable unless they have a specific reason to override the default system path. Doing so causes undefined behaviour in the standard utilities.

The environment variables *LANG*, *LC_ALL*, *LC_COLLATE*, *LC_CTYPE*, *LC_MESSAGES*, *LC_MONETARY*, *LC_NUMERIC*, *LC_TIME* (*LC_**) and *NLSPATH* provide for the support of internationalised applications. The standard utilities make use of these environment variables as described in this section and the individual **ENVIRONMENT VARIABLES** sections for the utilities. If these variables specify locale categories that are not based upon the same underlying codeset, the results are unspecified.

The values of locale categories are determined by a precedence order; the first condition met below determines the value:

1. If the *LC_ALL* environment variable is defined and is not null, the value of *LC_ALL* is used.
2. If the *LC_** environment variable (*LC_COLLATE*, *LC_CTYPE*, *LC_MESSAGES*, *LC_MONETARY*, *LC_NUMERIC*, *LC_TIME*) is defined and is not null, the value of the

environment variable is used to initialise the category that corresponds to the environment variable.

3. If the *LANG* environment variable is defined and is not null, the value of the *LANG* environment variable is used.
4. If the *LANG* environment variable is not set or is set to the empty string, the implementation-dependent default locale is used.

If the locale value is "C" or "POSIX", the POSIX locale is used and the standard utilities behave in accordance with the rules in Section 5.2 on page 50, for the associated category.

If the locale value begins with a slash, it is interpreted as the pathname of a file that was created in the output format used by the *localedef* utility; see **OUTPUT FILES** under *localedef*. Referencing such a pathname will result in that locale being used for the indicated category.

EX If the locale value has the form:

```
language[_territory][.codeset]
```

it refers to an implementation-provided locale, where settings of language, territory and codeset are implementation-dependent.

EX *LC_COLLATE*, *LC_CTYPE*, *LC_MESSAGES*, *LC_MONETARY*, *LC_NUMERIC* and *LC_TIME* are defined to accept an additional field “@*modifier*”, which allows the user to select a specific instance of localisation data within a single category (for example, for selecting the dictionary as opposed to the character ordering of data). The syntax for these environment variables is thus defined as:

```
[language[_territory][.codeset]][@modifier]]
```

For example, if a user wanted to interact with the system in French, but required to sort German text files, *LANG* and *LC_COLLATE* could be defined as:

```
LANG=Fr_FR
LC_COLLATE=De_DE
```

This could be extended to select dictionary collation (say) by use of the @*modifier* field; for example:

```
LC_COLLATE=De_DE@dict
```

An implementation may support other formats.

If the locale value is not recognised by the implementation, the behaviour is unspecified.

At run time, these values are bound to a program's locale by calling the *setlocale()* function.

Additional criteria for determining a valid locale name are implementation-dependent.

6.3 Other Environment Variables

COLUMNS

A decimal integer > 0 used to indicate the user's preferred width in column positions for the terminal screen or window. (See **column position** on page 10.) If this variable is unset or null, the implementation determines the number of columns, appropriate for the terminal or window, in an unspecified manner. When *COLUMNS* is set, any terminal-width information implied by *TERM* will be overridden. Users and portable applications should not set *COLUMNS* unless they wish to override the system selection and produce output unrelated to the terminal characteristics.

The default value for the number of column positions is unspecified because historical implementations use different methods to determine values corresponding to the size of the screen in which the utility is run. This size is typically known to the implementation through the value of *TERM*, or by more elaborate methods such as extensions to the *stty* utility, or knowledge of how the user is dynamically resizing windows on a bit-mapped display terminal. Users should not need to set this variable in the environment unless there is a specific reason to override the implementation's default behaviour, such as to display data in an area arbitrarily smaller than the terminal or window.

DATMSK

Indicates the pathname of the template file used by *getdate()*.

HOME

The system will initialise this variable at the time of login to be a pathname of the user's home directory. See <**pwd.h**>.

LINES

A decimal integer > 0 used to indicate the user's preferred number of lines on a page or the vertical screen or window size in lines. A line in this case is a vertical measure large enough to hold the tallest character in the character set being displayed. If this variable is unset or null, the implementation determines the number of lines, appropriate for the terminal or window (size, terminal baud rate, and so forth), in an unspecified manner. When *LINES* is set, any terminal-height information implied by *TERM* will be overridden. Users and portable applications should not set *LINES* unless they wish to override the system selection and produce output unrelated to the terminal characteristics.

The default value for the number of lines is unspecified because historical implementations use different methods to determine values corresponding to the size of the screen in which the utility is run. This size is typically known to the implementation through the value of *TERM*, or by more elaborate methods such as extensions to the *stty* utility, or knowledge of how the user is dynamically resizing windows on a bit-mapped display terminal. Users should not need to set this variable in the environment unless there is a specific reason to override the implementation's default behaviour, such as to display data in an area arbitrarily smaller than the terminal or window.

LOGNAME

The system will initialise this variable at the time of login to be the user's login name. See <**pwd.h**>. For a value of *LOGNAME* to be portable across implementations of the ISO POSIX-1 standard, the value should be composed of characters from the portable filename character set.

MSGVERB

Describes which message components are to be used in writing messages by *fmtmsg()*.

4056 **PATH** The sequence of path prefixes that certain functions and utilities apply in searching for
 4057 an executable file known only by a filename. The prefixes are separated by a colon (:)
 4058 When a non-zero-length prefix is applied to this filename, a slash is inserted between
 4059 the prefix and the filename. A zero-length prefix is a legacy feature that indicates
 4060 the current working directory. It appears as two adjacent colons (::), as an initial colon
 4061 preceding the rest of the list, or as a trailing colon following the rest of the list. A
 4062 portable application must use an actual pathname (such as `.`) to represent the current
 4063 working directory in *PATH*. The list is searched from beginning to end, applying the
 4064 filename to each prefix, until an executable file with the specified name and
 4065 appropriate execution permissions is found. If the pathname being sought contains a
 4066 slash, the search through the path prefixes will not be performed. If the pathname
 4067 begins with a slash, the specified path is resolved (see **pathname resolution** on page
 4068 22). If *PATH* is unset or is set to null, the path search is implementation-dependent.

4069 **SHELL** A pathname of the user's preferred command language interpreter. If this interpreter
 4070 does not conform to the XSI Shell Command Language in the XCU specification,
 4071 **Chapter 2, Shell Command Language**, utilities may behave differently from those
 4072 described in this specification set.

4073 **TMPDIR** A pathname of a directory made available for programs that need a place to create
 4074 temporary files.

4076 **TERM** The terminal type for which output is to be prepared. This information is used by
 4077 utilities and application programs wishing to exploit special capabilities specific to a
 4078 terminal. The format and allowable values of this environment variable are
 4079 unspecified.

4080 **TZ** Timezone information. The contents of the environment variable named *TZ* are used
 4081 by the *ctime()*, *localtime()*, *strftime()* and *mktime()* functions, and by various utilities, to
 4082 override the default timezone. The value of *TZ* has one of the two forms (spaces
 4083 inserted for clarity):

4084 *:characters*

4085 or:

4086 *std offset dst offset, rule*

4087 If *TZ* is of the first format (that is, if the first character is a colon), the characters
 4088 following the colon are handled in an implementation-dependent manner.

4089 The expanded format (for all *TZ*s whose value does not have a colon as the first
 4090 character) is as follows:

4091 *stdoffset[dst[offset][,start[/time],end[/time]]]*

4092 Where:

4093 *std* and *dst*

4094 Indicates no less than three, nor more than {TZNAME_MAX}, bytes that are
 4095 the designation for the standard (*std*) or the alternative (*dst* — such as
 4096 Daylight Savings Time) timezone. Only *std* is required; if *dst* is missing, then
 4097 the alternative time does not apply in this locale. Upper- and lower-case
 4098 letters are explicitly allowed. Any graphic characters except a leading colon (:)
 4099 or digits, the comma (,), the minus (–), the plus (+), and the null character are
 4100 permitted to appear in these fields, but their meaning is unspecified.

offset Indicates the value one must add to the local time to arrive at Coordinated Universal Time. The *offset* has the form:

$$hh[:mm[:ss]]$$

The minutes (*mm*) and seconds (*ss*) are optional. The hour (*hh*) is required and may be a single digit. The *offset* following *std* is required. If no *offset* follows *dst*, the alternative time is assumed to be one hour ahead of standard time. One or more digits may be used; the value is always interpreted as a decimal number. The hour is between zero and 24, and the minutes (and seconds) if present between zero and 59. Use of values outside these ranges causes undefined behaviour. If preceded by a $-$, the timezone is east of the Prime Meridian; otherwise it is west (which may be indicated by an optional preceding $+$).

rule Indicates when to change to and back from the alternative time. The *rule* has the form:

$$date[/time], date[/time]$$

where the first *date* describes when the change from standard to alternative time occurs and the second *date* describes when the change back happens. Each *time* field describes when, in current local time, the change to the other time is made.

The format of *date* is one of the following:

Jn The Julian day n ($1 \leq n \leq 365$). Leap days are not counted. That is, in all years including leap years February 28 is day 59 and March 1 is day 60. It is impossible to refer explicitly to the occasional February 29.

n The zero-based Julian day ($0 \leq n \leq 365$). Leap days are counted, and it is possible to refer to February 29.

Mm.n.d The d^{th} day ($0 \leq d \leq 6$) of week n of month m of the year ($1 \leq n \leq 5$, $1 \leq m \leq 12$, where week 5 means “the last d day in month m ” which may occur in either the fourth or the fifth week). Week 1 is the first week in which the d^{th} day occurs. Day zero is Sunday.

The *time* has the same format as *offset* except that no leading sign ($-$ or $+$) is allowed. The default, if *time* is not given, is 02:00:00.

Regular Expressions

Note: Two versions of regular expressions are supported in this specification set:

- the historical **Simple Regular Expressions**, which provide backward compatibility, but which may be withdrawn from a future issue of this specification set
- the improved internationalised version that complies with the ISO/IEC 9945-2: 1993 standard.

The first (historical) version is described as part of the *regexp()* function in the **XSH** specification. The second (improved) version is described in this chapter.

Regular Expressions (REs) provide a mechanism to select specific strings from a set of character strings.

Regular expressions are a context-independent syntax that can represent a wide variety of character sets and character set orderings, where these character sets are interpreted according to the current locale. While many regular expressions can be interpreted differently depending on the current locale, many features, such as character class expressions, provide for contextual invariance across locales.

The Basic Regular Expression (BRE) notation and construction rules in Section 7.3 on page 104 apply to most utilities supporting regular expressions. Some utilities, instead, support the Extended Regular Expressions (ERE) described in Section 7.4 on page 109; any exceptions for both cases are noted in the descriptions of the specific utilities using regular expressions. Both BREs and EREs are supported by the Regular Expression Matching interface in the **XSH** specification under *regcomp()*, *regex()* and related functions.

7.1 Regular Expression Definitions

For the purposes of this section, the following definitions apply:

entire regular expression

The concatenated set of one or more BREs or EREs that make up the pattern specified for string selection.

matched

A sequence of zero or more characters is said to be matched by a BRE or ERE when the characters in the sequence correspond to a sequence of characters defined by the pattern.

Matching is based on the bit pattern used for encoding the character, not on the graphic representation of the character. This means that if a character set contains two or more encodings for a graphic symbol, or if the strings searched contain text encoded in more than one codeset, no attempt is made to search for any other representation of the encoded symbol. If that is required, the user can specify equivalence classes containing all variations of the desired graphic symbol.

The search for a matching sequence starts at the beginning of a string and stops when the first sequence matching the expression is found, where *first* is defined to mean “begins earliest in the string”. If the pattern permits a variable number of matching characters and thus there is more than one such sequence starting at that point, the longest such sequence will be matched. For example: the BRE *bb** matches the second to fourth characters of *abbbc*, and the ERE *(wee|week)(knights|night)* matches all ten characters of *weeknights*.

Consistent with the whole match being the longest of the leftmost matches, each subpattern, from left to right, matches the longest possible string. For this purpose, a null string is considered to be longer than no match at all. For example, matching the BRE `\(.*)\.` against `abcdef`, the subexpression `(\1)` is `abcdef`, and matching the BRE `\(a*)\.` against `bc`, the subexpression `(\1)` is the null string.

It is possible to determine what strings correspond to subexpressions by recursively applying the leftmost longest rule to each subexpression, but only with the proviso that the overall match is leftmost longest. For example, matching `\(ac*)c*d[ac]*\1` against `acdacaaa` matches `acdacaaa` (with `\1=a`); simply matching the longest match for `\(ac*)\.` would yield `\1=ac`, but the overall match would be smaller (`acdac`). Conceptually, the implementation must examine every possible match and among those that yield the leftmost longest total matches, pick the one that does the longest match for the leftmost subexpression and so on. Note that this means that matching by subexpressions is context-dependent: a subexpression within a larger RE may match a different string from the one it would match as an independent RE, and two instances of the same subexpression within the same larger RE may match different lengths even in similar sequences of characters. For example, in the ERE `(a.*b)(a.*b)`, the two identical subexpressions would match four and six characters, respectively, of `acbbaccccb`.

When a multi-character collating element in a bracket expression (see Section 7.3.5 on page 105) is involved, the longest sequence will be measured in characters consumed from the string to be matched; that is, the collating element counts not as one element, but as the number of characters it matches.

BRE (ERE) matching a single character

A BRE or ERE that matches either a single character or a single collating element.

Only a BRE or ERE of this type that includes a bracket expression (see Section 7.3.5 on page 105) can match a collating element.

The definition of *single character* has been expanded to include also collating elements consisting of two or more characters; this expansion is applicable only when a bracket expression is included in the BRE or ERE. An example of such a collating element may be the Dutch `ij`, which collates as a `y`. In some encodings, a ligature “i with j” exists as a character and would represent a single-character collating element. In another encoding, no such ligature exists, and the two-character sequence `ij` is defined as a multi-character collating element. Outside brackets, the `ij` is treated as a two-character RE and matches the same characters in a string. Historically, a bracket expression only matched a single character. If, however, the bracket expression defines, for example, a range that includes `ij`, then this particular bracket expression will also match a sequence of the two characters `i` and `j` in the string.

BRE (ERE) matching multiple characters

A BRE or ERE that matches a concatenation of single characters or collating elements.

Such a BRE or ERE is made up from a BRE (ERE) matching a single character and BRE (ERE) special characters.

invalid

This section uses the term *invalid* for certain constructs or conditions. Invalid REs will cause the utility or function using the RE to generate an error condition. When *invalid* is not used, violations of the specified syntax or semantics for REs produce undefined results: this may entail an error, enabling an extended syntax for that RE, or using the construct in error as literal characters to be matched. For example, the BRE construct `\{1,2,3\}` does not comply with the grammar. A portable application cannot rely on it producing an error nor matching the literal characters `\{1,2,3\}`.

4222 7.2 Regular Expression General Requirements

4223 The requirements in this section apply to both basic and extended regular expressions.

4224 The use of regular expressions is generally associated with text processing. REs (BREs and
4225 EREs) operate on text strings; that is, zero or more characters followed by an end-of-string
4226 delimiter (typically NUL). Some utilities employing regular expressions limit the processing to
4227 lines; that is, zero or more characters followed by a newline character. In the regular expression
4228 processing described in this specification, the newline character is regarded as an ordinary
4229 character and both a period and a non-matching list can match one. The XCU specification
4230 specifies within the individual descriptions of those standard utilities employing regular
4231 expressions whether they permit matching of newline characters; if not stated otherwise, the use
4232 of literal newline characters or any escape sequence equivalent produces undefined results.
4233 Those utilities (like *grep*) that do not allow newline characters to match are responsible for
4234 eliminating any newline character from strings before matching against the RE. The *regcomp()*
4235 function in the XSH specification, however, can provide support for such processing without
4236 violating the rules of this section.

4237 The interfaces specified in this specification set do not permit the inclusion of a NUL character in
4238 an RE or in the string to be matched. If during the operation of a standard utility a NUL is
4239 included in the text designated to be matched, that NUL may designate the end of the text string
4240 for the purposes of matching.

4241 When a standard utility or function that uses regular expressions specifies that pattern matching
4242 will be performed without regard to the case (upper- or lower-) of either data or patterns, then
4243 when each character in the string is matched against the pattern, not only the character, but also
4244 its case counterpart (if any), will be matched. This definition of case-insensitive processing is
4245 intended to allow matching of multi-character collating elements as well as characters. For
4246 instance, as each character in the string is matched using both its cases, the RE `[.Ch.]` when
4247 matched against the string `char`, is in reality matched against `ch`, `Ch`, `cH` and `CH`.

4248 The implementation will support any regular expression that does not exceed 256 bytes in
4249 length.

4250 7.3 Basic Regular Expressions

4251 7.3.1 BREs Matching a Single Character or Collating Element

4252 A BRE ordinary character, a special character preceded by a backslash or a period matches a
4253 single character. A bracket expression matches a single character or a single collating element.

4254 7.3.2 BRE Ordinary Characters

4255 An ordinary character is a BRE that matches itself: any character in the supported character set,
4256 except for the BRE special characters listed in Section 7.3.3.

4257 The interpretation of an ordinary character preceded by a backslash (\) is undefined, except for:

- 4258 1. the characters), (, { and }
- 4259 2. the digits 1 to 9 inclusive (see Section 7.3.6 on page 107)
- 4260 3. a character inside a bracket expression.

4261 7.3.3 BRE Special Characters

4262 A *BRE special character* has special properties in certain contexts. Outside those contexts, or
4263 when preceded by a backslash, such a character will be a BRE that matches the special character
4264 itself. The BRE special characters and the contexts in which they have their special meaning are:

4265 . [\ The period, left-bracket and backslash is special except when used in a bracket
4266 expression (see Section 7.3.5 on page 105). An expression containing a [that is not
4267 preceded by a backslash and is not part of a bracket expression produces undefined
4268 results.

4269 * The asterisk is special except when used:
4270 • in a bracket expression
4271 • as the first character of an entire BRE (after an initial ^, if any)
4272 • as the first character of a subexpression (after an initial ^, if any); see Section 7.3.6 on
4273 page 107.

4274 ^ The circumflex is special when used:
4275 • as an anchor (see Section 7.3.8 on page 108)
4276 • as the first character of a bracket expression (see Section 7.3.5 on page 105).

4277 \$ The dollar sign is special when used as an anchor.

4278 7.3.4 Periods in BREs

4279 A period (.), when used outside a bracket expression, is a BRE that matches any character in the
4280 supported character set except NUL.

4281 **7.3.5 RE Bracket Expression**

4282 A bracket expression (an expression enclosed in square brackets, []) is an RE that matches a
 4283 single collating element contained in the non-empty set of collating elements represented by the
 4284 bracket expression.

4285 The following rules and definitions apply to bracket expressions:

- 4286 1. A *bracket expression* is either a matching list expression or a non-matching list expression. It
 4287 consists of one or more expressions: collating elements, collating symbols, equivalence
 4288 classes, character classes or range expressions. Portable applications must not use range
 4289 expressions, even though all implementations support them. The right-bracket (]) loses its
 4290 special meaning and represents itself in a bracket expression if it occurs first in the list
 4291 (after an initial circumflex (^), if any). Otherwise, it terminates the bracket expression,
 4292 unless it appears in a collating symbol (such as [.].) or is the ending right-bracket for a
 4293 collating symbol, equivalence class or character class. The special characters:

4294 . * [\

4295 (period, asterisk, left-bracket and backslash, respectively) lose their special meaning within
 4296 a bracket expression.

4297 The character sequences:

4298 [. [= [:

4299 (left-bracket followed by a period, equals-sign or colon) are special inside a bracket
 4300 expression and are used to delimit collating symbols, equivalence class expressions and
 4301 character class expressions. These symbols must be followed by a valid expression and the
 4302 matching terminating sequence .], =] or :], as described in the following items.

- 4303 2. A *matching list* expression specifies a list that matches any one of the expressions
 4304 represented in the list. The first character in the list must not be the circumflex. For
 4305 example, [abc] is an RE that matches any of the characters a, b or c.

- 4306 3. A *non-matching list* expression begins with a circumflex (^), and specifies a list that matches
 4307 any character or collating element except for the expressions represented in the list after
 4308 the leading circumflex. For example, [^abc] is an RE that matches any character or collating
 4309 element except the characters a, b or c. The circumflex will have this special meaning only
 4310 when it occurs first in the list, immediately following the left-bracket.

- 4311 4. A *collating symbol* is a collating element enclosed within bracket-period ([. .]) delimiters.
 4312 Collating elements are defined as described in **Collation Order** on page 63. Multi-
 4313 character collating elements must be represented as collating symbols when it is necessary
 4314 to distinguish them from a list of the individual characters that make up the multi-
 4315 character collating element. For example, if the string ch is a collating element in the
 4316 current collation sequence with the associated collating symbol <ch>, the expression
 4317 [[.ch.]] will be treated as an RE matching the character sequence ch, while [ch] will be
 4318 treated as an RE matching c or h. Collating symbols will be recognised only inside bracket
 4319 expressions. This implies that the RE [[.ch.]]*c matches the first to fifth character in the
 4320 string chchch. If the string is not a collating element in the current collating sequence
 4321 definition, or if the collating element has no characters associated with it (for example, see
 4322 the symbol <HIGH> in the example collation definition shown in **Collation Order** on page
 4323 63), the symbol will be treated as an invalid expression.

- 4324 5. An *equivalence class expression* represents the set of collating elements belonging to an
 4325 equivalence class, as described in **Collation Order**. Only primary equivalence classes will
 4326 be recognised. The class is expressed by enclosing any one of the collating elements in the

equivalence class within bracket-equal (`[= =]`) delimiters. For example, if `a`, `à` and `â` belong to the same equivalence class, then `[a=à]`, `[a=â]` and `[à=â]` will each be equivalent to `[aââ]`. If the collating element does not belong to an equivalence class, the equivalence class expression will be treated as a *collating symbol*.

6. A *character class expression* represents the set of characters belonging to a character class, as defined in the `LC_CTYPE` category in the current locale. All character classes specified in the current locale will be recognised. A character class expression is expressed as a character class name enclosed within bracket-colon (`[:]`) delimiters.

The following character class expressions are supported in all locales:

<code>[:alnum:]</code>	<code>[:cntrl:]</code>	<code>[:lower:]</code>	<code>[:space:]</code>
<code>[:alpha:]</code>	<code>[:digit:]</code>	<code>[:print:]</code>	<code>[:upper:]</code>
<code>[:blank:]</code>	<code>[:graph:]</code>	<code>[:punct:]</code>	<code>[:xdigit:]</code>

In addition, character class expressions of the form:

`[:name:]`

are recognised in those locales where the *name* keyword has been given a **charclass** definition in the `LC_CTYPE` category.

7. A *range expression* represents the set of collating elements that fall between two elements in the current collation sequence, inclusively. It is expressed as the starting point and the ending point separated by a hyphen (`-`).

Range expressions must not be used in portable applications because their behaviour is dependent on the collating sequence. Ranges will be treated according to the current collating sequence, and include such characters that fall within the range based on that collating sequence, regardless of character values. This, however, means that the interpretation will differ depending on collating sequence. If, for instance, one collating sequence defines `ä` as a variant of `a`, while another defines it as a letter following `z`, then the expression `[ä-z]` is valid in the first language and invalid in the second.

In the following, all examples assume the collation sequence specified for the POSIX locale, unless another collation sequence is specifically defined.

The starting range point and the ending range point must be a collating element or collating symbol. An equivalence class expression used as a starting or ending point of a range expression produces unspecified results. An equivalence class can be used portably within a bracket expression, but only outside the range. For example, the unspecified expression `[a-e]-f]` should be given as `[a-e]e-f]`. The ending range point must collate equal to or higher than the starting range point; otherwise, the expression will be treated as invalid. The order used is the order in which the collating elements are specified in the current collation definition. One-to-many mappings (see the description of `LC_COLLATE` in Chapter 5 on page 49) will not be performed. For example, assuming that the character *eszet* (`ß`) is placed in the collation sequence after `r` and `s`, but before `t` and that it maps to the sequence `ss` for collation purposes, then the expression `[r-s]` matches only `r` and `s`, but the expression `[s-t]` matches `s`, `ß` or `t`.

The interpretation of range expressions where the ending range point is also the starting range point of a subsequent range expression (for instance `[a-m-o]`) is undefined.

The hyphen character will be treated as itself if it occurs first (after an initial `^`, if any) or last in the list, or as an ending range point in a range expression. As examples, the expressions `[-ac]` and `[ac-]` are equivalent and match any of the characters `a`, `c` or `-`; `[^ac]` and `[^ac-]` are equivalent and match any characters except `a`, `c` or `-`; the expression `[%--]` matches

any of the characters between % and – inclusive; the expression [–@] matches any of the characters between – and @ inclusive; and the expression [a–@] is invalid, because the letter a follows the symbol – in the POSIX locale. To use a hyphen as the starting range point, it must either come first in the bracket expression or be specified as a collating symbol, for example: [[.–.]–0], which matches either a right bracket or any character or collating element that collates between hyphen and 0, inclusive.

If a bracket expression must specify both – and], the] must be placed first (after the ^, if any) and the – last within the bracket expression.

7.3.6 BREs Matching Multiple Characters

The following rules can be used to construct BREs matching multiple characters from BREs matching a single character:

1. The concatenation of BREs matches the concatenation of the strings matched by each component of the BRE.
2. A *subexpression* can be defined within a BRE by enclosing it between the character pairs \ (and \) . Such a subexpression matches whatever it would have matched without the \ (and \), except that anchoring within subexpressions is optional behaviour; see Section 7.3.8 on page 108. Subexpressions can be arbitrarily nested.
3. The *back-reference* expression \ *n* matches the same (possibly empty) string of characters as was matched by a subexpression enclosed between \ (and \) preceding the \ *n*. The character *n* must be a digit from 1 to 9 inclusive, *n*th subexpression (the one that begins with the *n*th \ (and ends with the corresponding paired \)). The expression is invalid if less than *n* subexpressions precede the \ *n*. For example, the expression ^ \ (. * \) \ 1 \$ matches a line consisting of two adjacent appearances of the same string, and the expression \ (a \) * \ 1 fails to match a. The limit of nine back-references to subexpressions in the RE is based on the use of a single digit identifier. This does not imply that only nine subexpressions are allowed in REs. The following is a valid BRE with ten subexpressions:
`\ (\ (\ (ab \) * c \) * d \) \ (ef \) * \ (gh \) \ { 2 \ } \ (ij \) * \ (kl \) * \ (mn \) * \ (op \) * \ (qr \) *`
4. When a BRE matching a single character, a subexpression or a back-reference is followed by the special character asterisk (*), together with that asterisk it matches what zero or more consecutive occurrences of the BRE would match. For example, [ab]* and [ab][ab] are equivalent when matching the string ab.
5. When a BRE matching a single character, a subexpression or a back-reference is followed by an *interval expression* of the format \ { *m* \ }, \ { *m* , \ } or \ { *m* , *n* \ }, together with that interval expression it matches what repeated consecutive occurrences of the BRE would match. The values of *m* and *n* will be decimal integers in the range $0 \leq m \leq n \leq \{RE_DUP_MAX\}$, where *m* specifies the exact or minimum number of occurrences and *n* specifies the maximum number of occurrences. The expression \ { *m* \ } matches exactly *m* occurrences of the preceding BRE, \ { *m* , \ } matches at least *m* occurrences and \ { *m* , *n* \ } matches any number of occurrences between *m* and *n*, inclusive.

For example, in the string abababcccccd the BRE c \ { 3 \ } is matched by characters seven to nine, the BRE \ (ab \) \ { 4 , \ } is not matched at all and the BRE c \ { 1 , 3 \ } d is matched by characters ten to thirteen.

The behaviour of multiple adjacent duplication symbols (* and intervals) produces undefined results.

4417 **7.3.7 BRE Precedence**

4418 The order of precedence is as shown in the following table:
 4419
 4420

BRE Precedence (from high to low)	
collation-related bracket symbols	[=] [: :] [. .]
escaped characters	\< <i>special character</i> >
bracket expression	[]
subexpressions/back-references	\(\) \n
single-character-BRE duplication	* \{m,n\}
concatenation	
anchoring	^ \$

4428 **7.3.8 BRE Expression Anchoring**

4429 A BRE can be limited to matching strings that begin or end a line; this is called *anchoring*. The
 4430 circumflex and dollar sign special characters will be considered BRE anchors in the following
 4431 contexts:

- 4432 1. A circumflex (^) is an anchor when used as the first character of an entire BRE. The
 4433 implementation may treat circumflex as an anchor when used as the first character of a
 4434 subexpression. The circumflex will anchor the expression (or optionally subexpression) to
 4435 the beginning of a string; only sequences starting at the first character of a string will be
 4436 matched by the BRE. For example, the BRE ^ab matches ab in the string abcdef, but fails to
 4437 match in the string cdefab. The BRE \(^ab\) may match the former string. A portable BRE
 4438 must escape a leading circumflex in a subexpression to match a literal circumflex.
- 4439 2. A dollar sign (\$) is an anchor when used as the last character of an entire BRE. The
 4440 implementation may treat a dollar sign as an anchor when used as the last character of a
 4441 subexpression. The dollar sign will anchor the expression (or optionally subexpression) to
 4442 the end of the string being matched; the dollar sign can be said to match the end-of-string
 4443 following the last character.
- 4444 3. A BRE anchored by both "^" and "\$" matches only an entire string. For example, the BRE
 4445 ^abcdef\$ matches strings consisting only of abcdef.

4446 7.4 Extended Regular Expressions

4447 The *extended regular expression* (ERE) notation and construction rules will apply to utilities
 4448 defined as using extended regular expressions; any exceptions to the following rules are noted in
 4449 the descriptions of the specific utilities using EREs.

4450 7.4.1 EREs Matching a Single Character or Collating Element

4451 An ERE ordinary character, a special character preceded by a backslash or a period matches a
 4452 single character. A bracket expression matches a single character or a single collating element.
 4453 An *ERE matching a single character* enclosed in parentheses matches the same as the ERE without
 4454 parentheses would have matched.

4455 7.4.2 ERE Ordinary Characters

4456 An *ordinary character* is an ERE that matches itself. An ordinary character is any character in the
 4457 supported character set, except for the ERE special characters listed in Section 7.4.3. The
 4458 interpretation of an ordinary character preceded by a backslash (\) is undefined.

4459 7.4.3 ERE Special Characters

4460 An *ERE special character* has special properties in certain contexts. Outside those contexts, or
 4461 when preceded by a backslash, such a character is an ERE that matches the special character
 4462 itself. The extended regular expression special characters and the contexts in which they have
 4463 their special meaning are:

4464 . [\ (The period, left-bracket, backslash and left-parenthesis are special except when
 4465 used in a bracket expression (see Section 7.3.5 on page 105). Outside a bracket
 4466 expression, a left-parenthesis immediately followed by a right-parenthesis
 4467 produces undefined results.

4468) The right-parenthesis is special when matched with a preceding left-parenthesis,
 4469 both outside a bracket expression.

4470 * + ? { The asterisk, plus-sign, question-mark and left-brace are special except when used
 4471 in a bracket expression (see Section 7.3.5 on page 105). Any of the following uses
 4472 produce undefined results:

- 4473 • if these characters appear first in an ERE, or immediately following a vertical-
 4474 line, circumflex or left-parenthesis
- 4475 • if a left-brace is not part of a valid interval expression.

4476 | The vertical-line is special except when used in a bracket expression (see Section
 4477 7.3.5 on page 105). A vertical-line appearing first or last in an ERE, or immediately
 4478 following a vertical-line or a left-parenthesis, or immediately preceding a right-
 4479 parenthesis, produces undefined results.

4480 ^ The circumflex is special when used:

- 4481 • as an anchor (see Section 7.4.9 on page 111)
- 4482 • as the first character of a bracket expression (see Section 7.3.5 on page 105).

4483 \$ The dollar sign is special when used as an anchor.

4484 **7.4.4 Periods in EREs**

4485 A period (.), when used outside a bracket expression, is an ERE that matches any character in the
 4486 supported character set except NUL.

4487 **7.4.5 ERE Bracket Expression**

4488 The rules for ERE Bracket Expressions are the same as for Basic Regular Expressions; see Section
 4489 7.3.5 on page 105.

4490 **7.4.6 EREs Matching Multiple Characters**

4491 The following rules will be used to construct EREs matching multiple characters from EREs
 4492 matching a single character:

- 4493 1. A *concatenation of EREs* matches the concatenation of the character sequences matched by
 4494 each component of the ERE. A concatenation of EREs enclosed in parentheses matches
 4495 whatever the concatenation without the parentheses matches. For example, both the ERE
 4496 `cd` and the ERE `(cd)` are matched by the third and fourth character of the string
 4497 `abcdefabcdef`.
- 4498 2. When an ERE matching a single character or an ERE enclosed in parentheses is followed by
 4499 the special character plus-sign (+), together with that plus-sign it matches what one or
 4500 more consecutive occurrences of the ERE would match. For example, the ERE `b+(bc)`
 4501 matches the fourth to seventh characters in the string `acabbbcd`. And, `[ab]+` and `[ab][ab]*`
 4502 are equivalent.
- 4503 3. When an ERE matching a single character or an ERE enclosed in parentheses is followed by
 4504 the special character asterisk (*), together with that asterisk it matches what zero or more
 4505 consecutive occurrences of the ERE would match. For example, the ERE `b*c` matches the
 4506 first character in the string `cabbbcd`, and the ERE `b*cd` matches the third to seventh
 4507 characters in the string `cabbbcd`. And, `[ab]*` and `[ab][ab]` are equivalent when
 4508 matching the string `ab`.
- 4509 4. When an ERE matching a single character or an ERE enclosed in parentheses is followed by
 4510 the special character question-mark (?), together with that question-mark it matches what
 4511 zero or one consecutive occurrences of the ERE would match. For example, the ERE `b?c`
 4512 matches the second character in the string `acabbbcd`.
- 4513 5. When an ERE matching a single character or an ERE enclosed in parentheses is followed by
 4514 an *interval expression* of the format `{m}`, `{m,}` or `{m,n}`, together with that interval expression
 4515 it matches what repeated consecutive occurrences of the ERE would match. The values of
 4516 *m* and *n* will be decimal integers in the range $0 \leq m \leq n \leq \{\text{RE_DUP_MAX}\}$, where *m*
 4517 specifies the exact or minimum number of occurrences and *n* specifies the maximum
 4518 number of occurrences. The expression `{m}` matches exactly *m* occurrences of the
 4519 preceding ERE, `{m,}` matches at least *m* occurrences and `{m,n}` matches any number of
 4520 occurrences between *m* and *n*, inclusive.

4521 For example, in the string `abababcccccd` the ERE `c{3}` is matched by characters seven to
 4522 nine and the ERE `(ab){2,}` is matched by characters one to six.

4523 The behaviour of multiple adjacent duplication symbols (+, *, ? and intervals) produces
 4524 undefined results.

4525 7.4.7 ERE Alternation

4526 Two EREs separated by the special character vertical-line (|) match a string that is matched by
 4527 either. For example, the ERE `a((bc)|d)` matches the string `abc` and the string `ad`. Single
 4528 characters, or expressions matching single characters, separated by the vertical bar and enclosed
 4529 in parentheses, will be treated as an ERE matching a single character.

4530 7.4.8 ERE Precedence

4531 The order of precedence will be as shown in the following table:

ERE Precedence (from high to low)	
4534 collation-related bracket symbols	<code>[= =] [: :] [. .]</code>
4535 escaped characters	<code>\<special character></code>
4536 bracket expression	<code>[]</code>
4537 grouping	<code>()</code>
4538 single-character-ERE duplication	<code>* + ? {m,n}</code>
4539 concatenation	
4540 anchoring	<code>^ \$</code>
4541 alternation	<code> </code>

4542 For example, the ERE `abba | cde` matches either the string `abba` or the string `cde` (rather than the
 4543 string `abbade` or `abbcde`, because concatenation has a higher order of precedence than
 4544 alternation).

4545 7.4.9 ERE Expression Anchoring

4546 An ERE can be limited to matching strings that begin or end a line; this is called *anchoring*. The
 4547 circumflex and dollar sign special characters are considered ERE anchors when used anywhere
 4548 outside a bracket expression. This has the following effects:

- 4549 1. A circumflex (^) outside a bracket expression anchors the expression or subexpression it
 4550 begins to the beginning of a string; such an expression or subexpression can match only a
 4551 sequence starting at the first character of a string. For example, the EREs `^ab` and `(^ab)`
 4552 match `ab` in the string `abcdef`, but fail to match in the string `cdefab`, and the ERE `a^b` is
 4553 valid, but can never match because the `a` prevents the expression `^b` from matching starting
 4554 at the first character.
- 4555 2. A dollar sign (\$) outside a bracket expression anchors the expression or subexpression it
 4556 ends to the end of a string; such an expression or subexpression can match only a sequence
 4557 ending at the last character of a string. For example, the EREs `ef$` and `(ef$)` match `ef` in the
 4558 string `abcdef`, but fail to match in the string `cdefab`, and the ERE `e$f` is valid, but can never
 4559 match because the `f` prevents the expression `e$` from matching ending at the last character.

4560 7.5 Regular Expression Grammar

Grammars describing the syntax of both basic and extended regular expressions are presented in this section. The grammar takes precedence over the text. See the XCU specification, **Section 1.8, Grammar Conventions**.

4564 7.5.1 BRE/ERE Grammar Lexical Conventions

4565 The lexical conventions for regular expressions are as described in this section.

4566 Except as noted, the longest possible token or delimiter beginning at a given point will be
4567 recognised.

4568 The following tokens will be processed (in addition to those string constants shown in the
4569 grammar):

4570	COLL_ELEM	Any single-character collating element, unless it is a META_CHAR.
------	-----------	---

4571	BACKREF	Applicable only to basic regular expressions. The character string
4572		consisting of "\" followed by a single-digit numeral, 1 to 9.

4573	DUP_COUNT	Represents a numeric constant. It is an integer in the range $0 \leq$
4574		DUP_COUNT \leq {RE_DUP_MAX}. This token will only be recognised
4575		when the context of the grammar requires it. At all other times, digits not
4576		preceded by "\" will be treated as ORD_CHAR.

4577 META_CHAR One of the characters:

4578 ^ when found first in a bracket expression

4579 – when found anywhere but first (after an initial "^", if any) or last in a
4580 bracket expression, or as the ending range point in a range
4581 expression

4582] when found anywhere but first (after an initial "" if any) in a bracket
4583 expression.

4584	L_ANCHOR	Applicable only to basic regular expressions. The character "^" when it
4585		appears as the first character of a basic regular expression and when not
4586		QUOTED_CHAR. The "^" may be recognised as an anchor elsewhere; see
4587		Section 7.3.8 on page 108.

4588	ORD_CHAR	A character, other than one of the special characters in SPEC_CHAR.
------	----------	---

4589 QUOTED_CHAR In a BRE, one of the character sequences:

4590 \^ \. * \[\\$ \

4591 In an ERE, one of the character sequences:

4592 \^ \. \| \\$ \(\) |

4593 * \+ \? \{ \\

4594 R_ANCHOR (Applicable only to basic regular expressions.) The character "\$" when it
4595 appears as the last character of a basic regular expression and when not
4596 QUOTED_CHAR. The "\$" may be recognised as an anchor elsewhere; see
4597 Section 7.3.8 on page 108.

4598 SPEC_CHAR For basic regular expressions, will be one of the following special
4599 characters:

4600 . anywhere outside bracket expressions
 4601 \ anywhere outside bracket expressions
 4602 [anywhere outside bracket expressions
 4603 ^ when used as an anchor (see Section 7.3.8 on page 108) or when first
 4604 in a bracket expression
 4605 \$ when used as an anchor
 4606 * anywhere except: first in an entire RE; anywhere in a bracket
 4607 expression; directly following \(: directly following an anchoring "^.
 4608 For extended regular expressions, will be one of the following special
 4609 characters found anywhere outside bracket expressions:
 4610 ^ . [\$ () | * + ? { \\
 4611 (The close-parenthesis is considered special in this context only if
 4612 matched with a preceding open-parenthesis.)

4613 7.5.2 RE and Bracket Expression Grammar

4614 This section presents the grammar for basic regular expressions, including the bracket
 4615 expression grammar that is common to both BREs and EREs.

```

4616 %token    ORD_CHAR QUOTED_CHAR DUP_COUNT
4617 %token    BACKREF L_ANCHOR R_ANCHOR
4618 %token    Back_open_paren Back_close_paren
4619 /*      '\('      '\)'      */
4620 %token    Back_open_brace Back_close_brace
4621 /*      '\{'      '\}'      */
4622 /* The following tokens are for the Bracket Expression
4623    grammar common to both RES and ERES. */
4624 %token    COLL_ELEM META_CHAR
4625 %token    Open_equal Equal_close Open_dot Dot_close Open_colon Colon_close
4626 /*      '['      '='      '['      '.'      '['      ':'      ':'      */
4627 %token    class_name
4628 /* class_name is a keyword to the LC_CTYPE locale category */
4629 /* (representing a character class) in the current locale */
4630 /* and is only recognised between [: and :] */
4631 %start    basic_reg_exp
4632 %%
4633 /* -----
4634    Basic Regular Expression
4635    -----
4636 */
4637 basic_reg_exp : RE_expression
4638              | L_ANCHOR
4639              | L_ANCHOR R_ANCHOR
4640              | L_ANCHOR R_ANCHOR
4641              | L_ANCHOR RE_expression
  
```

```

4642          | RE_expression R_ANCHOR
4643          | L_ANCHOR RE_expression R_ANCHOR
4644          ;
4645 RE_expression : simple_RE
4646          | RE_expression simple_RE
4647          ;
4648 simple_RE : nondupl_RE
4649          | nondupl_RE RE_dupl_symbol
4650          ;
4651 nondupl_RE : one_character_RE
4652          | Back_open_paren RE_expression Back_close_paren
4653          | Back_open_paren Back_close_paren
4654          | BACKREF
4655          ;
4656 one_character_RE : ORD_CHAR
4657          | QUOTED_CHAR
4658          | '.'
4659          | bracket_expression
4660          ;
4661 RE_dupl_symbol : '*'
4662          | Back_open_brace DUP_COUNT Back_close_brace
4663          | Back_open_brace DUP_COUNT ',' Back_close_brace
4664          | Back_open_brace DUP_COUNT ',' DUP_COUNT Back_close_brace
4665          ;

4666 /* -----
4667    Bracket Expression
4668    -----
4669 */
4670 bracket_expression : '[' matching_list ']'
4671          | '[' nonmatching_list ']'
4672          ;
4673 matching_list : bracket_list
4674          ;
4675 nonmatching_list : '^' bracket_list
4676          ;
4677 bracket_list : follow_list
4678          | follow_list '-'
4679          ;
4680 follow_list : expression_term
4681          | follow_list expression_term
4682          ;
4683 expression_term : single_expression
4684          | range_expression
4685          ;
4686 single_expression : end_range
4687          | character_class
4688          | equivalence_class
4689          ;
4690 range_expression : start_range end_range
4691          | start_range '-'
4692          ;
4693 start_range : end_range '-'

```

```

4694          ;
4695  end_range      : COLL_ELEM
4696                | collating_symbol
4697          ;
4698  collating_symbol : Open_dot COLL_ELEM Dot_close
4699                | Open_dot META_CHAR Dot_close
4700          ;
4701  equivalence_class : Open_equal COLL_ELEM Equal_close
4702          ;
4703  character_class : Open_colon class_name Colon_close
4704          ;

```

4705 The BRE grammar does not permit `L_ANCHOR` or `R_ANCHOR` inside `\(` and `\)` (which implies
4706 that `^` and `$` are ordinary characters). This reflects the semantic limits on the application, as
4707 noted in Section 7.3.8 on page 108. Implementations are permitted to extend the language to
4708 interpret `"^"` and `"$"` as anchors in these locations, and as such, portable applications cannot use
4709 unescaped `"^"` and `"$"` in positions inside `\(` and `\)` that might be interpreted as anchors.

4710 7.5.3 ERE Grammar

4711 This section presents the grammar for extended regular expressions, excluding the bracket
4712 expression grammar.

4713 **Note:** The bracket expression grammar and the associated `%token` lines are identical between
4714 BREs and EREs. It has been omitted from the ERE section to avoid unnecessary
4715 editorial duplication.

```

4716 %token  ORD_CHAR QUOTED_CHAR DUP_COUNT
4717 %start  extended_reg_exp
4718 %%
4719 /* -----
4720    Extended Regular Expression
4721    ----- */
4722
4723 extended_reg_exp : ERE_branch
4724                 | extended_reg_exp ' | ' ERE_branch
4725                 ;
4726 ERE_branch      : ERE_expression
4727                 | ERE_branch ERE_expression
4728                 ;
4729 ERE_expression  : one_character_ERE
4730                 | '^'
4731                 | '$'
4732                 | '(' extended_reg_exp ')'
4733                 | ERE_expression ERE_dupl_symbol
4734                 ;
4735 one_character_ERE : ORD_CHAR
4736                 | QUOTED_CHAR
4737                 | '.'
4738                 | bracket_expression
4739                 ;
4740 ERE_dupl_symbol  : '*'
4741                 | '+'
4742                 | '?'

```

```

4743         | ' { ' DUP_COUNT ' } '
4744         | ' { ' DUP_COUNT ' , ' ' } '
4745         | ' { ' DUP_COUNT ' , ' DUP_COUNT ' } '
4746         ;

```

4747 The ERE grammar does not permit several constructs that previous sections specify as having
 4748 undefined results:

- 4749 • ORD_CHAR preceded by "\"
- 4750 • one or more ERE_dupl_symbols appearing first in an ERE, or immediately following "|", "^"
 4751 or "("
- 4752 • "{" not part of a valid ERE_dupl_symbol
- 4753 • "|" appearing first or last in an ERE, or immediately following "|" or "(", or immediately
 4754 preceding ")".

4755 Implementations are permitted to extend the language to allow these. Portable applications
 4756 cannot use such constructs.

Directory Structure and Devices

4757

8.1 Directory Structure and Files

The following directories exist on conforming systems and must be used as described. Portable applications cannot assume the ability to create files in any of these directories.

/ The root directory.

/dev Contains **/dev/console**, **/dev/null** and **/dev/tty**, described below.

The following directory exists on conforming systems and is used as described.

/tmp A directory made available for programs that need a place to create temporary files. Applications are allowed to create files in this directory, but cannot assume that such files are preserved between invocations of the application.

The **/tmp** directory is defined to accommodate historical applications that assume its availability. Applications are encouraged to use the contents of *TMPDIR* for creating temporary files rather than the specific name **/tmp**. See *tempnam()* in the *XSH* specification.

The following files exist on conforming systems and are both readable and writable.

/dev/null An infinite data source and data sink. Data written to **/dev/null** is discarded. Reads from **/dev/null** always return end-of-file (EOF).

/dev/tty In each process, a synonym for the controlling terminal associated with the process group of that process, if any. It is useful for programs or shell procedures that wish to be sure of writing messages to or reading data from the terminal no matter how output has been redirected. It can also be used for programs that demand the name of a file for output, when typed output is desired and it is tiresome to find out what terminal is currently in use.

The following file exists on conforming systems and need not be readable or writable:

/dev/console The **/dev/console** file is a generic name given to the system console. It is usually linked to a particular machine-dependent special file. It provides a basic I/O interface to the system console.

8.2 Output Devices and Terminal Types

The utilities in the XCU specification historically have been implemented on a wide range of terminal types, but a conforming implementation need not support all features of all utilities on every conceivable terminal. This specification set states which features are optional for certain classes of terminals in the individual utility description sections. The implementation will document which terminal types it supports and which of these features and utilities are not supported by each terminal.

When a feature or utility is not supported on a specific terminal type, as allowed by this specification set, and the implementation considers such a condition to be an error preventing use of the feature or utility, the implementation will indicate such conditions through diagnostic messages or exit status values or both (as appropriate to the specific utility description) that inform the user that the terminal type lacks the appropriate capability.

This specification set uses a notational convention based on historical practice that identifies some of the control characters defined in Section 4.1 on page 43 in a manner easily remembered by users on many terminals. The correspondence between this “control-*char*” notation and the actual control characters is shown in the following table. When this specification set refers to a character by its control- name, it is referring to the actual control character shown in the Value column of the table, which is not necessarily the exact control key sequence on all terminals. Some terminals have keyboards that do not allow the direct transmission of all the non-alphanumeric characters shown. In such cases, the system documentation will describe which data sequences transmitted by the terminal are interpreted by the system as representing the special characters.

Name	Value	Name	Value	Name	Value
control-A	<SOH>	control-L	<FF>	control-W	<ETB>
control-B	<STX>	control-M	<CR>	control-X	<CAN>
control-C	<ETX>	control-N	<SO>	control-Y	
control-D	<EOT>	control-O	<SI>	control-Z	<SUB>
control-E	<ENQ>	control-P	<DLE>	control-[<ESC>
control-F	<ACK>	control-Q	<DC1>	control-\	<FS>
control-G	<BEL>	control-R	<DC2>	control-]	<GS>
control-H	<BS>	control-S	<DC3>	control-^	<RS>
control-I	<HT>	control-T	<DC4>	control- <u></u>	<US>
control-J	<LF>	control-U	<NAK>	control-?	
control-K	<VT>	control-V	<SYN>		

Table 8-1 Control Character Names

Note: The notation uses upper-case letters for arbitrary editorial reasons. There is no implication that the keystrokes represent control-shift-letter sequences.

General Terminal Interface

4822

4823 This chapter describes a general terminal interface that is provided to control asynchronous
 4824 communications ports. It is implementation-dependent whether it supports network
 4825 connections or synchronous ports or both.

4826 9.1 Interface Characteristics

4827 9.1.1 Opening a Terminal Device File

4828 When a terminal device file is opened, it normally causes the thread to wait until a connection is
 4829 established. In practice, application programs seldom open these files; they are opened by
 4830 special programs and become an application's standard input, output and error files.

4831 As described in *open()*, opening a terminal device file with the *O_NONBLOCK* flag clear causes
 4832 the thread to block until the terminal device is ready and available. If *CLOCAL* mode is not set,
 4833 this means blocking until a connection is established. If *CLOCAL* mode is set in the terminal, or
 4834 the *O_NONBLOCK* flag is specified in the *open()*, the *open()* function returns a file descriptor
 4835 without waiting for a connection to be established.

4836 9.1.2 Process Groups

4837 A terminal may have a foreground process group associated with it. This foreground process
 4838 group plays a special role in handling signal-generating input characters, as discussed in Section
 4839 9.1.9 on page 123.

4840 A command interpreter process supporting job control can allocate the terminal to different jobs,
 4841 or process groups, by placing related processes in a single process group and associating this
 4842 process group with the terminal. A terminal's foreground process group may be set or
 4843 examined by a process, assuming the permission requirements are met; see *tcgetpgrp()* and
 4844 *tcsetpgrp()*. The terminal interface aids in this allocation by restricting access to the terminal by
 4845 processes that are not in the current process group; see Section 9.1.4 on page 120.

4846 When there is no longer any process whose process ID or process group ID matches the process
 4847 group ID of the foreground process group, the terminal will have no foreground process group.
 4848 It is unspecified whether the terminal has a foreground process group when there is a process
 4849 whose process ID matches the foreground process ID, but whose process group ID does not. No
 4850 actions defined in this specification set, other than allocation of a controlling terminal or a
 4851 successful call to *tcsetpgrp()*, will cause a process group to become the foreground process group
 4852 of the terminal.

4853 9.1.3 The Controlling Terminal

4854 A terminal may belong to a process as its controlling terminal. Each process of a session that has
 4855 a controlling terminal has the same controlling terminal. A terminal may be the controlling
 4856 terminal for at most one session. The controlling terminal for a session is allocated by the
 4857 session leader in an implementation-dependent manner. If a session leader has no controlling
 4858 terminal, and opens a terminal device file that is not already associated with a session without
 4859 using the *O_NOCTTY* option (see *open()*), it is implementation-dependent whether the terminal
 4860 becomes the controlling terminal of the session leader. If a process which is not a session leader
 4861 opens a terminal file, or the *O_NOCTTY* option is used on *open()*, then that terminal does not

become the controlling terminal of the calling process. When a controlling terminal becomes associated with a session, its foreground process group is set to the process group of the session leader.

The controlling terminal is inherited by a child process during a *fork()* function call. A process relinquishes its controlling terminal when it creates a new session with the *setsid()* function; other processes remaining in the old session that had this terminal as their controlling terminal continue to have it. Upon the close of the last file descriptor in the system (whether or not it is in the current session) associated with the controlling terminal, it is unspecified whether all processes that had that terminal as their controlling terminal cease to have any controlling terminal. Whether and how a session leader can reacquire a controlling terminal after the controlling terminal has been relinquished in this fashion is unspecified. A process does not relinquish its controlling terminal simply by closing all of its file descriptors associated with the controlling terminal if other processes continue to have it open.

When a controlling process terminates, the controlling terminal is dissociated from the current session, allowing it to be acquired by a new session leader. Subsequent access to the terminal by other processes in the earlier session may be denied, with attempts to access the terminal treated as if a modem disconnect had been sensed.

9.1.4 Terminal Access Control

If a process is in the foreground process group of its controlling terminal, read operations are allowed, as described in Section 9.1.5. Any attempts by a process in a background process group to read from its controlling terminal cause its process group to be sent a SIGTTIN signal unless one of the following special cases applies: if the reading process is ignoring or blocking the SIGTTIN signal, or if the process group of the reading process is orphaned, the *read()* returns -1, with *errno* set to [EIO] and no signal is sent. The default action of the SIGTTIN signal is to stop the process to which it is sent. See <signal.h>.

If a process is in the foreground process group of its controlling terminal, write operations are allowed as described in Section 9.1.8 on page 122. Attempts by a process in a background process group to write to its controlling terminal will cause the process group to be sent a SIGTTOU signal unless one of the following special cases applies: if TOSTOP is not set, or if TOSTOP is set and the process is ignoring or blocking the SIGTTOU signal, the process is allowed to write to the terminal and the SIGTTOU signal is not sent. If TOSTOP is set, and the process group of the writing process is orphaned, and the writing process is not ignoring or blocking the SIGTTOU signal, the *write()* returns -1, with *errno* set to [EIO] and no signal is sent.

Certain calls that set terminal parameters are treated in the same fashion as *write()*, except that TOSTOP is ignored; that is, the effect is identical to that of terminal writes when TOSTOP is set (see Section 9.2.5 on page 129, *tcdrain()*, *tcflow()*, *tcflush()*, *tcsendbreak()* and *tcsetattr()*).

9.1.5 Input Processing and Reading Data

A terminal device associated with a terminal device file may operate in full-duplex mode, so that data may arrive even while output is occurring. Each terminal device file has an *input queue*, associated with it, into which incoming data is stored by the system before being read by a process. The system may impose a limit, {MAX_INPUT}, on the number of bytes that may be stored in the input queue. The behaviour of the system when this limit is exceeded is implementation-dependent.

Two general kinds of input processing are available, determined by whether the terminal device file is in canonical mode or non-canonical mode. These modes are described in Section 9.1.6 on page 121 and Section 9.1.7 on page 121. Additionally, input characters are processed according to the *c_iflag* (see Section 9.2.2 on page 125) and *c_lflag* (see Section 9.2.5 on page 129) fields.

Such processing can include *echoing*, which in general means transmitting input characters immediately back to the terminal when they are received from the terminal. This is useful for terminals that can operate in full-duplex mode.

The manner in which data is provided to a process reading from a terminal device file is dependent on whether the terminal file is in canonical or non-canonical mode, and on whether or not the `O_NONBLOCK` flag is set by `open()` or `fcntl()`.

If the `O_NONBLOCK` flag is clear, then the read request is blocked until data is available or a signal has been received. If the `O_NONBLOCK` flag is set, then the read request is completed, without blocking, in one of three ways:

1. If there is enough data available to satisfy the entire request, the `read()` completes successfully and returns the number of bytes read.
2. If there is not enough data available to satisfy the entire request, the `read()` completes successfully, having read as much data as possible, and returns the number of bytes it was able to read.
3. If there is no data available, the `read()` returns `-1`, with `errno` set to `[EAGAIN]`.

When data is available depends on whether the input processing mode is canonical or non-canonical. The following sections, Section 9.1.6 and Section 9.1.7 describe each of these input processing modes.

9.1.6 Canonical Mode Input Processing

In canonical mode input processing, terminal input is processed in units of lines. A line is delimited by a newline character (NL), an end-of-file character (EOF), or an end-of-line (EOL) character. See Section 9.1.9 on page 123 for more information on EOF and EOL. This means that a read request will not return until an entire line has been typed or a signal has been received. Also, no matter how many bytes are requested in the `read()` call, at most one line will be returned. It is not, however, necessary to read a whole line at once; any number of bytes, even one, may be requested in a `read()` without losing information.

If `{MAX_CANON}` is defined for this terminal device, it is a limit on the number of bytes in a line. The behaviour of the system when this limit is exceeded is implementation-dependent. If `{MAX_CANON}` is not defined, there is no such limit; see `pathconf()`.

Erase and kill processing occur when either of two special characters, the ERASE and KILL characters (see Section 9.1.9 on page 123), is received. This processing affects data in the input queue that has not yet been delimited by a newline (NL), EOF or EOL character. This undelimited data makes up the current line. The ERASE character deletes the last character in the current line, if there is one. The KILL character deletes all data in the current line, if there are any. The ERASE and KILL characters have no effect if there is no data in the current line. The ERASE and KILL characters themselves are not placed in the input queue.

9.1.7 Non-canonical Mode Input Processing

In non-canonical mode input processing, input bytes are not assembled into lines, and erase and kill processing does not occur. The values of the MIN and TIME members of the `c_cc` array are used to determine how to process the bytes received. The ISO POSIX-1 standard does not specify whether the setting of `O_NONBLOCK` takes precedence over MIN or TIME settings. Therefore, if `O_NONBLOCK` is set, `read()` may return immediately, regardless of the setting of MIN or TIME. Also, if no data is available, `read()` may either return 0, or return `-1` with `errno` set to `[EAGAIN]`.

MIN represents the minimum number of bytes that should be received when the *read()* function returns successfully. TIME is a timer of 0.1 second granularity that is used to time out bursty and short-term data transmissions. If MIN is greater than {MAX_INPUT}, the response to the request is undefined. The four possible values for MIN and TIME and their interactions are described below.

Case A: MIN > 0, TIME > 0

In this case TIME serves as an inter-byte timer and is activated after the first byte is received. Since it is an inter-byte timer, it is reset after a byte is received. The interaction between MIN and TIME is as follows. As soon as one byte is received, the inter-byte timer is started. If MIN bytes are received before the inter-byte timer expires (remember that the timer is reset upon receipt of each byte), the read is satisfied. If the timer expires before MIN bytes are received, the characters received to that point are returned to the user. Note that if TIME expires at least one byte is returned because the timer would not have been enabled unless a byte was received. In this case (MIN > 0, TIME > 0) the read blocks until the MIN and TIME mechanisms are activated by the receipt of the first byte, or a signal is received. If the data is in the buffer at the time of the *read()*, the result will be as if the data has been received immediately after the *read()*.

Case B: MIN > 0, TIME = 0

In this case, since the value of TIME is zero, the timer plays no role and only MIN is significant. A pending read is not satisfied until MIN bytes are received (that is, the pending read blocks until MIN bytes are received), or a signal is received. A program that uses this case to read record-based terminal I/O may block indefinitely in the read operation.

Case C: MIN = 0, TIME > 0

In this case, since MIN = 0, TIME no longer represents an inter-byte timer. It now serves as a read timer that is activated as soon as the *read()* function is processed. A read is satisfied as soon as a single byte is received or the read timer expires. Note that in this case if the timer expires, no bytes are returned. If the timer does not expire, the only way the read can be satisfied is if a byte is received. In this case the read will not block indefinitely waiting for a byte; if no byte is received within TIME*0.1 seconds after the read is initiated, the *read()* returns a value of zero, having read no data. If the data is in the buffer at the time of the *read()*, the timer is started as if the data has been received immediately after the *read()*.

Case D: MIN = 0, TIME = 0

The minimum of either the number of bytes requested or the number of bytes currently available is returned without waiting for more bytes to be input. If no characters are available, *read()* returns a value of zero, having read no data.

9.1.8 Writing Data and Output Processing

When a process writes one or more bytes to a terminal device file, they are processed according to the *c_oflag* field (see Section 9.2.3 on page 126). The implementation may provide a buffering mechanism; as such, when a call to *write()* completes, all of the bytes written have been scheduled for transmission to the device, but the transmission will not necessarily have completed. See *write()* for the effects of O_NONBLOCK on *write()*.

4993 9.1.9 Special Characters

4994 Certain characters have special functions on input or output or both. These functions are
4995 summarised as follows:

4996 INTR Special character on input, which is recognised if the ISIG flag is set. Generates a
4997 SIGINT signal which is sent to all processes in the foreground process group for which
4998 the terminal is the controlling terminal. If ISIG is set, the INTR character is discarded
4999 when processed.

5000 QUIT Special character on input, which is recognised if the ISIG flag is set. Generates a
5001 SIGQUIT signal which is sent to all processes in the foreground process group for
5002 which the terminal is the controlling terminal. If ISIG is set, the QUIT character is
5003 discarded when processed.

5004 ERASE Special character on input, which is recognised if the ICANON flag is set. Erases the
5005 last character in the current line; see Section 9.1.6 on page 121. It will not erase beyond
5006 the start of a line, as delimited by an NL, EOF or EOL character. If ICANON is set, the
5007 ERASE character is discarded when processed.

5008 KILL Special character on input, which is recognised if the ICANON flag is set. Deletes the
5009 entire line, as delimited by an NL, EOF or EOL character. If ICANON is set, the KILL
5010 character is discarded when processed.

5011 EOF Special character on input, which is recognised if the ICANON flag is set. When
5012 received, all the bytes waiting to be read are immediately passed to the process without
5013 waiting for a newline, and the EOF is discarded. Thus, if there are no bytes waiting
5014 (that is, the EOF occurred at the beginning of a line), a byte count of zero is returned
5015 from the *read()*, representing an end-of-file indication. If ICANON is set, the EOF
5016 character is discarded when processed.

5017 NL Special character on input, which is recognised if the ICANON flag is set. It is the line
5018 delimiter newline. It cannot be changed.

5019 EOL Special character on input, which is recognised if the ICANON flag is set. It is an
5020 additional line delimiter, like NL.

5021 SUSP If the ISIG flag is set, receipt of the SUSP character causes a SIGTSTP signal to be sent
5022 to all processes in the foreground process group for which the terminal is the
5023 controlling terminal, and the SUSP character is discarded when processed.

5024 STOP Special character on both input and output, which is recognised if the IXON (output
5025 control) or IXOFF (input control) flag is set. Can be used to suspend output
5026 temporarily. It is useful with CRT terminals to prevent output from disappearing
5027 before it can be read. If IXON is set, the STOP character is discarded when processed.

5028 START Special character on both input and output, which is recognised if the IXON (output
5029 control) or IXOFF (input control) flag is set. Can be used to resume output that has
5030 been suspended by a STOP character. If IXON is set, the START character is discarded
5031 when processed.

5032 CR Special character on input, which is recognised if the ICANON flag is set; it is the
5033 carriage-return character. When ICANON and ICRNL are set and IGNCR is not set,
5034 this character is translated into an NL, and has the same effect as an NL character.

5035 The NL and CR characters cannot be changed. It is implementation-dependent whether the
5036 START and STOP characters can be changed. The values for INTR, QUIT, ERASE, KILL, EOF,
5037 FIPS EOL and SUSP are changeable to suit individual tastes. Special character functions associated
5038 with changeable special control characters can be disabled individually.

5039 If two or more special characters have the same value, the function performed when that
 5040 character is received is undefined.

5041 A special character is recognised not only by its value, but also by its context; for example, an
 5042 implementation may support multi-byte sequences that have a meaning different from the
 5043 meaning of the bytes when considered individually. Implementations may also support
 5044 additional single-byte functions. These implementation-dependent multi-byte or single-byte
 5045 functions are recognised only if the IEXTEN flag is set; otherwise, data is received without
 5046 interpretation, except as required to recognise the special characters defined in this section.

5047 EX If IEXTEN is set, the ERASE, KILL and EOF characters can be escaped by a preceding \
 5048 character, in which case no special function occurs.

5049 **9.1.10 Modem Disconnect**

5050 If a modem disconnect is detected by the terminal interface for a controlling terminal, and if
 5051 CLOCAL is not set in the **c_cflag** field for the terminal (see Section 9.2.4 on page 128), the
 5052 SIGHUP signal is sent to the controlling process for which the terminal is the controlling
 5053 terminal. Unless other arrangements have been made, this causes the controlling process to
 5054 terminate (see *exit()*). Any subsequent read from the terminal device returns the value of zero,
 5055 indicating end-of-file. (See *read()*.) Thus, processes that read a terminal file and test for end-of-
 5056 file can terminate appropriately after a disconnect. If the EIO condition as specified in *read()*
 5057 also exists, it is unspecified whether on EOF condition or the [EIO] is returned. Any subsequent
 5058 *write()* to the terminal device returns -1, with *errno* set to [EIO], until the device is closed.

5059 **9.1.11 Closing a Terminal Device File**

5060 The last process to close a terminal device file causes any output to be sent to the device and any
 5061 input to be discarded. If HUPCL is set in the control modes and the communications port
 5062 supports a disconnect function, the terminal device will perform a disconnect.

9.2 Parameters that Can be Set

9.2.1 The `termios` Structure

Routines that need to control certain terminal I/O characteristics do so by using the **`termios`** structure as defined in the header `<termios.h>`. The members of this structure include (but are not limited to):

Member Type	Array Size	Member Name	Description
<code>tcflag_t</code>	NCCS	<code>c_iflag</code>	Input modes.
<code>tcflag_t</code>		<code>c_oflag</code>	Output modes.
<code>tcflag_t</code>		<code>c_cflag</code>	Control modes.
<code>tcflag_t</code>		<code>c_lflag</code>	Local modes.
<code>cc_t</code>		<code>c_cc []</code>	Control characters.

The types `tcflag_t` and `cc_t` are defined in the header `<termios.h>`. They are unsigned integral types.

9.2.2 Input Modes

Values of the `c_iflag` field describe the basic terminal input control, and are composed of the bitwise inclusive OR of the masks shown, which will be bitwise distinct. The mask name symbols in this table are defined in `<termios.h>`:

Mask Name	Description
BRKINT	Signal interrupt on break.
ICRNL	Map CR to NL on input.
IGNBRK	Ignore break condition.
IGNCR	Ignore CR.
IGNPAR	Ignore characters with parity errors.
INLCR	Map NL to CR on input.
INPCK	Enable input parity check.
ISTRIP	Strip character.
IUCLC	Map upper case to lower case on input. (LEGACY)
IXANY	Enable any character to restart output.
IXOFF	Enable start/stop input control.
IXON	Enable start/stop output control.
PARMRK	Mark parity errors.

In the context of asynchronous serial data transmission, a break condition is defined as a sequence of zero-valued bits that continues for more than the time to send one byte. The entire sequence of zero-valued bits is interpreted as a single break condition, even if it continues for a time equivalent to more than one byte. In contexts other than asynchronous serial data transmission, the definition of a break condition is implementation-dependent.

If `IGNBRK` is set, a break condition detected on input is ignored that is, not put on the input queue and therefore not read by any process. If `IGNBRK` is not set and `BRKINT` is set, the break condition will flush the input and output queues, and if the terminal is the controlling terminal of a foreground process group, the break condition will generate a single `SIGINT` signal to that foreground process group. If neither `IGNBRK` nor `BRKINT` is set, a break condition is read as a single `0x00`, or if `PARMRK` is set, as `0xff 0x00 0x00`.

5108		If IGNPAR is set, a byte with a framing or parity error (other than break) is ignored.
5109		If PARMRK is set, and IGNPAR is not set, a byte with a framing or parity error (other than
5110		break) is given to the application as the three-byte sequence 0xff 0x00 X, where 0xff 0x00 is a
5111		two-byte flag preceding each sequence and X is the data of the byte received in error. To avoid
5112		ambiguity in this case, if ISTRIP is not set, a valid byte of 0xff is given to the application as 0xff
5113		0xff. If neither PARMRK nor IGNPAR is set, a framing or parity error (other than break) is given
5114		to the application as a single byte 0x00.
5115		If INPCK is set, input parity checking is enabled. If INPCK is not set, input parity checking is
5116		disabled, allowing output parity generation without input parity errors. Note that whether
5117		input parity checking is enabled or disabled is independent of whether parity detection is
5118		enabled or disabled (see Section 9.2.4 on page 128). If parity detection is enabled but input parity
5119		checking is disabled, the hardware to which the terminal is connected will recognise the parity
5120		bit but the terminal special file will not check whether or not this bit is correctly set.
5121		If ISTRIP is set, valid input bytes are first stripped to seven bits, otherwise all eight bits are
5122		processed.
5123		If INLCR is set, a received NL character is translated into a CR character. If IGNCR is set, a
5124		received CR character is ignored (not read). If IGNCR is not set and ICRNL is set, a received CR
5125		character is translated into an NL character.
5126	EX	If IUCLC is set, upper- to lower-case mappings are performed on the received character. In
5127		locales other than the POSIX locale, the mapping is unspecified. (LEGACY)
5128		If IXANY is set, any input character will restart output that has been suspended.
5129		If IXON is set, start/stop output control is enabled. A received STOP character suspends output
5130		and a received START character restarts output. When IXON is set, START and STOP characters
5131		are not read, but merely perform flow control functions. When IXON is not set, the START and
5132		STOP characters are read.
5133		If IXOFF is set, start/stop input control is enabled. The system transmits STOP characters,
5134		which are intended to cause the terminal device to stop transmitting data, as needed to prevent
5135		the input queue from overflowing and causing undefined behaviour, and transmits START
5136		characters, which are intended to cause the terminal device to resume transmitting data, as soon
5137		as the device can continue transmitting data without risk of overflowing the input queue. The
5138		precise conditions under which STOP and START characters are transmitted are
5139		implementation-dependent.
5140		The initial input control value after <i>open()</i> is implementation-dependent.

5141 9.2.3 Output Modes

5142 The **c_oflag** field specifies the terminal interface's treatment of output, and is composed of the
5143 bitwise inclusive OR of the masks shown, which will be bitwise distinct. The mask name
5144 symbols in this table are defined in **<termios.h>**:

5145		
5146		
5147		
5148	EX	
5149		
5150		
5151		
5152		
5153		
5154		
5155		
5156		
5157		
5158		
5159		
5160		
5161		
5162		
5163		
5164		
5165		
5166		
5167		
5168		
5169		
5170		
5171		
5172		
5173		
5174		
5175		
5176		

Mask Name	Description
OPOST	Perform output processing.
OLCUC	Map lower case to upper on output. (LEGACY)
ONLCR	Map NL to CR-NL on output.
OCRNL	Map CR to NL on output.
ONOCR	No CR output at column 0.
ONLRET	NL performs CR function.
OFILL	Use fill characters for delay.
OFDEL	Fill is DEL, else NUL.
NLDLY	Select newline delays:
NL0	Newline character type 0
NL1	Newline character type 1.
CRDLY	Select carriage-return delays:
CR0	Carriage-return delay type 0
CR1	Carriage-return delay type 1
CR2	Carriage-return delay type 2
CR3	Carriage-return delay type 3.
TABDLY	Select horizontal-tab delays:
TAB0	Horizontal-tab delay type 0
TAB1	Horizontal-tab delay type 1
TAB2	Horizontal-tab delay type 2.
TAB3	Expand tabs to spaces.
BSDLY	Select backspace delays:
BS0	Backspace-delay type 0
BS1	Backspace-delay type 1.
VTDLY	Select vertical-tab delays:
VT0	Vertical-tab delay type 0
VT1	Vertical-tab delay type 1.
FFDLY	Select form-feed delays:
FF0	Form-feed delay type 0
FF1	Form-feed delay type 1.

5177 If OPOST is set, output data is post-processed as described below, so that lines of text are
5178 modified to appear appropriately on the terminal device; otherwise, characters are transmitted
5179 without change.

5180 EX If OLCUC is set, lower- to upper-case mappings are performed on the characters before they are
5181 transmitted. In locales other than the POSIX locale, the mapping is unspecified. (**LEGACY**)

5182 If ONLCR is set, the NL character is transmitted as the CR-NL character pair. If OCRNL is set,
5183 the CR character is transmitted as the NL character. If ONOCR is set, no CR character is
5184 transmitted when at column 0 (first position). If ONLRET is set, the NL character is assumed to
5185 do the carriage-return function; the column pointer will be set to 0 and the delays specified for
5186 CR will be used. Otherwise the NL character is assumed to do just the line-feed function; the
5187 column pointer will remain unchanged. The column pointer is also set to 0 if the CR character is
5188 actually transmitted.

5189 The delay bits specify how long transmission stops to allow for mechanical or other movement
5190 when certain characters are sent to the terminal. In all cases a value of 0 indicates no delay. If
5191 OFILL is set, fill characters will be transmitted for delay instead of a timed delay. This is useful
5192 for high baud rate terminals which need only a minimal delay. If OFDEL is set, the fill character
5193 is DEL, otherwise NUL.

If a form-feed or vertical-tab delay is specified, it lasts for about 2 seconds.

New-line delay lasts about 0.10 seconds. If ONLRET is set, the carriage-return delays are used instead of the newline delays. If OFILL is set, two fill characters will be transmitted.

Carriage-return delay type 1 is dependent on the current column position, type 2 is about 0.10 seconds, and type 3 is about 0.15 seconds. If OFILL is set, delay type 1 transmits two fill characters, and type 2, four fill characters.

Horizontal-tab delay type 1 is dependent on the current column position. Type 2 is about 0.10 seconds. Type 3 specifies that tabs are to be expanded into spaces. If OFILL is set, two fill characters will be transmitted for any delay.

Backspace delay lasts about 0.05 seconds. If OFILL is set, one fill character will be transmitted.

The actual delays depend on line speed and system load.

The initial output control value after *open()* is implementation-dependent.

9.2.4 Control Modes

The **c_cflag** field describes the hardware control of the terminal, and is composed of the bitwise inclusive OR of the masks shown, which will be bitwise distinct. The mask name symbols in this table are defined in `<termios.h>`; not all values specified are required to be supported by the underlying hardware:

Mask Name	Description
CLOCAL	Ignore modem status lines.
CREAD	Enable receiver.
CSIZE	Number of bits transmitted or received per byte:
CS5	5 bits
CS6	6 bits
CS7	7 bits
CS8	8 bits.
CSTOPB	Send two stop bits, else one.
HUPCL	Hang up on last close.
PARENB	Parity enable.
PARODD	Odd parity, else even.

In addition, the input and output baud rates are stored in the **termios** structure. The following values are supported:

Name	Description	Name	Description
B0	Hang up	B600	600 baud
B50	50 baud	B1200	1200 baud
B75	75 baud	B1800	1800 baud
B110	110 baud	B2400	2400 baud
B134	134.5 baud	B4800	4800 baud
B150	150 baud	B9600	9600 baud
B200	200 baud	B19200	19200 baud
B300	300 baud	B38400	38400 baud

The following interfaces are provided for getting and setting the values of the input and output baud rates in the **termios** structure: *cfgetispeed()*, *cfgetospeed()*, *cfsetispeed()* and *cfsetospeed()*. The effects on the terminal device do not become effective and not all errors are detected until the *tcsetattr()* function is successfully called.

5240 The CSIZE bits specify the number of transmitted or received bits per byte. If ISTRIP is not set,
 5241 the value of all the other bits is unspecified. If ISTRIP is set, the value of all but the 7 low-order
 5242 bits is zero, but the value of any other bits beyond CSIZE is unspecified when read. CSIZE does
 5243 not include the parity bit, if any. If CSTOPB is set, two stop bits are used, otherwise one stop bit.
 5244 For example, at 110 baud, two stop bits are normally used.

5245 If CREAD is set, the receiver is enabled. Otherwise, no characters will be received.

5246 If PARENB is set, parity generation and detection is enabled and a parity bit is added to each
 5247 byte. If parity is enabled, PARODD specifies odd parity if set, otherwise even parity is used.

5248 If HUPCL is set, the modem control lines for the port are lowered when the last process with the
 5249 port open closes the port or the process terminates. The modem connection is broken.

5250 If CLOCAL is set, a connection does not depend on the state of the modem status lines. If
 5251 CLOCAL is clear, the modem status lines are monitored.

5252 Under normal circumstances, a call to the *open()* function waits for the modem connection to
 5253 complete. However, if the O_NONBLOCK flag is set (see *open()*) or if CLOCAL has been set,
 5254 the *open()* function returns immediately without waiting for the connection.

5255 If the object for which the control modes are set is not an asynchronous serial connection, some
 5256 of the modes may be ignored; for example, if an attempt is made to set the baud rate on a
 5257 network connection to a terminal on another host, the baud rate may or may not be set on the
 5258 connection between that terminal and the machine to which it is directly connected.

5259 The initial hardware control value after *open()* is implementation-dependent.

5260 9.2.5 Local Modes

5261 The **c_lflag** field of the argument structure is used to control various functions. It is composed
 5262 of the bitwise inclusive OR of the masks shown, which will be bitwise distinct. The mask name
 5263 symbols in this table are defined in **<termios.h>**; not all values specified are required to be
 5264 supported by the underlying hardware:

Mask Name	Description
ECHO	Enable echo.
ECHOE	Echo ERASE as an error correcting backspace.
ECHOK	Echo KILL.
ECHONL	Echo <newline>.
ICANON	Canonical input (erase and kill processing).
IEXTEN	Enable extended (implementation-dependent) functions.
ISIG	Enable signals.
NOFLSH	Disable flush after interrupt, quit or suspend.
TOSTOP	Send SIGTTOU for background output.
XCASE	Canonical upper/lower presentation. (LEGACY)

5277 If ECHO is set, input characters are echoed back to the terminal. If ECHO is clear, input
 5278 characters are not echoed.

5279 If ECHOE and ICANON are set, the ERASE character causes the terminal to erase, if possible,
 5280 the last character in the current line from the display. If there were no character to erase, an
 5281 implementation might echo an indication that this was the case, or do nothing.

5282 If ECHOK and ICANON are set, the KILL character causes the terminal to erase the line from
 5283 the display or echoes the newline character after the KILL character.

5284		If ECHONL and ICANON are set, the newline character is echoed even if ECHO is not set.	
5285		If ICANON is set, canonical processing is enabled. This enables the erase and kill edit functions,	
5286		and the assembly of input characters into lines delimited by NL, EOF and EOL, as described in	
5287		Section 9.1.6 on page 121.	
5288		If ICANON is not set, read requests are satisfied directly from the input queue. A read is not	
5289		satisfied until at least MIN bytes have been received or the timeout value TIME expired between	
5290		bytes. The time value represents tenths of a second. See Section 9.1.7 on page 121 for more	
5291		details.	
5292		If IEXTEN is set, implementation-dependent functions are recognised from the input data. It is	
5293		implementation-dependent how IEXTEN being set interacts with ICANON, ISIG, IXON or	
5294		IXOFF. If IEXTEN is not set, implementation-dependent functions are not recognised and the	
5295		corresponding input characters are processed as described for ICANON, ISIG, IXON and IXOFF.	
5296		If ISIG is set, each input character is checked against the special control characters INTR, QUIT	
5297		and SUSP. If an input character matches one of these control characters, the function associated	
5298		with that character is performed. If ISIG is not set, no checking is done. Thus these special input	
5299		functions are possible only if ISIG is set.	
5300		If NOFLSH is set, the normal flush of the input and output queues associated with the INTR,	
5301		QUIT and SUSP characters is not done.	
5302		If TOSTOP is set, the signal SIGTTOU is sent to the process group of a process that tries to write	
5303		to its controlling terminal if it is not in the foreground process group for that terminal. This	
5304		signal, by default, stops the members of the process group. Otherwise, the output generated by	
5305		that process is output to the current output stream. Processes that are blocking or ignoring	
5306		SIGTTOU signals are excepted and allowed to produce output, and the SIGTTOU signal is not	
5307		sent.	
5308	EX	If XCASE is set, canonical lower and canonical upper presentation are performed. In locales	
5309		other than the POSIX locale, the effect is unspecified. (LEGACY)	
5310		The initial local control value after <i>open()</i> is implementation-dependent.	

5311 9.2.6 Special Control Characters

5312 The special control characters values are defined by the array `c_cc`. The subscript name and
5313 description for each element in both canonical and non-canonical modes are as follows:

Subscript Usage		Description
Canonical Mode	Non-canonical Mode	
VEOF	VINTR	EOF character
VEOL		EOL character
VERASE		ERASE character
VINTR		INTR character
VKILL		KILL character
	VMIN	MIN value
VQUIT	VQUIT	QUIT character
VSUSP	VSUSP	SUSP character
	VTIME	TIME value
VSTART	VSTART	START character
VSTOP	VSTOP	STOP character

5329 The subscript values are unique, except that the VMIN and VTIME subscripts may have the
5330 same values as the VEOF and VEOL subscripts, respectively.

5331 The number of elements in the `c_cc` array, NCCS, is unspecified.

5332 Implementations that do not support changing the START and STOP characters may ignore the
5333 character values in the `c_cc` array indexed by the VSTART and VSTOP subscripts when
5334 `tcsetattr()` is called, but will return the value in use when `tcgetattr()` is called.

5335 The initial values of all control characters are implementation-dependent.

5336 If the value of one of the changeable special control characters (see Section 9.1.9 on page 123) is
5337 `{_POSIX_VDISABLE}`, that function is disabled; that is, no input data will be recognised as the
5338 disabled special character. If ICANON is not set, the value of `{_POSIX_VDISABLE}` has no
5339 special meaning for the VMIN and VTIME entries of the `c_cc` array.

Utility Conventions

5341

5342 10.1 Utility Argument Syntax

5343 This section describes the argument syntax of the standard utilities and introduces terminology
 5344 used throughout this specification set for describing the arguments processed by the utilities.

5345 Within this specification set, a special notation is used for describing the syntax of a utility's
 5346 arguments. Unless otherwise noted, all utility descriptions use this notation, which is illustrated
 5347 by this example (see the XCU specification, **Section 2.9.1, Simple Commands**):

5348 `utility_name[-a][-b][-c option_argument][-d|-e][-f option_argument][operand ...]`

5349 The notation used for the **SYNOPSIS** sections imposes requirements on the implementors of the
 5350 standard utilities and provides a simple reference for the application developer or system user.

- 5351 1. The utility in the example is named *utility_name*. It is followed by *options*, *option-arguments*
 5352 and *operands*. The arguments that consist of hyphens and single letters or digits, such as
 5353 `-a`, are known as *options* (or, historically, *flags*). Certain options are followed by an *option-*
 5354 *argument*, as shown with `[-c option_argument]`. The arguments following the last options
 5355 and option-arguments are named *operands*.
- 5356 2. Option-arguments are sometimes shown separated from their options by blank characters,
 5357 sometimes directly adjacent. This reflects the situation that in some cases an option-
 5358 argument is included within the same argument string as the option; in most cases it is the
 5359 next argument. The Utility Syntax Guidelines in Section 10.2 on page 136 require that the
 5360 option be a separate argument from its option-argument, but there are some exceptions in
 5361 this specification set to ensure continued operation of historical applications:
 - 5362 a. If the **SYNOPSIS** of a standard utility shows a space character between an option
 5363 and option-argument (as with `[-c option_argument]` in the example), a portable
 5364 application must use separate arguments for that option and its option-argument.
 - 5365 b. If a space character is not shown (as with `[-f option_argument]` in the example), a
 5366 portable application must place an option and its option-argument directly adjacent
 5367 in the same argument string, without intervening blank characters.
 - 5368 c. Notwithstanding the preceding requirements on portable applications, X/Open
 5369 systems permit, but do not require, an application to specify options and option-
 5370 arguments as separate arguments whether or not a space character is shown on the
 5371 synopsis line, except in those cases (marked with the EX portability warning) where
 5372 an option-argument is optional and no separation can be used.
 - 5373 d. A standard utility may also be implemented to operate correctly when the required
 5374 separation into multiple arguments is violated by a non-portable application.

In summary, the following table shows allowable combinations:

	SYNOPSIS Shows:		
	<i>-a arg</i>	<i>-barg</i>	<i>-c[arg]</i>
Portable application must use:	<i>-a arg</i>	<i>-barg</i>	n/a
System will support:	<i>-a arg</i>	<i>-barg</i>	<i>-carg</i> or <i>-c</i>
System may support:	<i>-aarg</i>	<i>-b arg</i>	

- Options are usually listed in alphabetical order unless this would make the utility description more confusing. There are no implied relationships between the options based upon the order in which they appear, unless otherwise stated in the **OPTIONS** section, or unless the exception in Section 10.2 on page 136 guideline 11 applies. If an option that does not have option-arguments is repeated, the results are undefined, unless otherwise stated.
- Frequently, names of parameters that require substitution by actual values are shown with embedded underscores. Alternatively, parameters are shown as follows:

<parameter name>

The angle brackets are used for the symbolic grouping of a phrase representing a single parameter and must never be included in data submitted to the utility.

- When a utility has only a few permissible options, they are sometimes shown individually, as in the example. Utilities with many flags generally show all of the individual flags (that do not take option-arguments) grouped, as in:

utility_name [-abcDxyz][-p arg][operand]

Utilities with very complex arguments may be shown as follows:

utility_name [options][operands]

- Unless otherwise specified, whenever an operand or option-argument is, or contains, a numeric value:

- The number is interpreted as a decimal integer.
- Numerals in the range 0 to 2 147 483 647 are syntactically recognised as numeric values.
- When the utility description states that it accepts negative numbers as operands or option-arguments, numerals in the range -2 147 483 647 to 2 147 483 647 are syntactically recognised as numeric values.
- Ranges greater than those listed here are allowed.

This does not mean that all numbers within the allowable range are necessarily semantically correct. A standard utility that accepts an option-argument or operand that is to be interpreted as a number, and for which a range of values smaller than that shown above is permitted by the XCU specification, describes that smaller range along with the description of the option-argument or operand. If an error is generated, the utility's diagnostic message will indicate that the value is out of the supported range, not that it is syntactically incorrect.

For example, the specification of *dd obs=3000000000* would yield undefined behaviour for the application and could be a syntax error because the number 3 000 000 000 is outside of the range -2 147 483 647 to +2 147 483 647. On the other hand, *dd obs=2000000000* may

cause some error, such as “blocksize too large”, rather than a syntax error.

7. Arguments or option-arguments enclosed in the [and] notation are optional and can be omitted. The [and] symbols must never be included in data submitted to the utility.
8. Arguments separated by the | vertical bar notation are mutually exclusive. The | symbols must never be included in data submitted to the utility. Alternatively, mutually exclusive options and operands may be listed with multiple synopsis lines. For example:

```
utility_name -d[-a][-c option_argument][operand...]
```

```
utility_name[-a][-b][operand...]
```

When multiple synopsis lines are given for a utility, it is an indication that the utility has mutually exclusive arguments. These mutually exclusive arguments alter the functionality of the utility so that only certain other arguments are valid in combination with one of the mutually exclusive arguments. Only one of the mutually exclusive arguments is allowed for invocation of the utility. Unless otherwise stated in an accompanying **OPTIONS** section, the relationships between arguments depicted in the **SYNOPSIS** sections are mandatory requirements placed on portable applications. The use of conflicting mutually exclusive arguments produces undefined results, unless a utility description specifies otherwise. When an option is shown without the [] brackets, it means that option is required for that version of the **SYNOPSIS**. However, it is not required to be the first argument, as shown in the example above, unless otherwise stated.

The use of *undefined* for conflicting argument usage and for repeated usage of the same option is meant to prevent portable applications from using conflicting arguments or repeated options, unless specifically allowed, as is the case with *ls* (which allows simultaneous, repeated use of the *-C*, *-l* and *-1* options). Many historical implementations will tolerate this usage, choosing either the first or the last applicable argument, and this tolerance may continue, but portable applications cannot rely upon it. (Other implementations may choose to print usage messages instead.)

The use of *undefined* for conflicting argument usage also allows an implementation to make reasonable extensions to utilities where the implementor considers mutually exclusive options according to the **XCU** specification to have a sensible meaning and result.

9. Ellipses (...) are used to denote that one or more occurrences of an option or operand are allowed. When an option or an operand followed by ellipses is enclosed in brackets, zero or more options or operands can be specified. The forms:

```
utility_name -f option_argument...[operand...]
```

```
utility_name [-g option_argument]...[operand...]
```

indicate that multiple occurrences of the option and its option-argument preceding the ellipses are valid, with semantics as indicated in the **OPTIONS** section of the utility. (See also Guideline 11 in Section 10.2 on page 136.) In the first example, each option-argument requires a preceding *-f* and at least one *-f option_argument* must be given.

The **XCU** specification does not define the result of a utility when an option-argument or operand is not followed by ellipses and the application specifies more than one of that option-argument or operand. This allows an implementation to define valid (although non-standard) behaviour for the utility when more than one such option or operand are specified.

10. When the synopsis line is too long to be printed on a single line in the **XCU** specification, the indented lines following the initial line are continuation lines. An actual use of the command would appear on a single logical line.

5464 10.2 Utility Syntax Guidelines

5465 The following guidelines are established for the naming of utilities and for the specification of
 5466 options, option-arguments and operands. The *getopt()* function in the **XSH** specification assists
 5467 utilities in handling options and operands that conform to these guidelines.

5468 Operands and option-arguments can contain characters not specified in the portable character
 5469 set.

5470 The guidelines are intended to provide guidance to the authors of future utilities, such as those
 5471 written specific to a local system or that are to be components of a larger application. Some of
 5472 the standard utilities do not conform to all of these guidelines; in those cases, the **OPTIONS**
 5473 sections describe the deviations.

5474 **Guideline 1:** Utility names should be between two and nine characters, inclusive.

5475 **Guideline 2:** Utility names should include lower-case letters (the **lower** character
 5476 classification) and digits only from the portable character set.

5477 Guidelines 1 and 2 are offered as guidance for locales using Latin alphabets. |
 5478 No recommendations are made by this specification set concerning utility
 5479 naming in other locales.

5480 In the **XCU** specification, **Section 2.9.1, Simple Commands**, it is further stated
 5481 that a command used in the XSI Shell Command Language cannot be named
 5482 with a trailing colon.

5483 **Guideline 3:** Each option name should be a single alphanumeric character (the **alnum**
 5484 character classification) from the portable character set.

5485 Multi-digit options are not allowed. Instances of historical utilities that used
 5486 them have been marked **LEGACY** in the **XCU** specification, with the numbers |
 5487 being changed from option names to option-arguments.

5488 **Guideline 4:** All options should be preceded by the "-" delimiter character. |

5489 **Guideline 5:** Options without option-arguments should be accepted when grouped behind |
 5490 one "-" delimiter.

5491 **Guideline 6:** Each option and option-argument should be a separate argument, except as
 5492 noted in Section 10.1 on page 133, item (2).

5493 **Guideline 7:** Option-arguments should not be optional.

5494 **Guideline 8:** When multiple option-arguments are specified to follow a single option, they
 5495 should be presented as a single argument, using commas within that
 5496 argument or blank characters within that argument to separate them.

5497 It is up to the utility to parse a comma-separated list itself because *getopt()*
 5498 just returns a single string. This situation was retained so that certain
 5499 historical utilities would not violate the guidelines. Applications preparing
 5500 for international use should be aware of an occasional problem with comma-
 5501 separated lists: in some locales, the comma is used as the radix character.
 5502 Thus, if an application is preparing operands for a utility that expects a
 5503 comma-separated list, it should avoid generating non-integer values through
 5504 one of the means that is influenced by setting the *LC_NUMERIC* variable
 5505 (such as *awk*, *bc*, *printf* or *printf()*).

5506 **Guideline 9:** All options should precede operands on the command line.

5507 **Guideline 10:** The argument `--` should be accepted as a delimiter indicating the end of
 5508 options. Any following arguments should be treated as operands, even if they
 5509 begin with the `"-"` character. The `--` argument should not be used as an
 5510 option or as an operand.

5511 Applications calling any utility with a first operand starting with `-` should
 5512 usually specify `--`, as indicated by Guideline 10, to mark the end of the
 5513 options. This is true even if the **SYNOPSIS** in the **XCU** specification does not
 5514 specify any options; implementations may provide options as extensions to
 5515 the **XCU** specification. The standard utilities that do not support Guideline 10
 5516 indicate that fact in the **OPTIONS** section of the utility description.

5517 **Guideline 11:** The order of different options relative to one another should not matter,
 5518 unless the options are documented as mutually exclusive and such an option
 5519 is documented to override any incompatible options preceding it. If an option
 5520 that has option-arguments is repeated, the option and option-argument
 5521 combinations should be interpreted in the order specified on the command
 5522 line.

5523 The order of repeated options that also have option-arguments may be
 5524 significant; therefore, such options are required to be interpreted in the order
 5525 that they are specified. The *make* utility is an instance of a historical utility that
 5526 uses repeated options in which the order is significant. Multiple files are
 5527 specified by giving multiple instances of the `-f` option, for example:

5528 `make -f common_header -f specific_rules target`

5529 **Guideline 12:** The order of operands may matter and position-related interpretations should
 5530 be determined on a utility-specific basis.

5531 **Guideline 13:** For utilities that use operands to represent files to be opened for either reading
 5532 or writing, the `"-"` operand should be used only to mean standard input (or
 5533 standard output when it is clear from context that an output file is being
 5534 specified).

5535 Guideline 13 does not imply that all of the standard utilities automatically
 5536 accept the operand `"-"` to mean standard input or output, nor does it specify
 5537 the actions of the utility upon encountering multiple `"-"` operands. It simply
 5538 says that, by default, `"-"` operands are not used for other purposes in the file
 5539 reading or writing (but not when using *stat()*, *unlink()*, *touch*, and so forth)
 5540 utilities. All information concerning actual treatment of the `"-"` operand is
 5541 found in the individual utility sections.

5542 The utilities in the **XCU** specification that claim conformance to these guidelines were written as
 5543 if the term *should* imposed a specific requirement on their interface and applications and users
 5544 can rely on the behaviour stated here; the Guidelines are rules for the standard utilities that
 5545 claim conformance to them. On some systems, the utilities will accept usage in violation of
 5546 these guidelines for backward compatibility as well as accepting the required form.

5547 It is recommended that all future utilities and applications use these guidelines to enhance user
 5548 portability. The fact that some historical utilities could not be changed (to avoid breaking
 5549 existing applications) should not deter this future goal.

